Agenda for today

- Homework questions
- Suffix Trees
  - Introduction and definitions
  - Linear-time construction
  - Deterministic exact string matching
- Time permitting, further related topics
  - Suffix automata
  - Suffix arrays

Suffix Trees

- Given a string S of length m, a suffix tree T encodes all suffixes, i.e., $S[i, m]$ for $1 \leq i \leq m$
- Concatenation of symbols labeling edges along the path from root to leaf spells out suffix
- Leaves are numbered with i, the start index of the suffix
- Edges can be labeled with one or more symbols
- Internal (non-leaf) nodes must have more than one child
- No two children leaving a node can start with the same symbol
- Tremendously useful data structure for many problems

x t p x t d $

Start	suffix
1	x t p x t d $
2	t p x t d $
3	p x t d $

Start	suffix
4	x t d $
5	t d $
6	d $

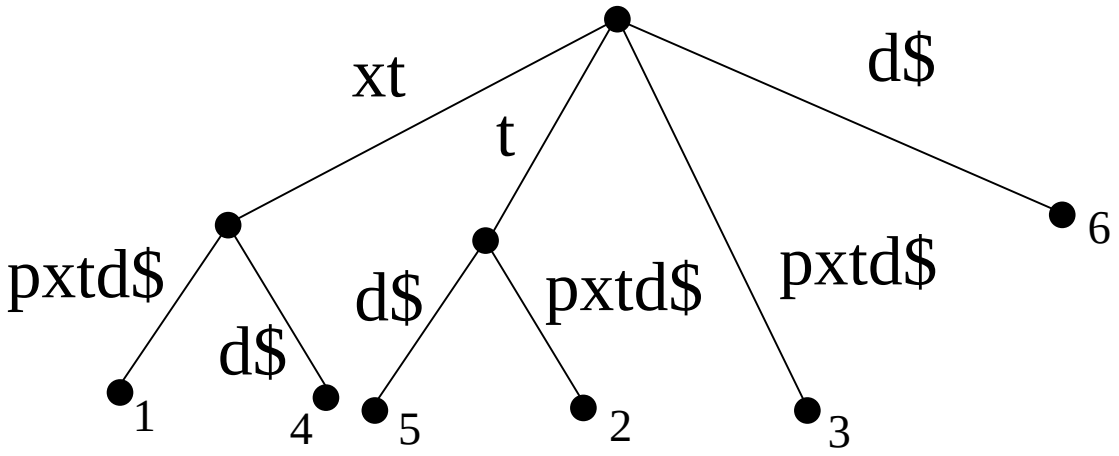

Problem: how to build

- Naive method too expensive:
  - Put suffix $S[1, m]$ into the tree
  - Then put $S[i, m]$ into the tree for $2 \leq i \leq m$
  - m^2 complexity
- Several linear-time suffix tree construction algorithms
- We will consider Ukkonen's algorithm, which has several key ideas, including
  - Suffix links (similar to failure links)
  - String indices instead of characters

Basic idea

- **Build suffix tree incrementally from left-to-right**
- **Build "implicit" suffix trees (no end-of-string marker)**
- **Extend the implicit suffix trees from the previous step**
- **Convert implicit suffix tree to explicit in final step**

Step 1: x t p x t d

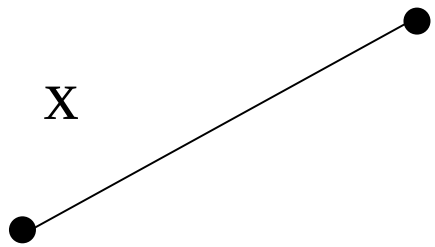

Step 2: x t p x t d

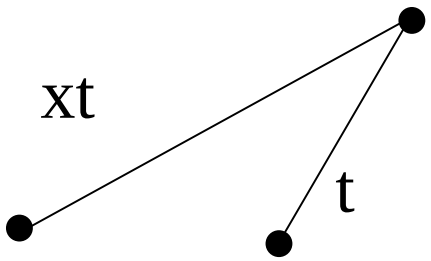

Step 3: x t p x t d

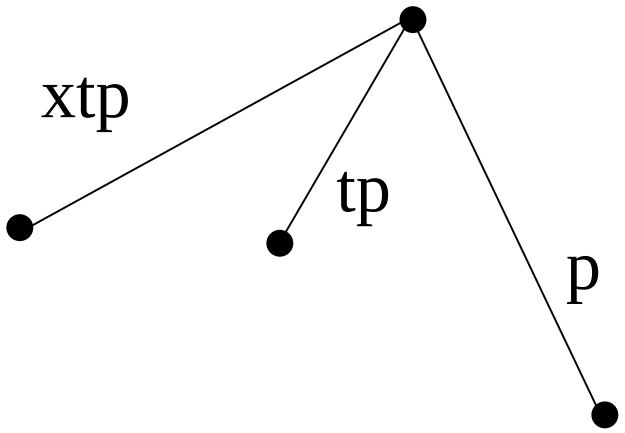

Step 4: x t p x t d

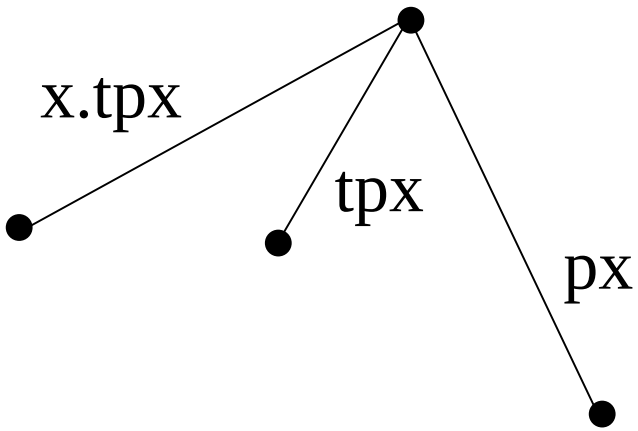

Step 5: x t p x t d

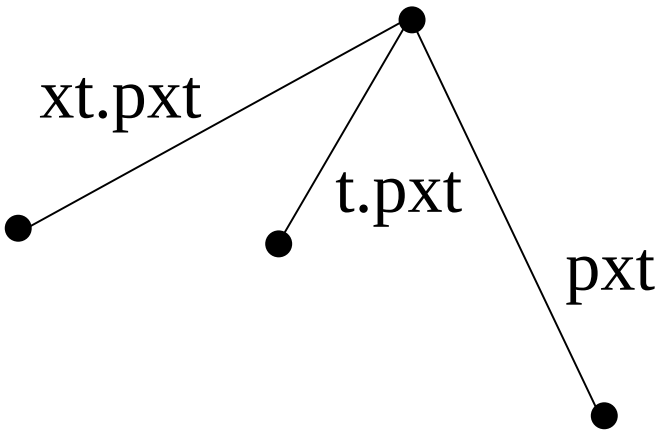

Step 6: x t p x t d

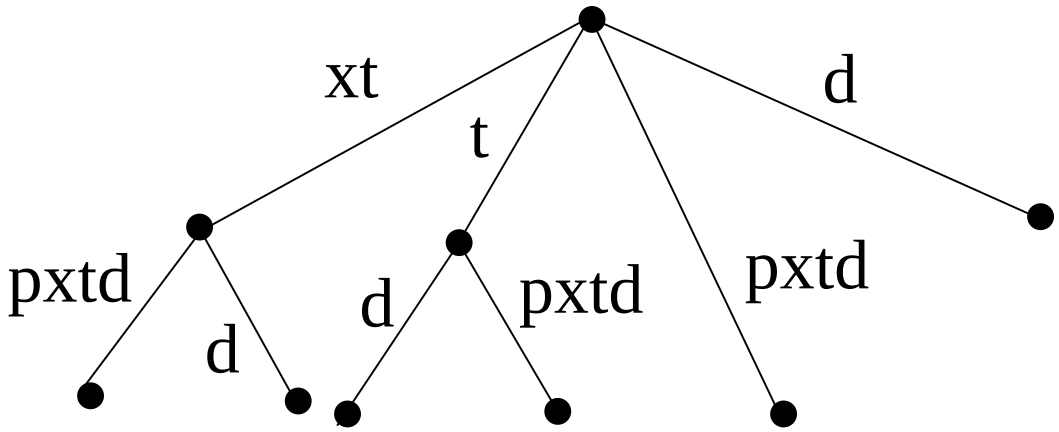

Step 7: finalize

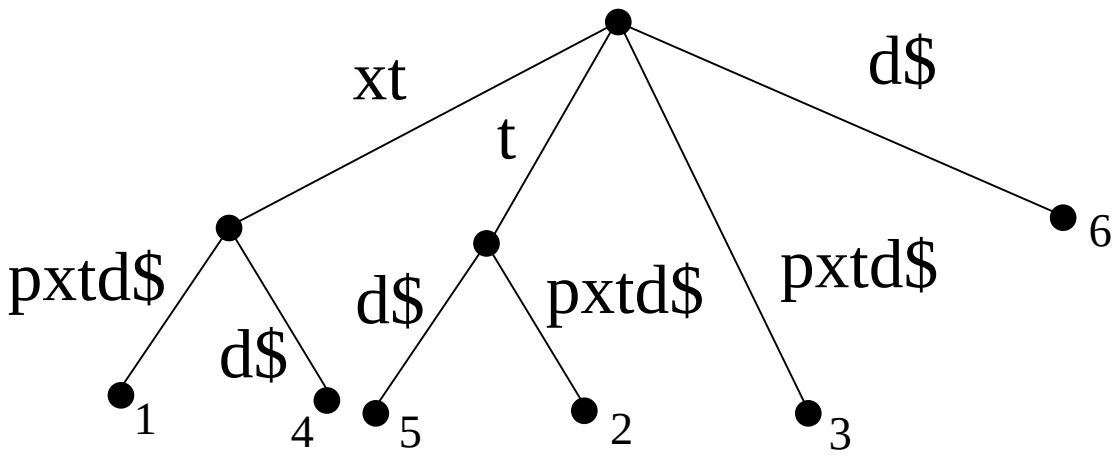

How can we use it?

- **Exact match: e.g., how many times does *xt* occur?**
- **Algorithm: order n (size of pattern) rather than order m (size of text)**
  - **Start at the root**
  - **Match pattern along branches in tree**
  - **Number of leaf nodes is the number of instances**
  - **Indices give you the locations**
- **Useful in scenarios where text is given in advance, and used for many pattern matches**

Efficient construction

- To really make this tractable, several clever "tricks"
  - Suffix links (like failure links)
  - Represent sequences as beginning and ending indices
  - Automatic extension for leaves (via global "end" index)
  - Early stopping of suffix link extensions
  - Check only first letter when following suffix link extensions ("Skip/count" trick)
- In aggregate, these make for tractable suffix tree construction

Suffix links

- **Note that every suffix of a suffix in the suffix tree must also be in the suffix tree**
- **For example: if *abcd* is in the tree, so are:**
  bcd, *cd*, *d*
- **If we extend a suffix, its suffixes will also be extended**
  - **e.g., if *abcd*+*e*, then *bcd*+*e*, *cd*+*e*, *d*+*e*, *e***
- **Need to find them (via suffix links) to extend them**
- **Can use the suffix links of parents in tree for quick access**

Suffix links: xyz

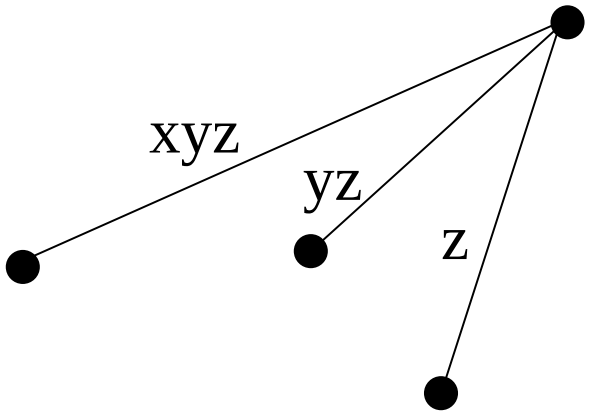

Suffix links: xyz

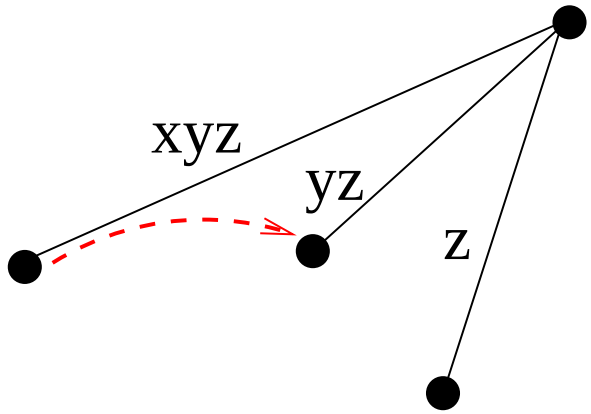

Suffix links: xyz

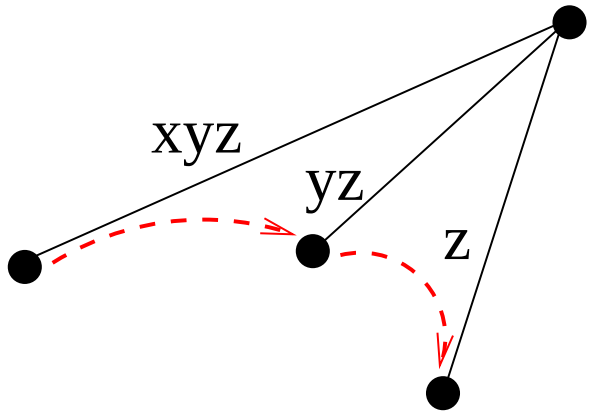

Suffix links: xyz

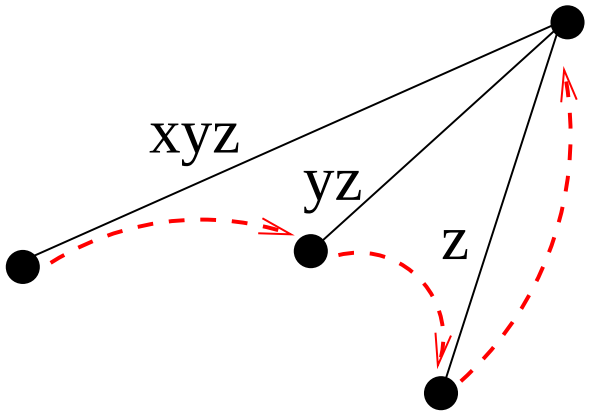

Suffix links: xyzabc

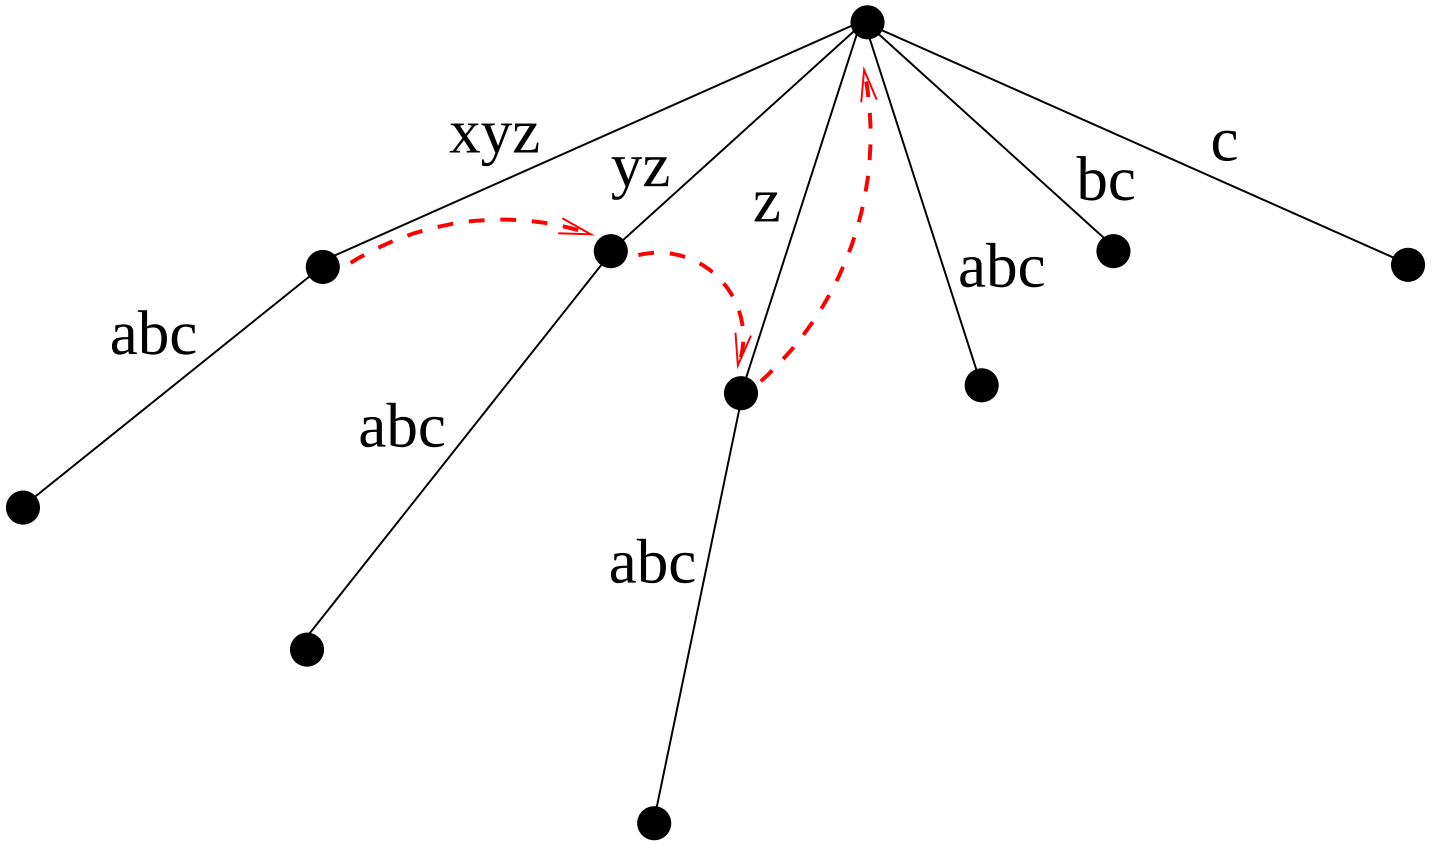

Suffix links: xyzabc

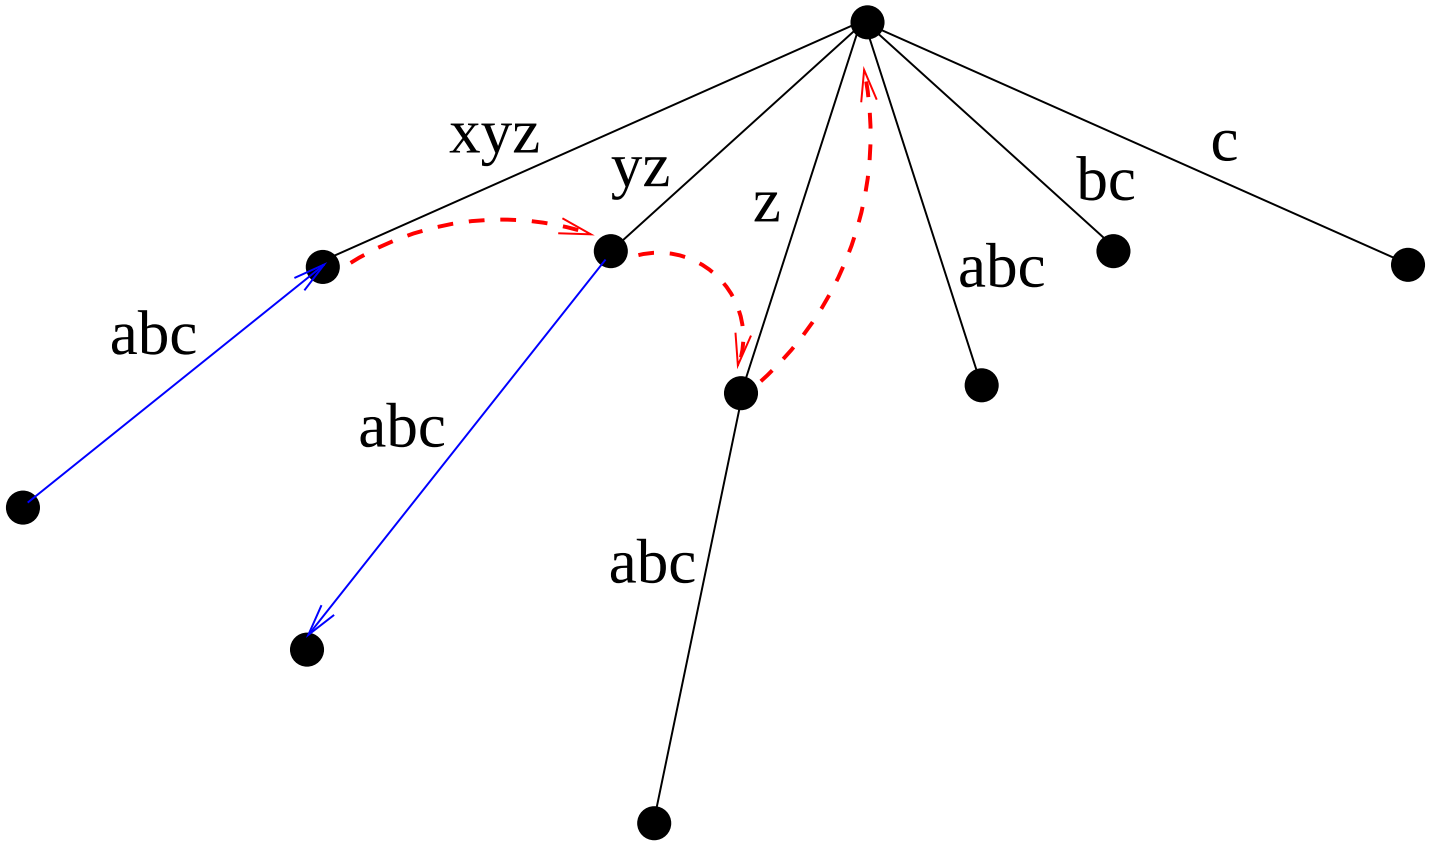

Suffix links: xyzabc

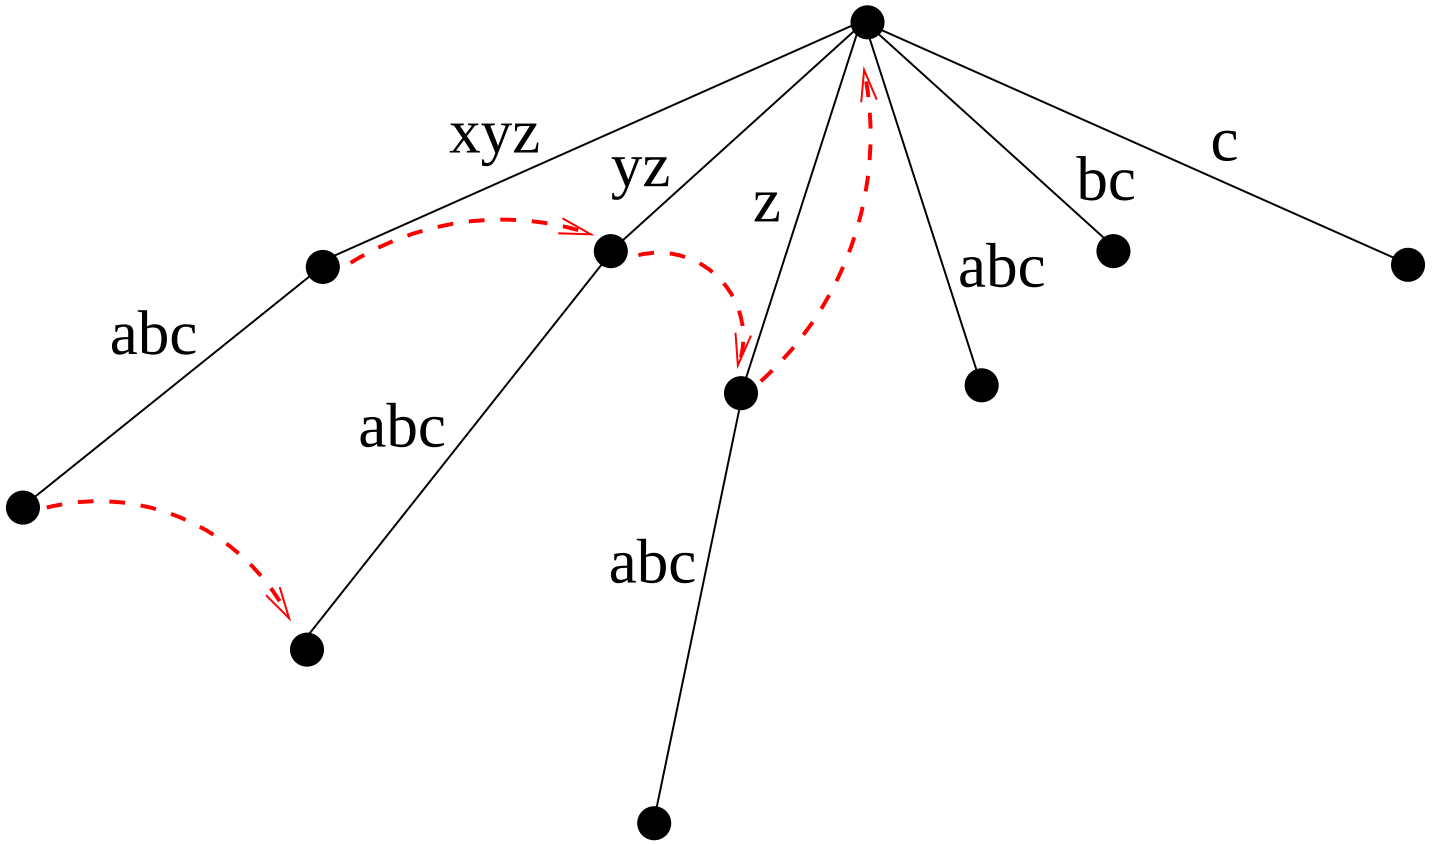

Suffix links: xyzabc

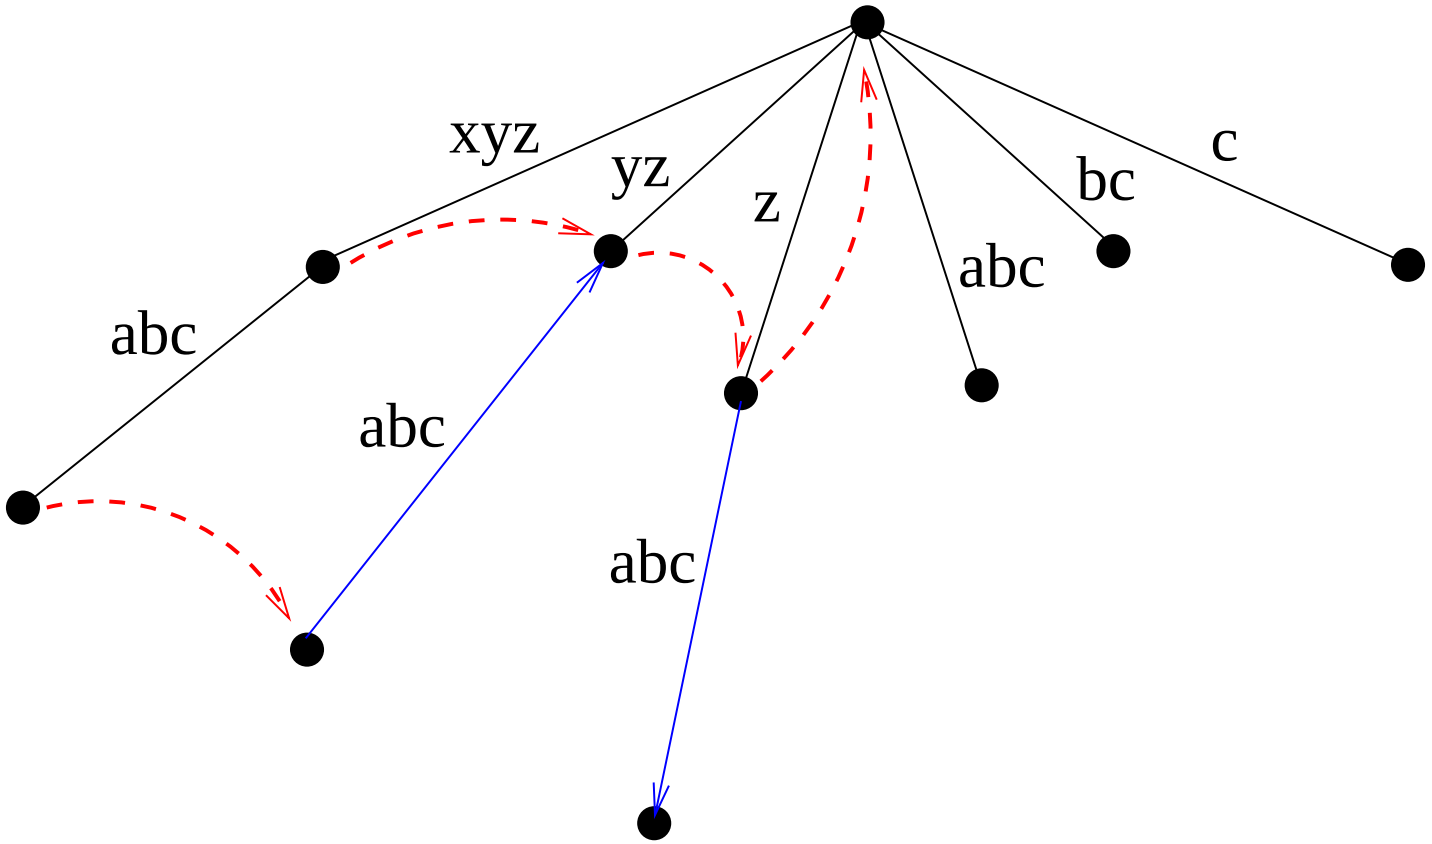

Suffix links: xyzabc

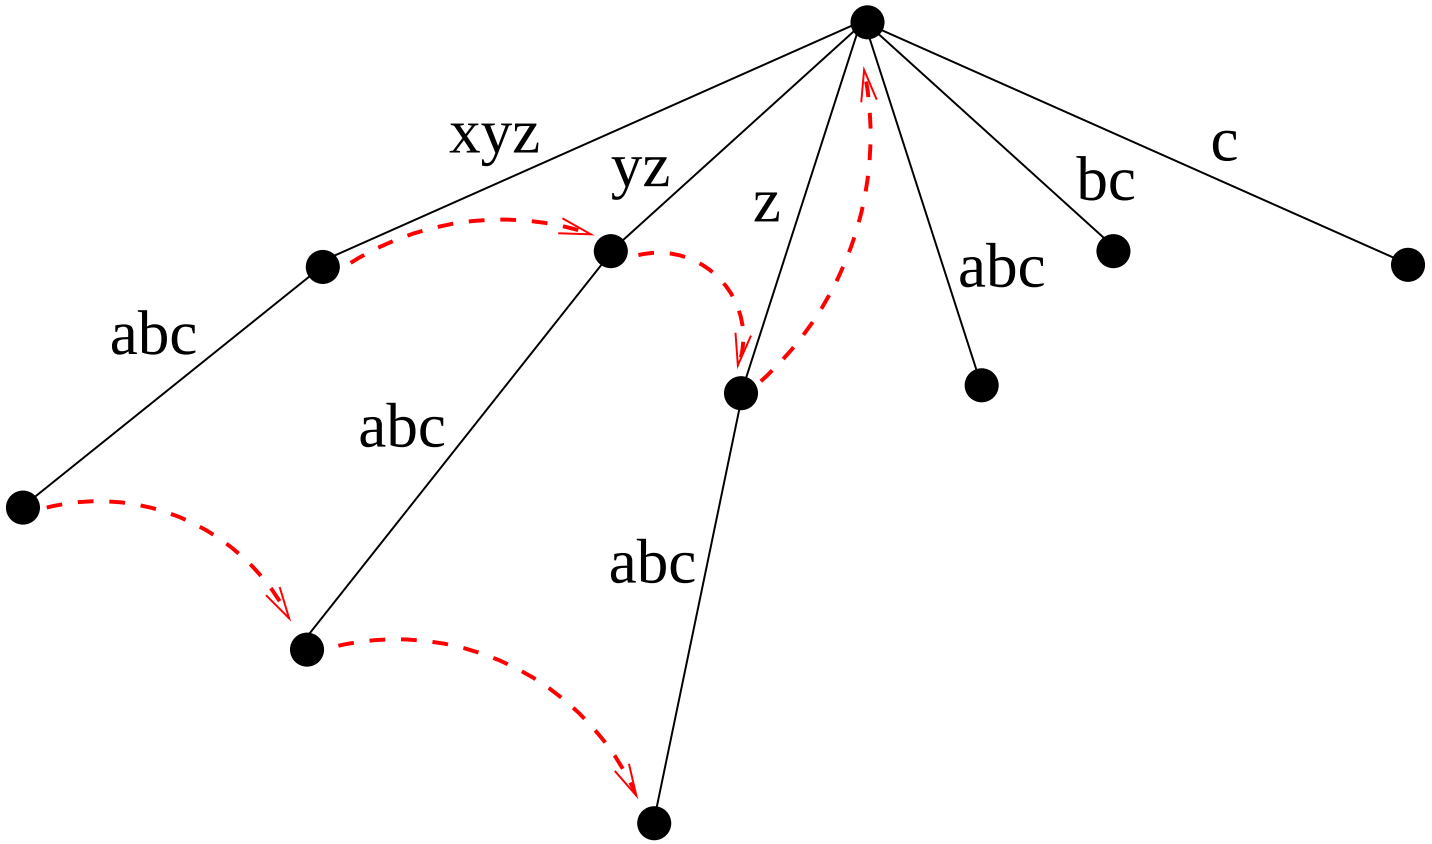

Suffix links: xyzabc

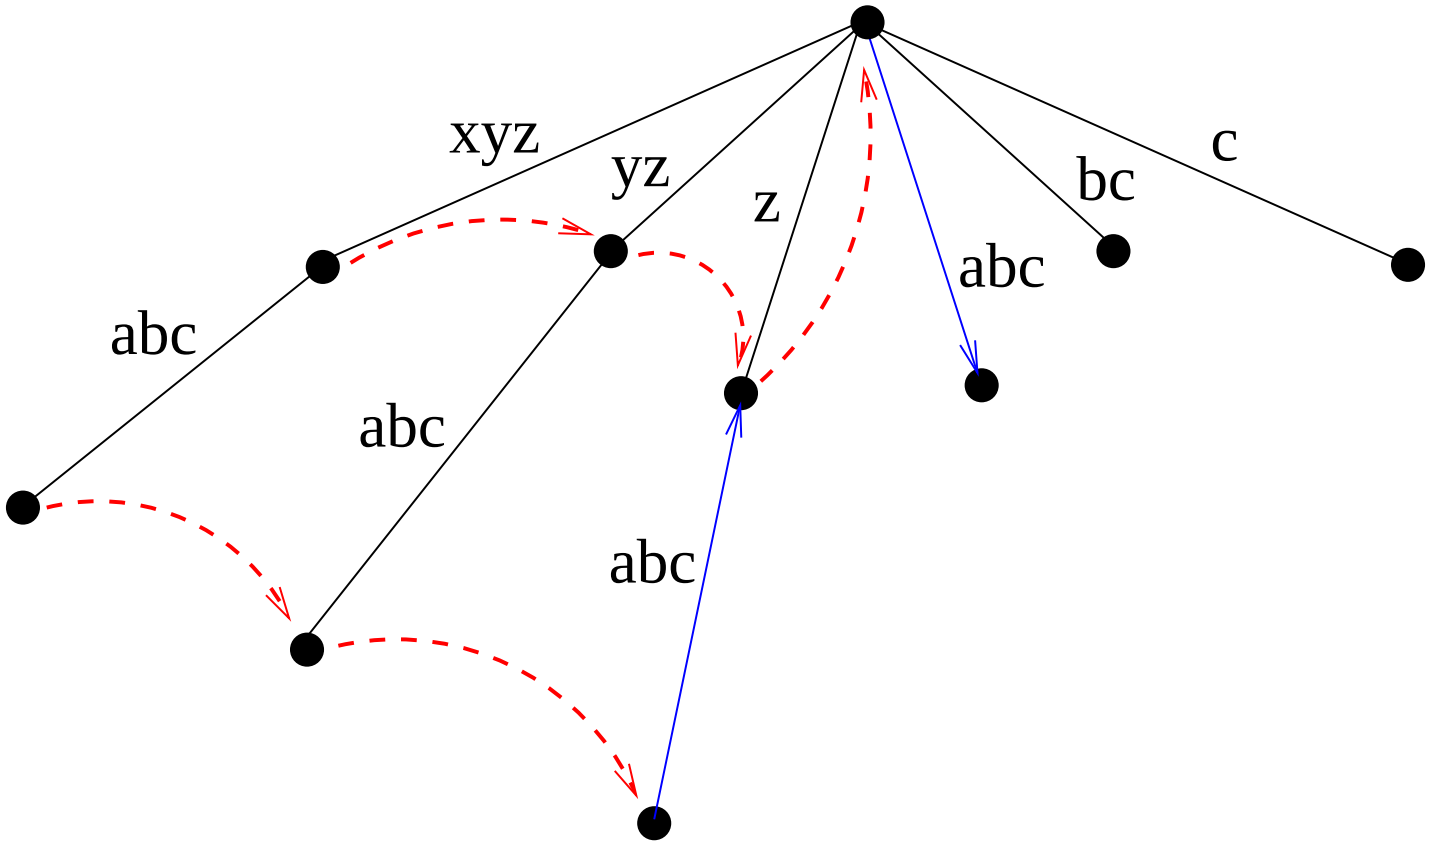

Suffix links: xyzabc

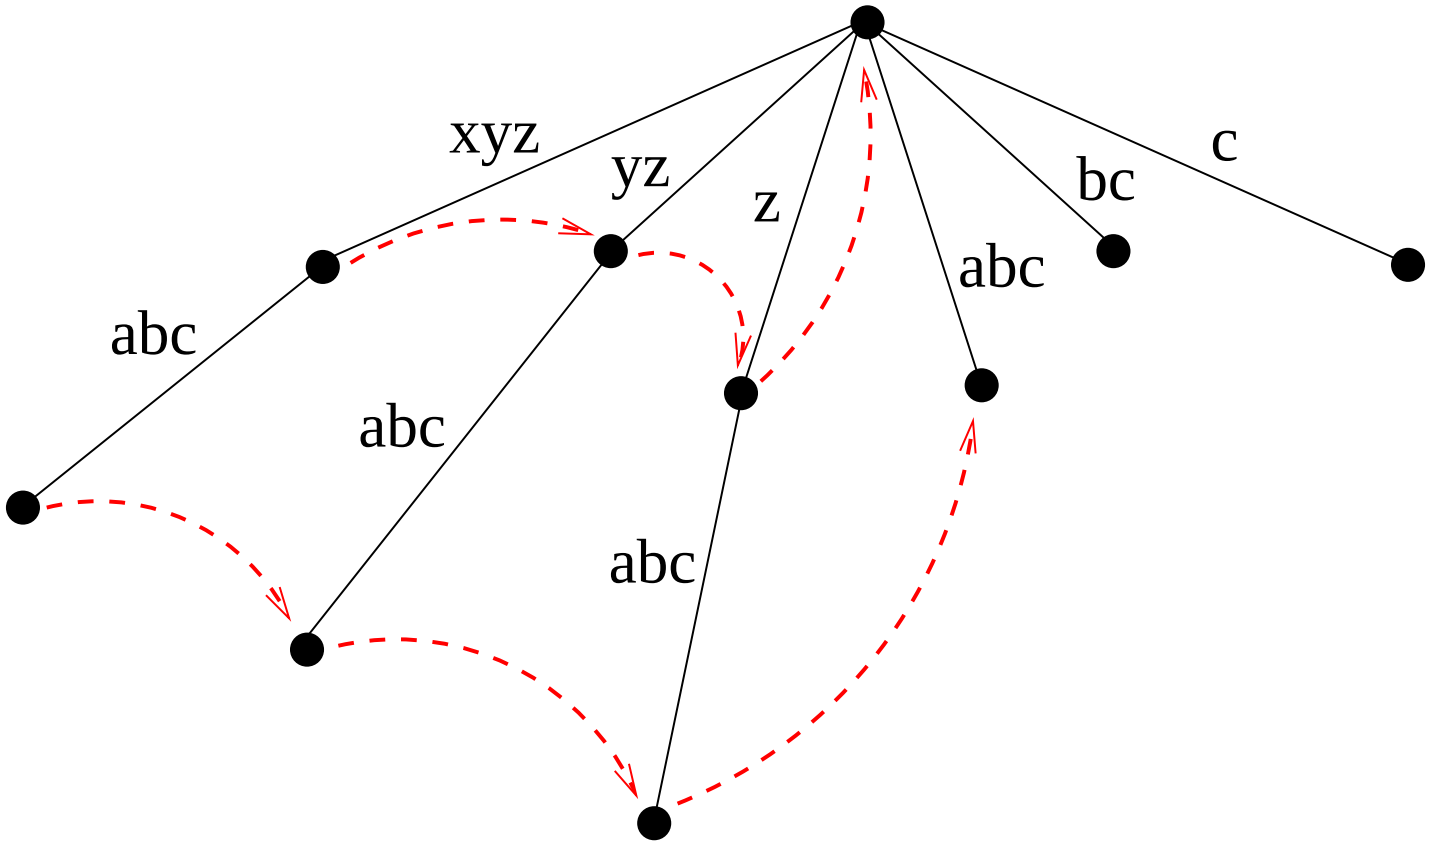

Suffix links: xyzabc

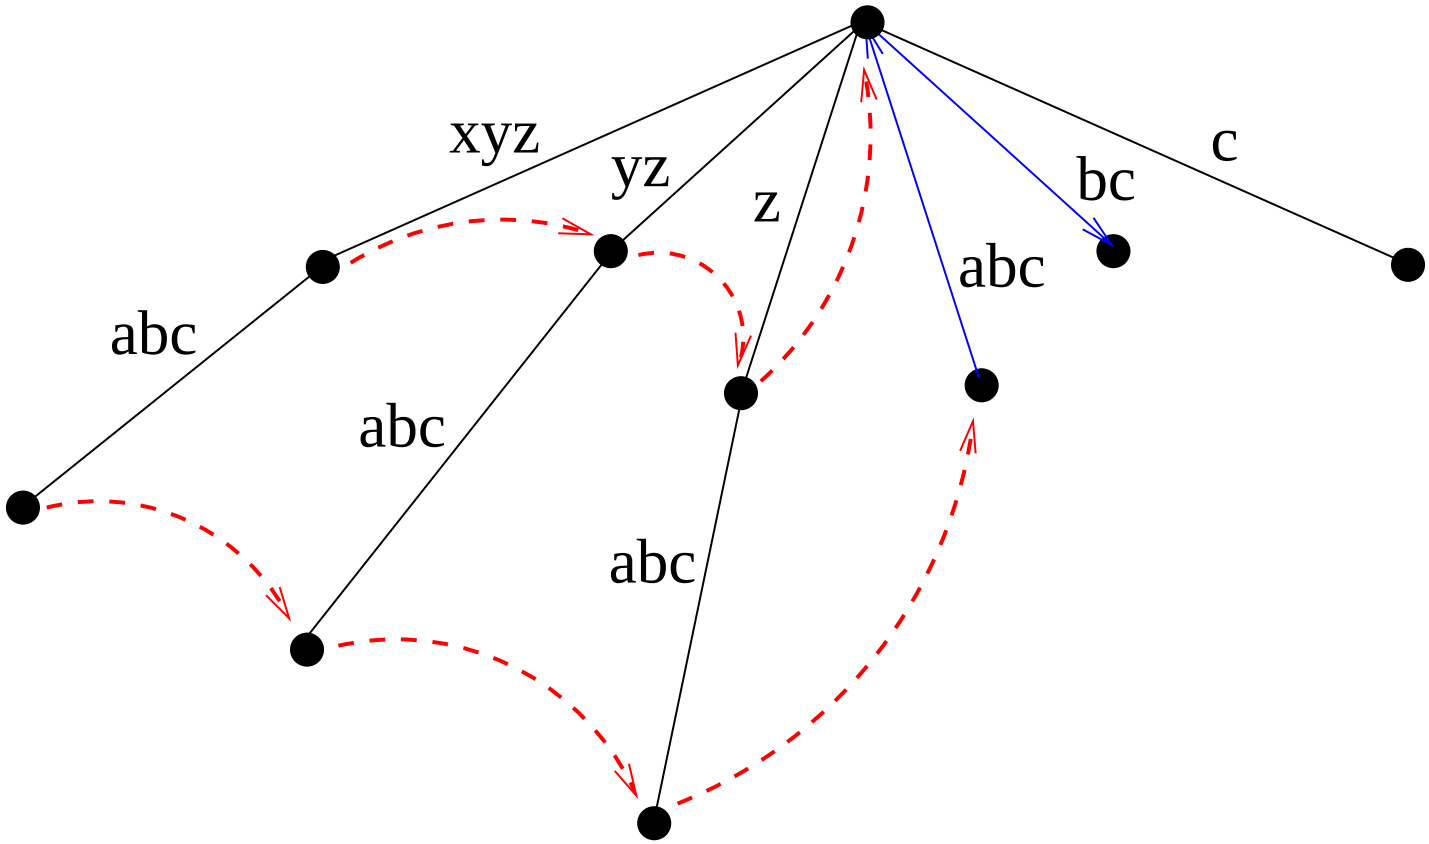

Suffix links: xyzabc

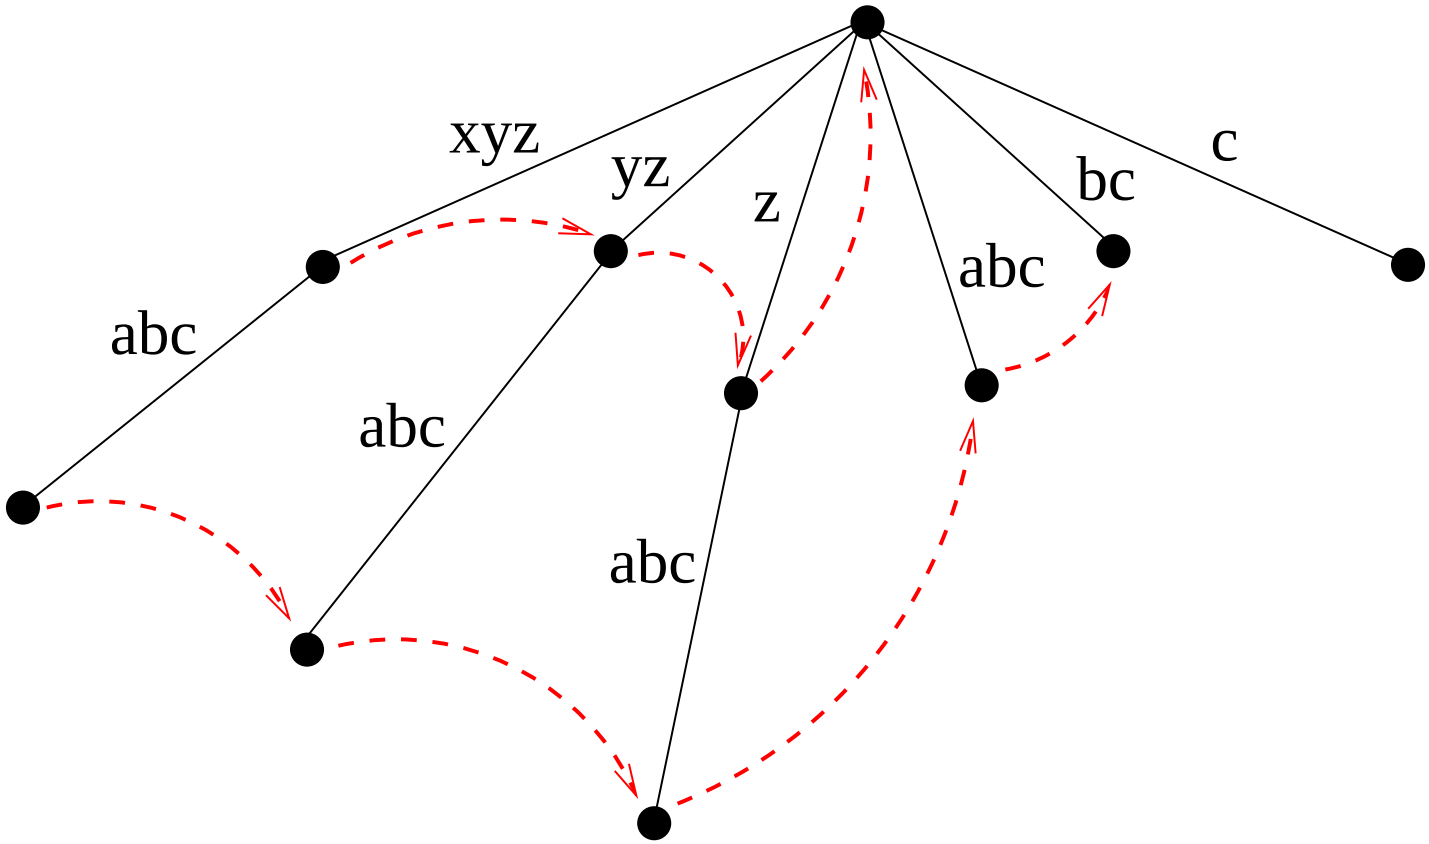

Suffix links: xyzabc

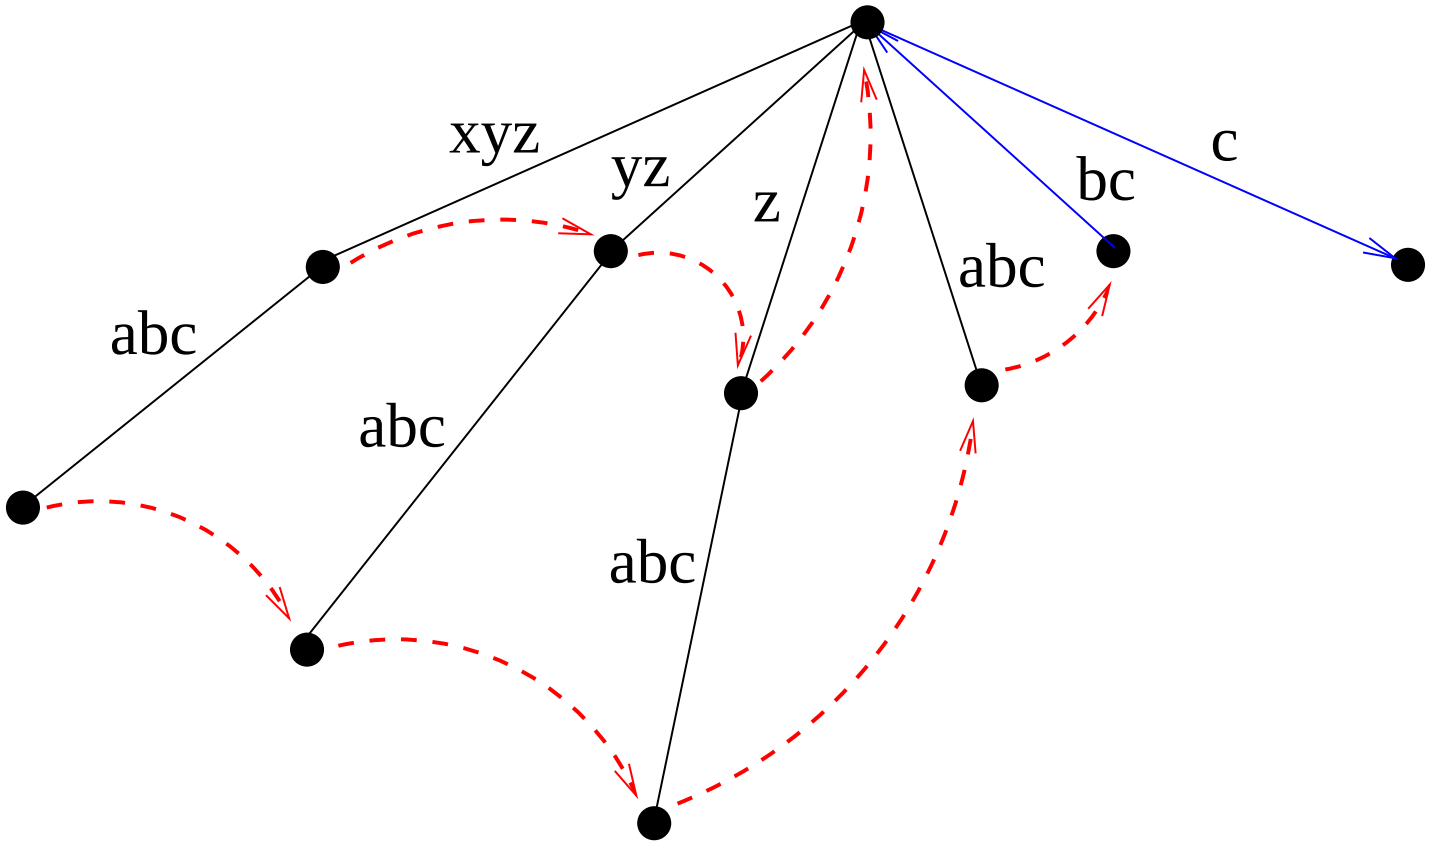

Suffix links: xyzabc

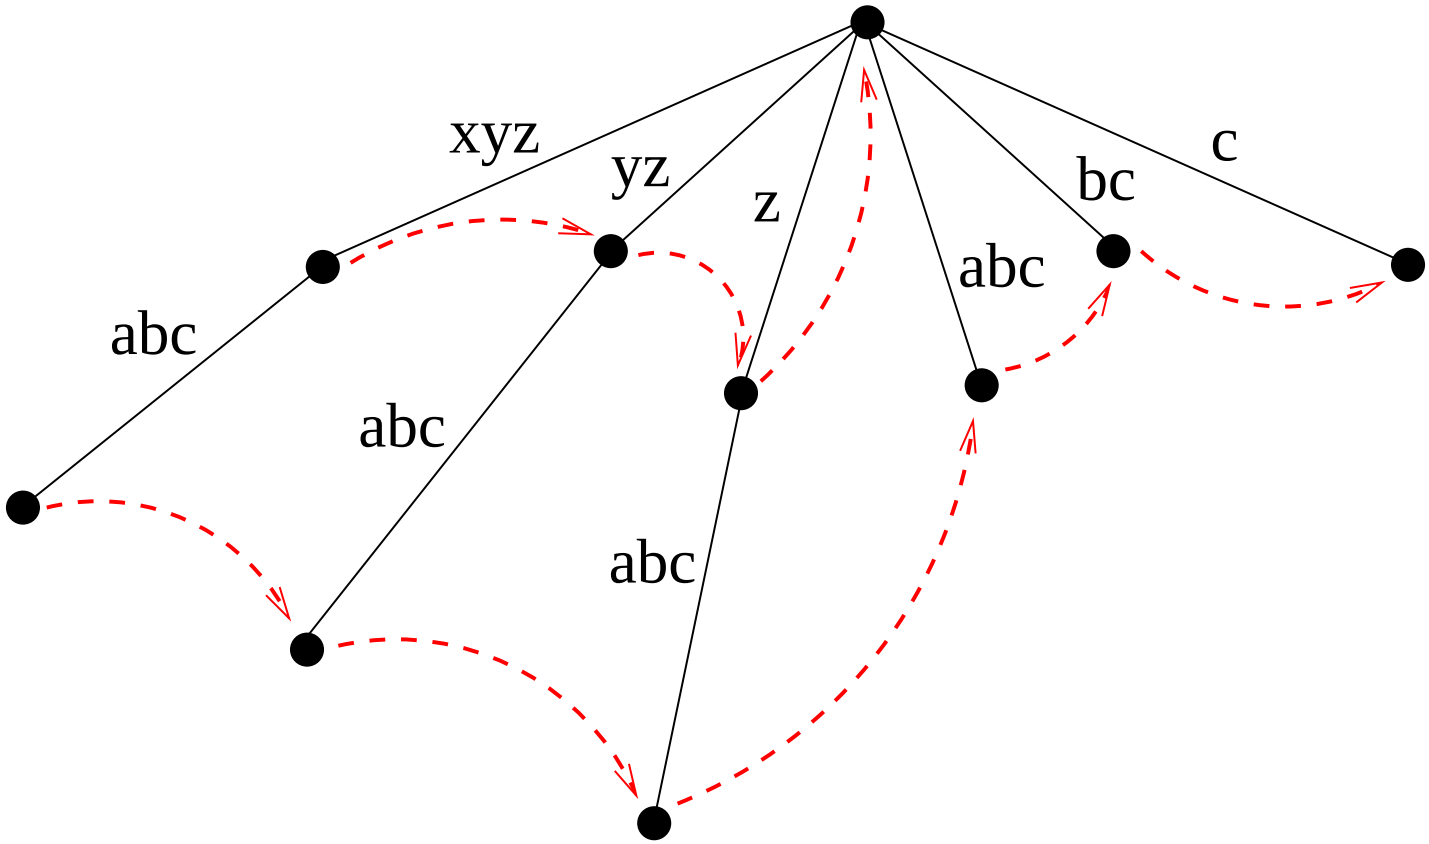

Suffix links: xyzabc

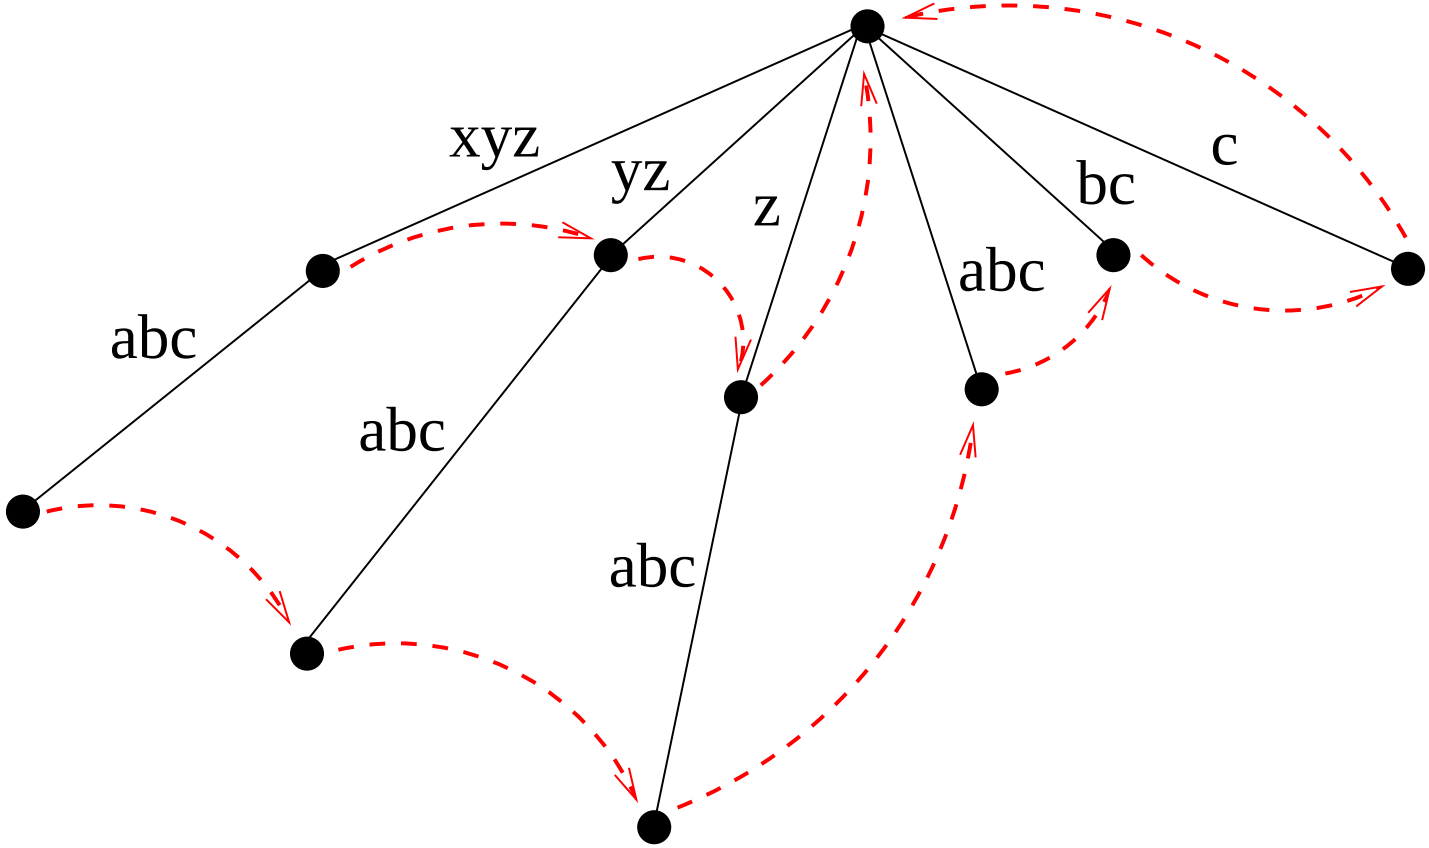

Finding suffix link for node X

- Go to parent node Y
- Denote $x\alpha$ string labeling $Y \rightarrow X$, $x \in \Sigma$, $\alpha \in \Sigma^*$
- If Y is the root of the tree
  - If $\alpha = \epsilon$, set Z to Y
  - Else set Z to state reached from Y with transitions labeled α
- Else
  - Set Y to destination state of suffix link from Y
  - Set Z to state reached from Y with transitions labeled $x\alpha$
- Create suffix link from X to Z

Some notes about suffix links

- Every previously created node (except root) has suffix link
- May require multiple transitions to generate all of string
- String is guaranteed to exist from suffix link of the parent (or root)
  - Only one transition leaving each state per letter
  - Hence only need to check first letter (skip/count)

Begin/End indices: xyzabc

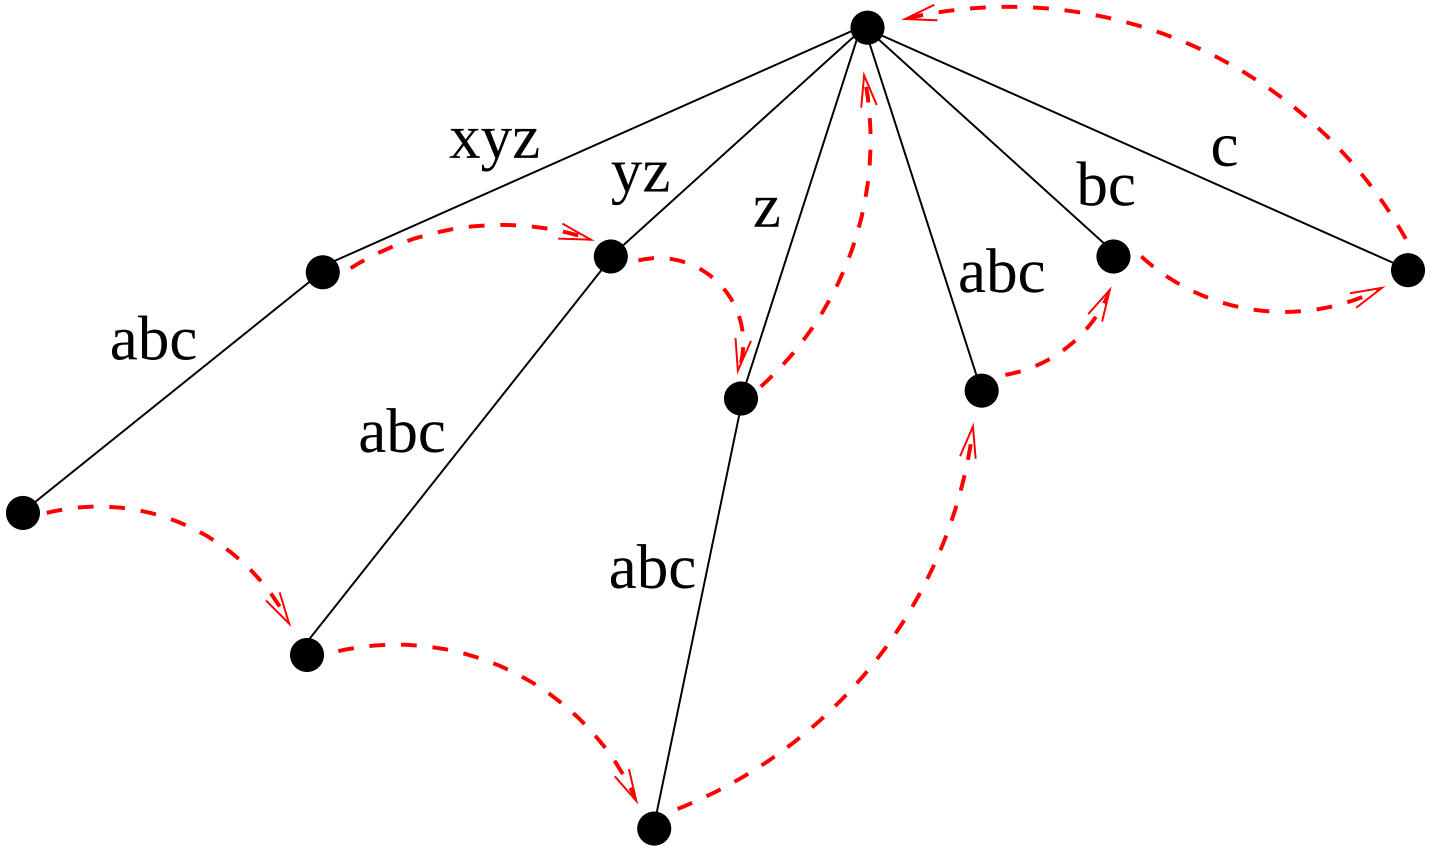

Begin/End indices: xyzabc

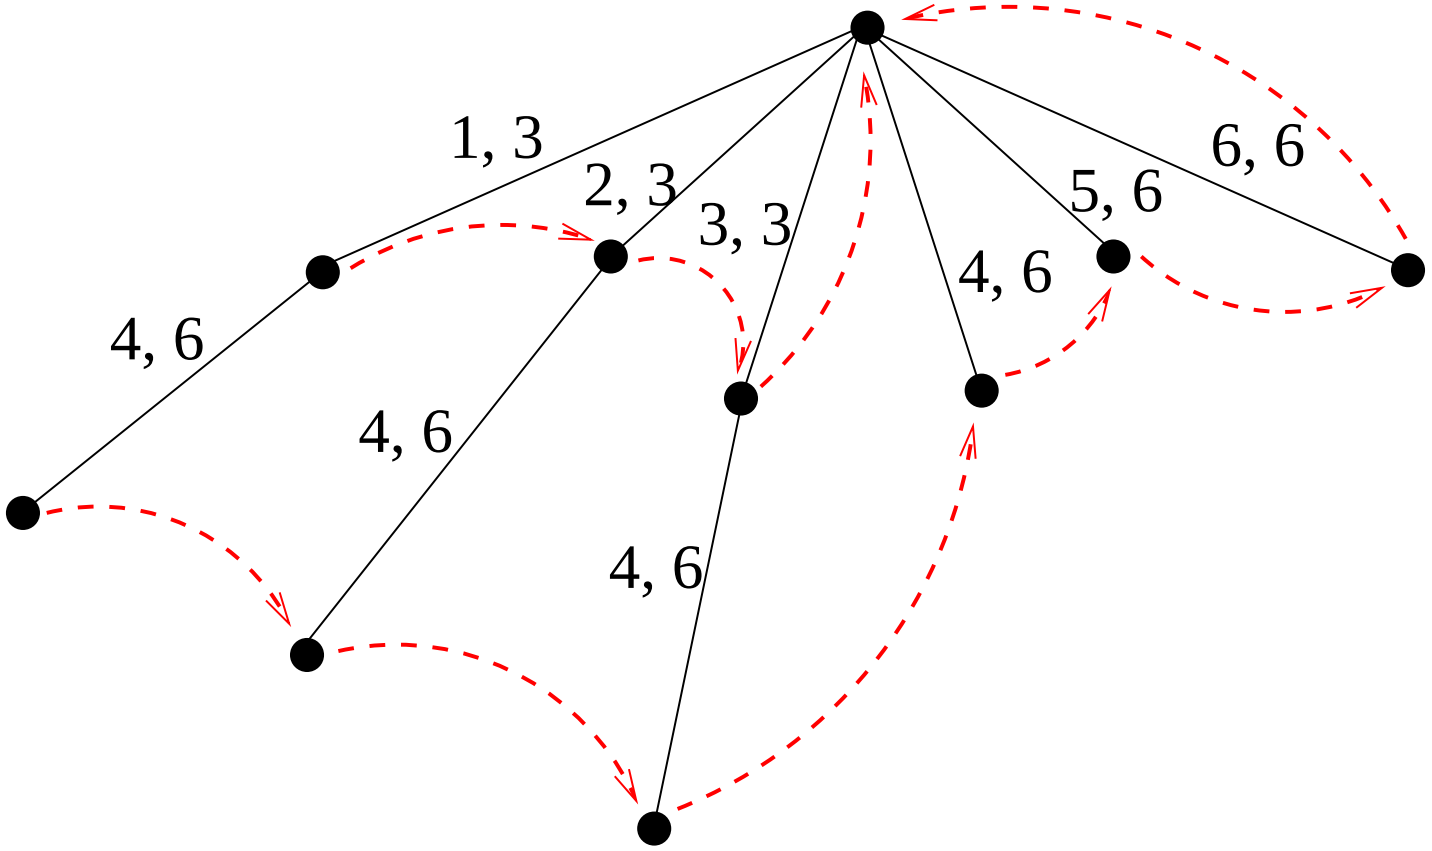

Global end index: xyzabc

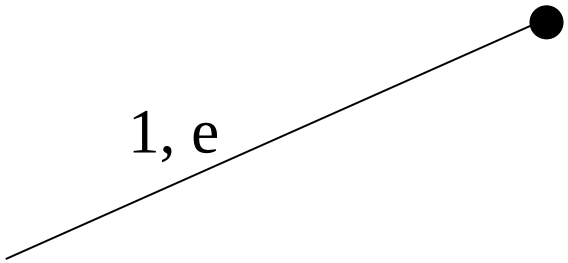

Global end index: xyzabc

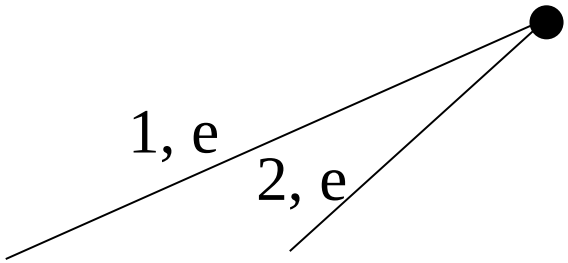

Global end index: xyzabc

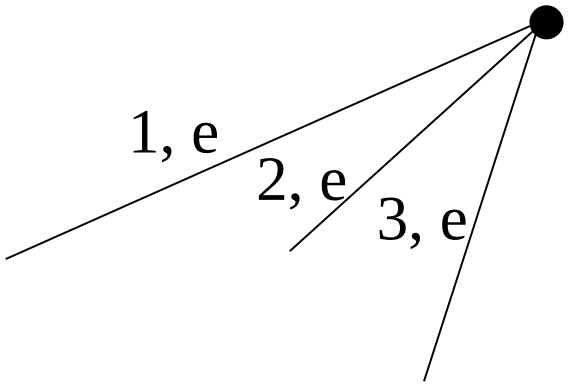

Global end index: xyzabc

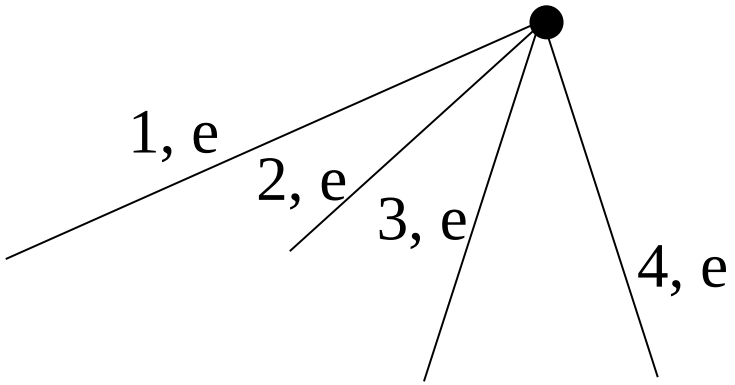

Global end index: xyzabc

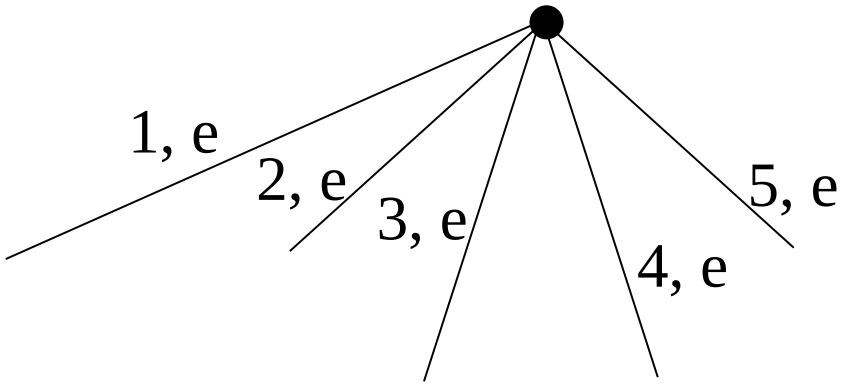

Global end index: xyzabc

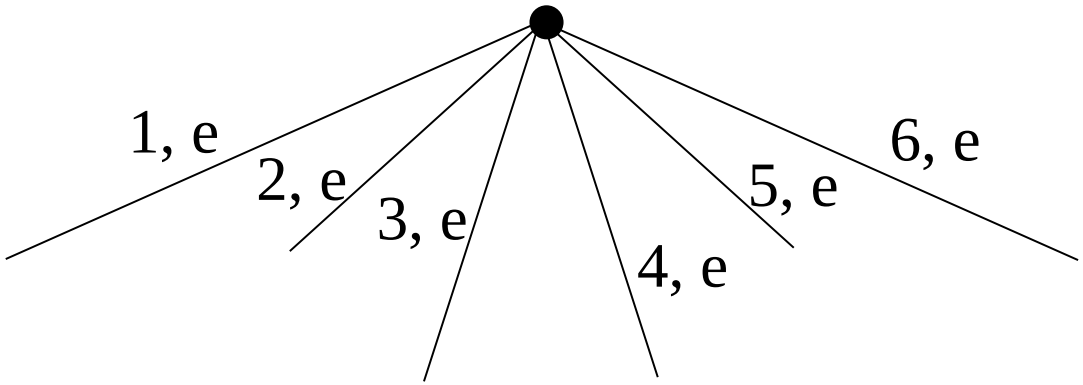

Global end index: xyzabcxyzr

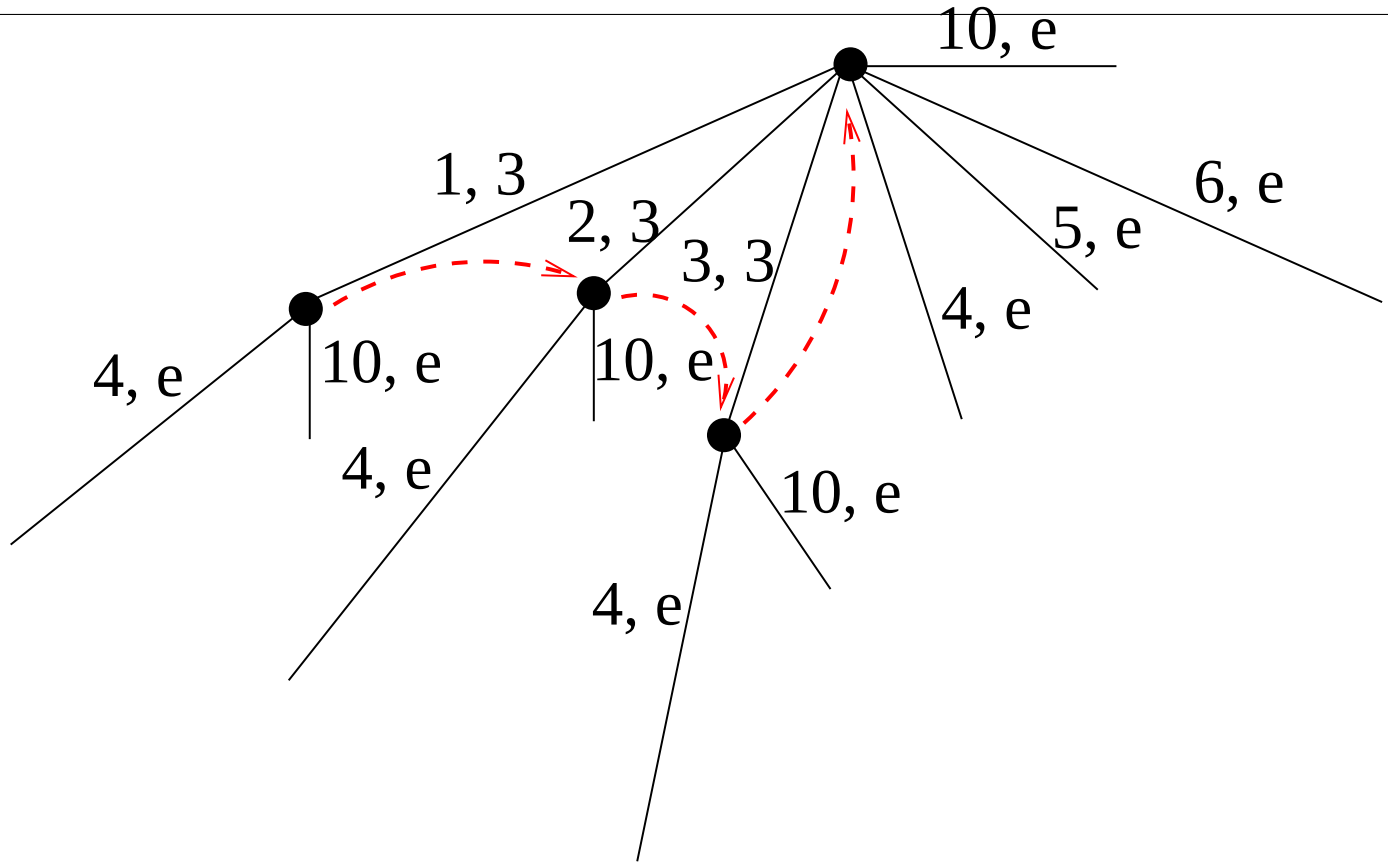

Final "tricks"

- Two situations as yet unobserved
- Early stopping
  - The first node where an extension is found is sufficient
  - All subsequent suffix links will have a match
- Skip count
  - When finding suffix link, just match first character
  - Rest of characters on arc MUST match

Early stopping and Skip/count: xyztyzrxyzr

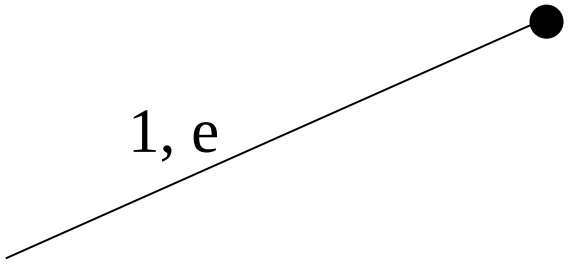

Early stopping and Skip/count: xyztyzrxyzr

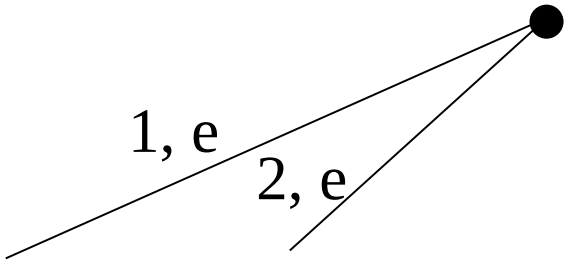

Early stopping and Skip/count: xyztyzrxyzr

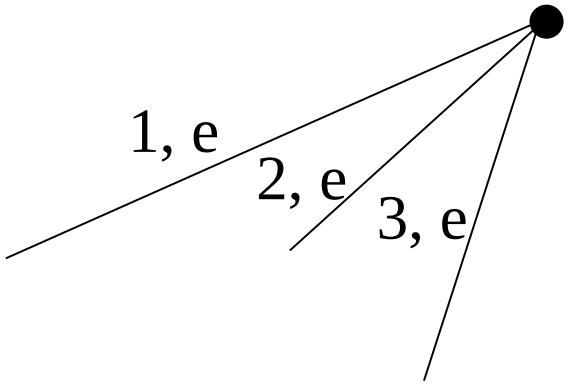

Early stopping and Skip/count: xyztyzrxyzr

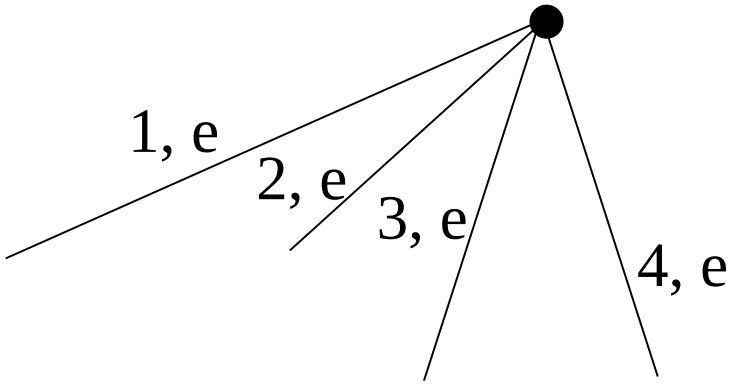

Early stopping and Skip/count: xyztyzrxyzr

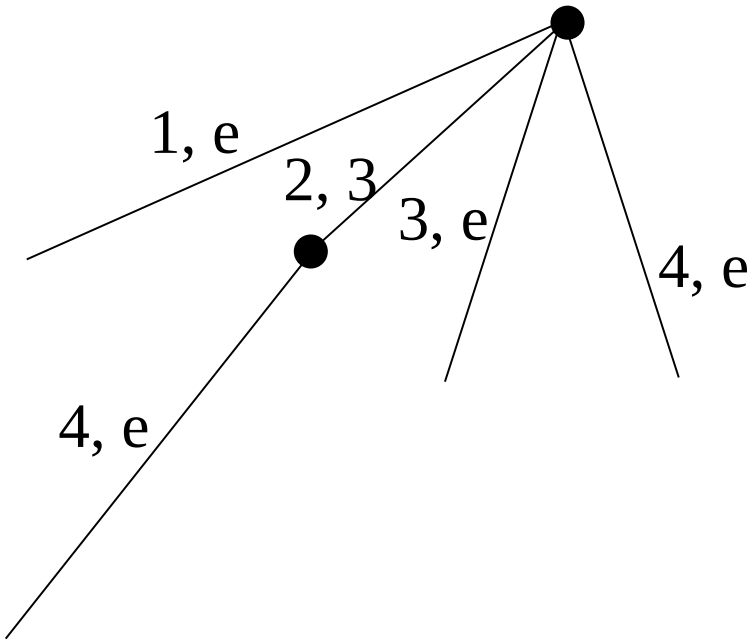

Early stopping and Skip/count: xyztyzrxyzr

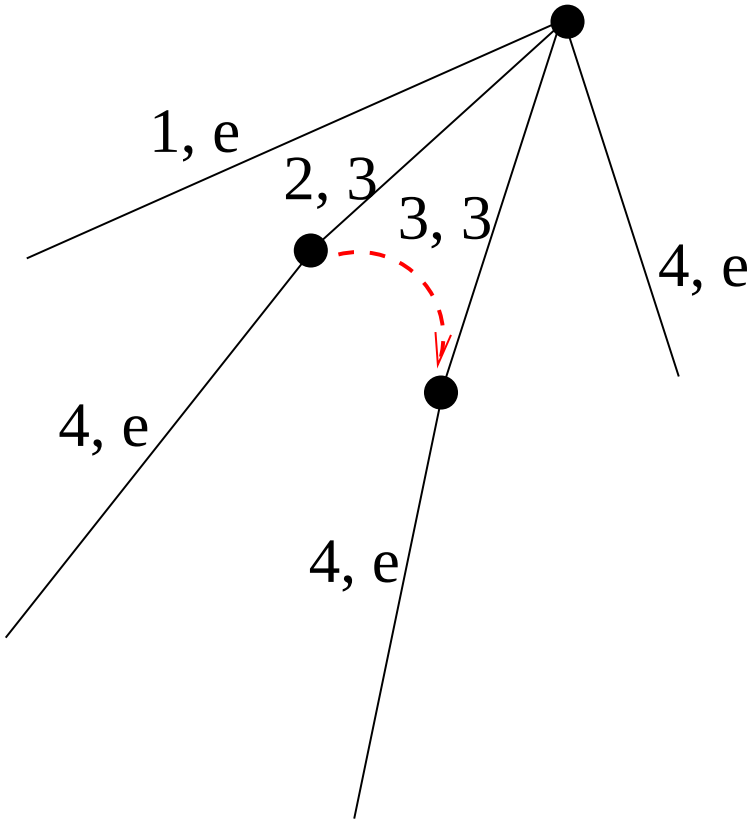

Early stopping and Skip/count: xyztyzrxyzr

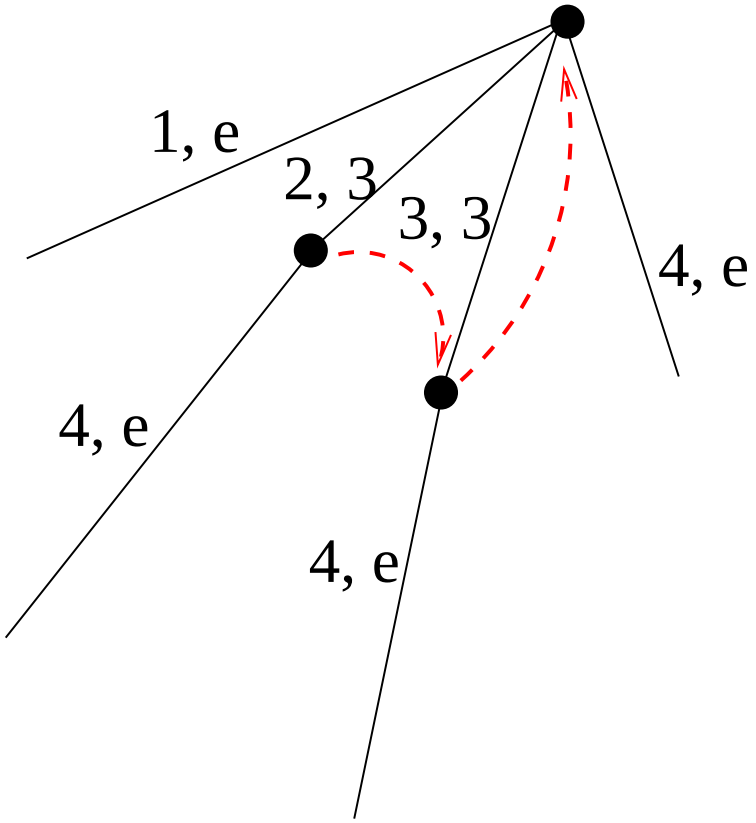

Early stopping and Skip/count: xyztyzrxyzr

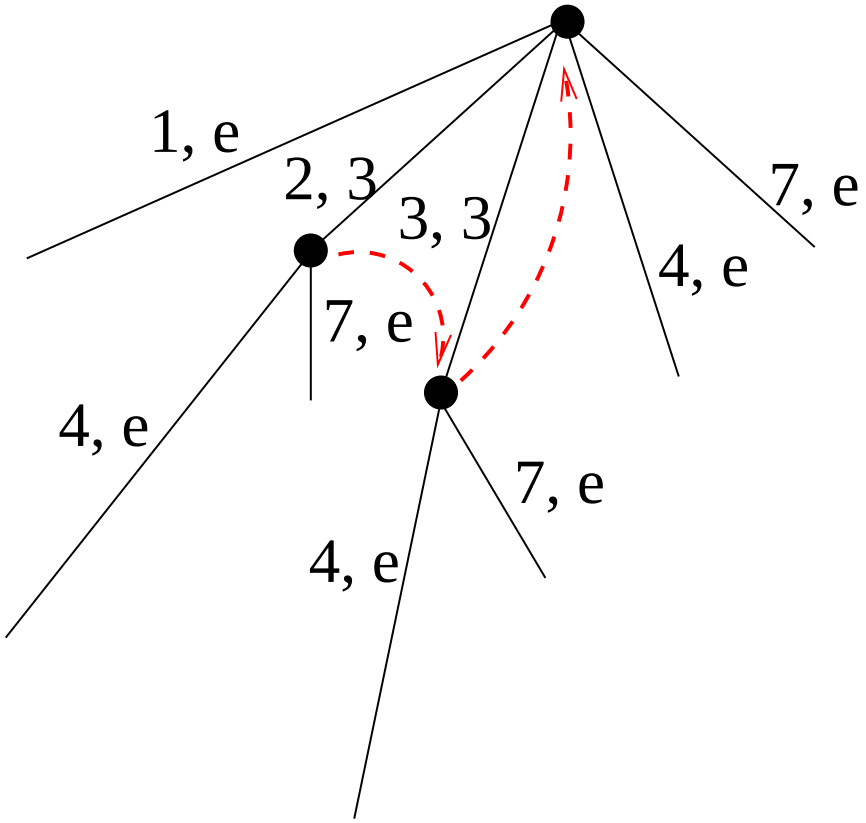

Early stopping and Skip/count: xyztyzrxyzr

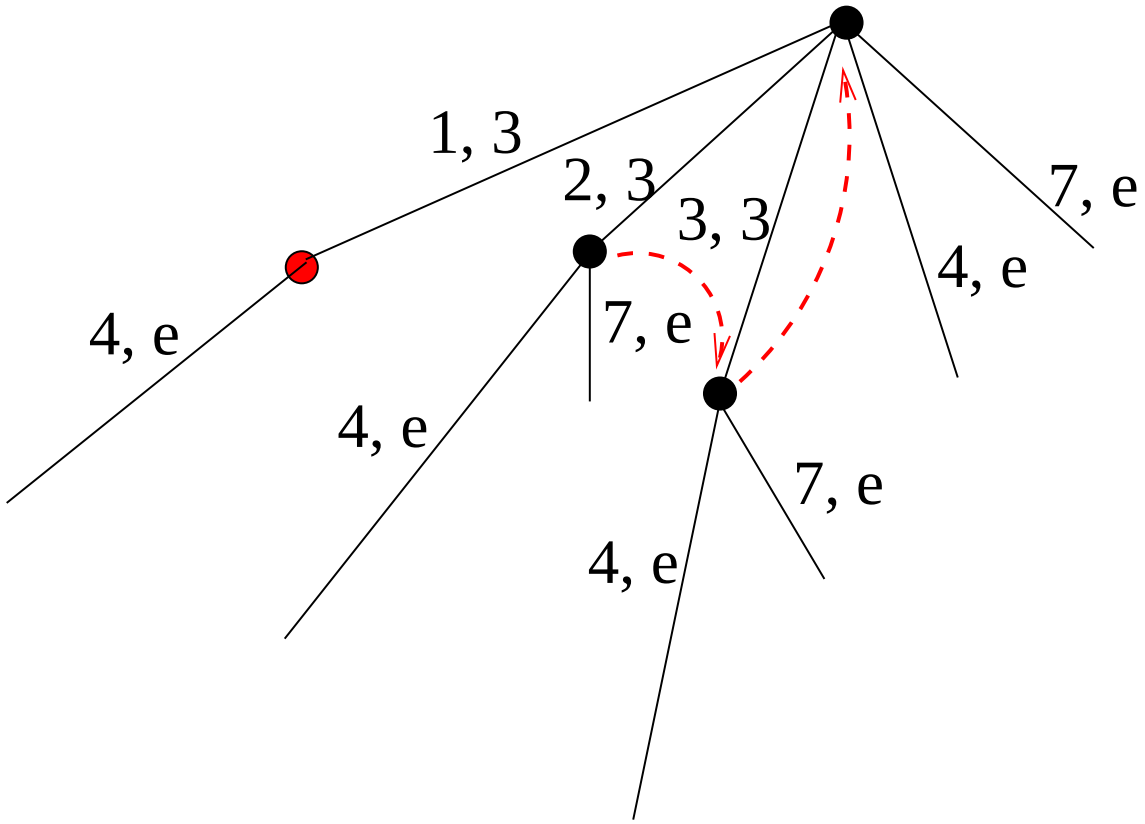

Early stopping and Skip/count: xyztyzrxyzr

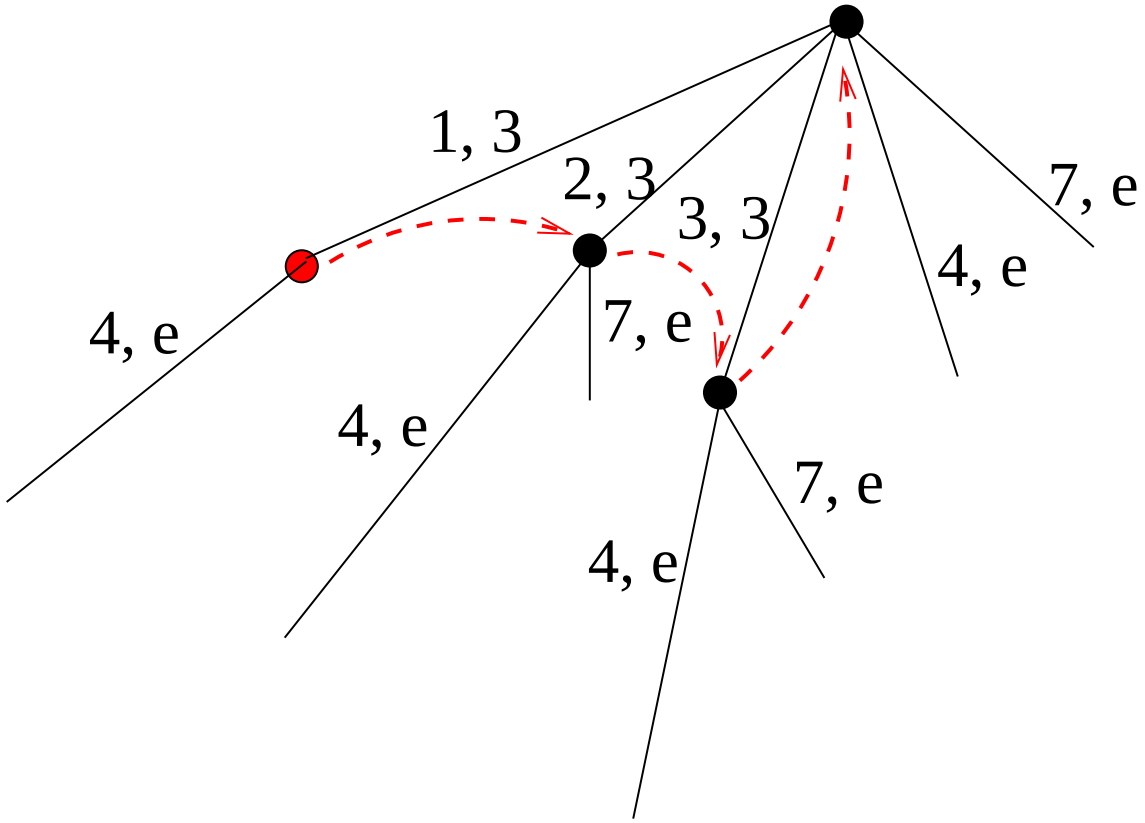

Early stopping and Skip/count: xyztyzrxyzr

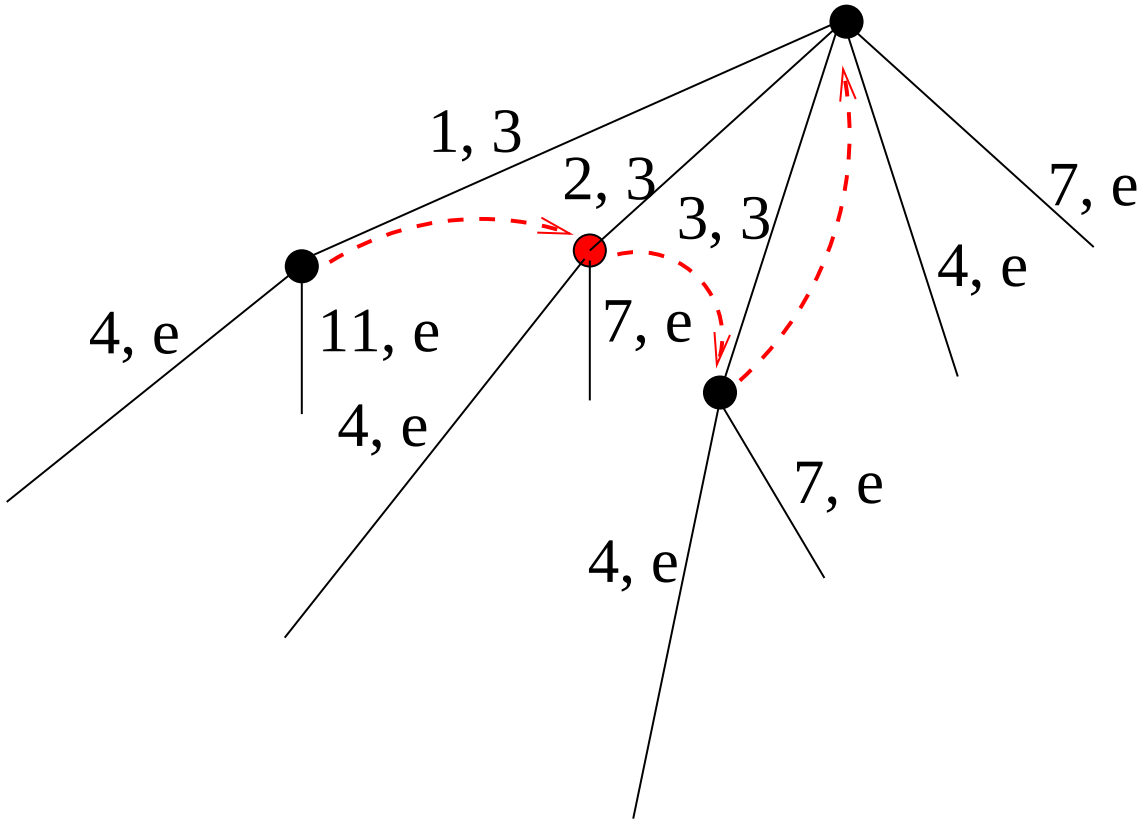

Example:

1	2	3	4	5	6	7	8	9	0	1	2	3	4	5	6
x	t	p	y	x	t	p	z	x	t	p	y	x	t	p	r

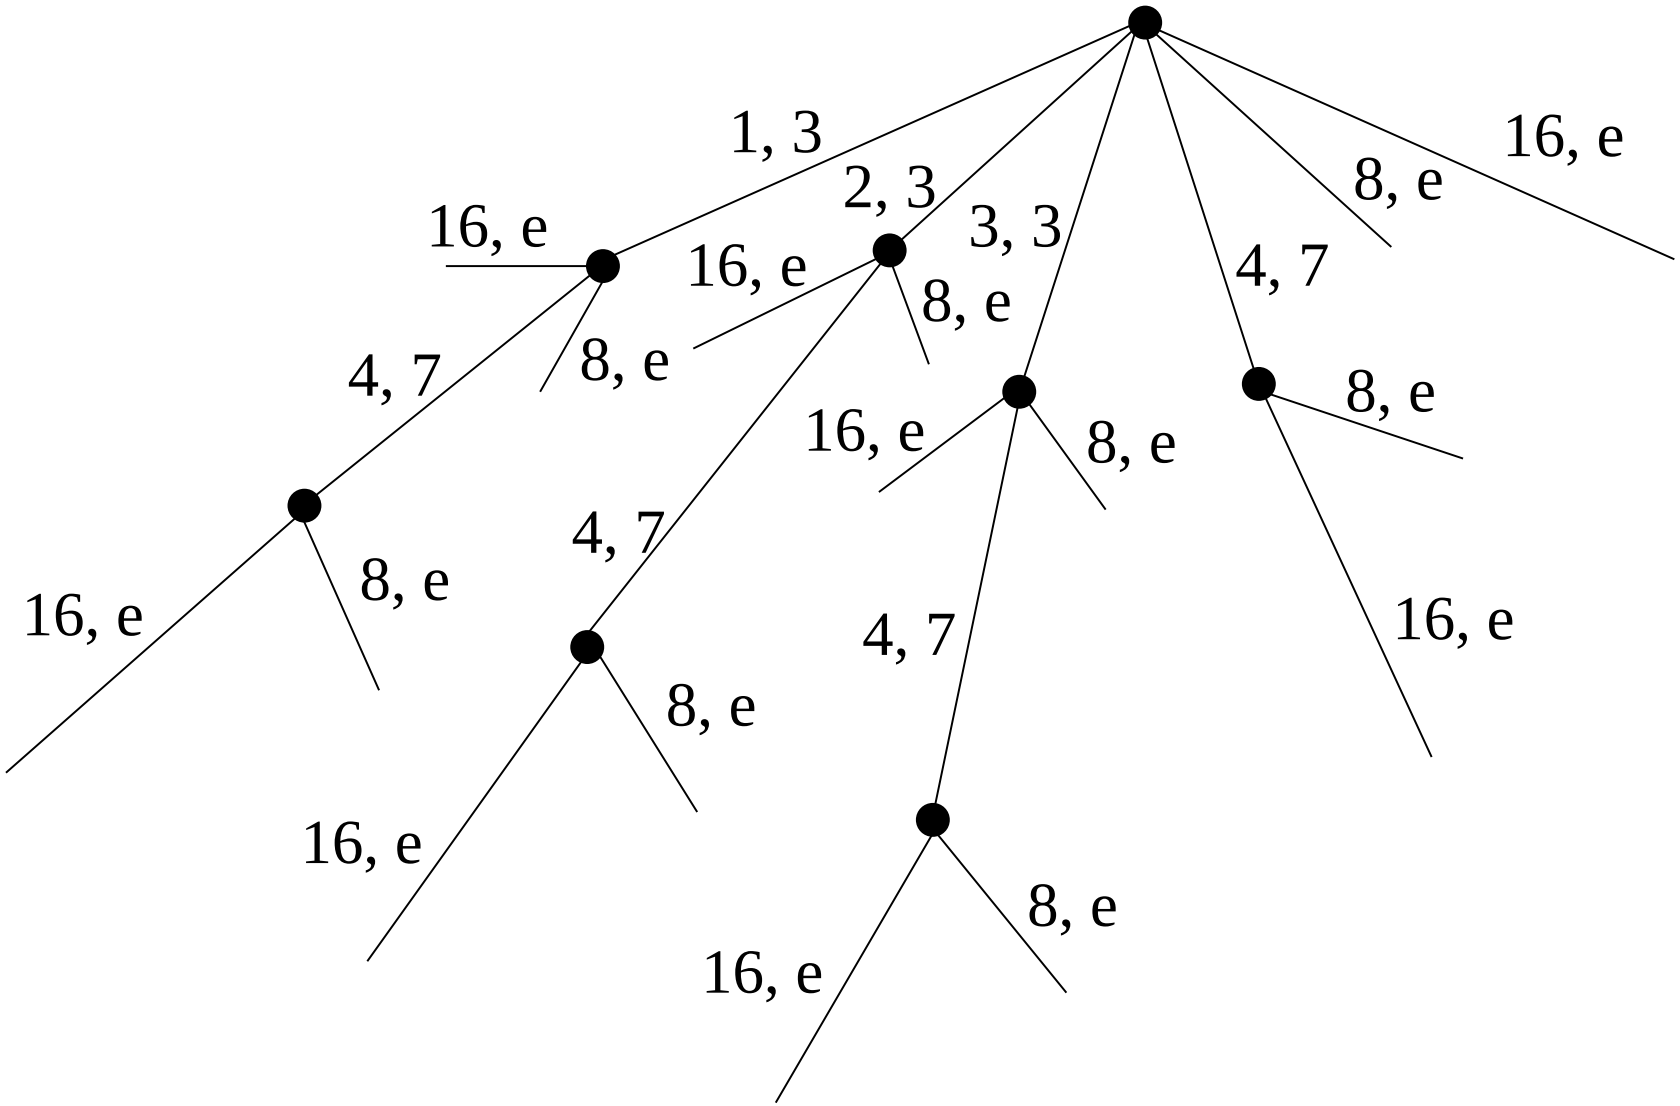

Example:	1	2	3	4	5	6	7	8	9	0	1	2	3	4	5	6
	x	t	p	y	x	t	p	z	x	t	p	y	x	t	p	r

Example:	1	2	3	4	5	6	7	8	9	0	1	2	3	4	5	6
	x	t	p	y	x	t	p	z	x	t	p	y	x	t	p	r

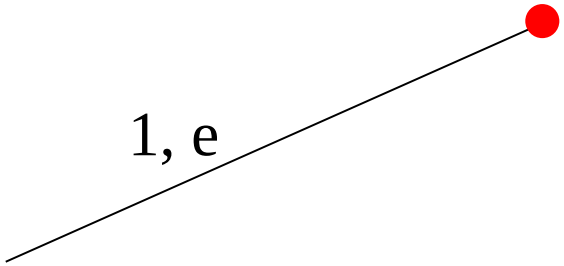

Example:	1	2	3	4	5	6	7	8	9	0	1	2	3	4	5	6
	x	t	p	y	x	t	p	z	x	t	p	y	x	t	p	r

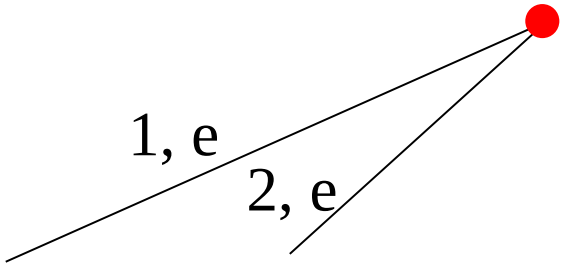

Example:	1	2	3	4	5	6	7	8	9	0	1	2	3	4	5	6
	x	t	p	y	x	t	p	z	x	t	p	y	x	t	p	r

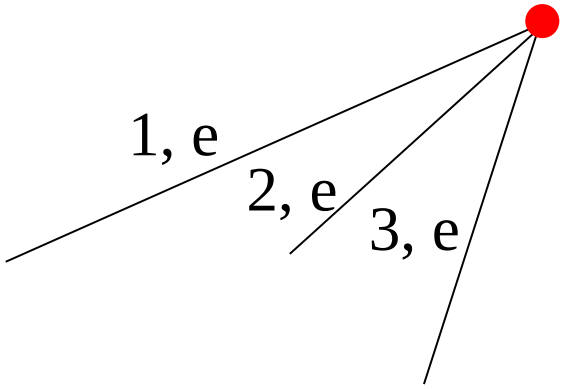

Example:	1	2	3	4	5	6	7	8	9	0	1	2	3	4	5	6
	x	t	p	y	x	t	p	z	x	t	p	y	x	t	p	r

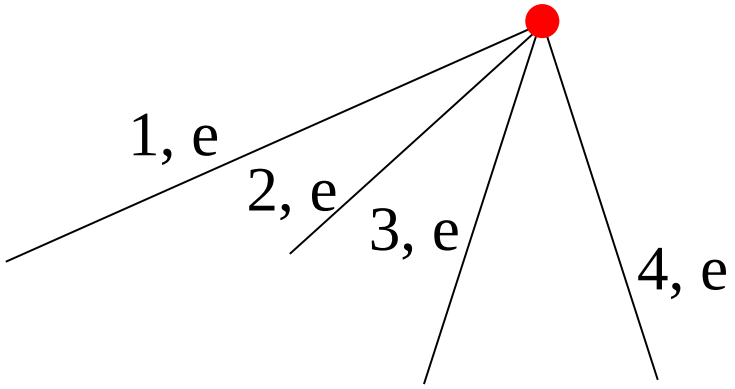

Example:

1	2	3	4	5	6	7	8	9	0	1	2	3	4	5	6
x	t	p	y	x	t	p	z	x	t	p	y	x	t	p	r

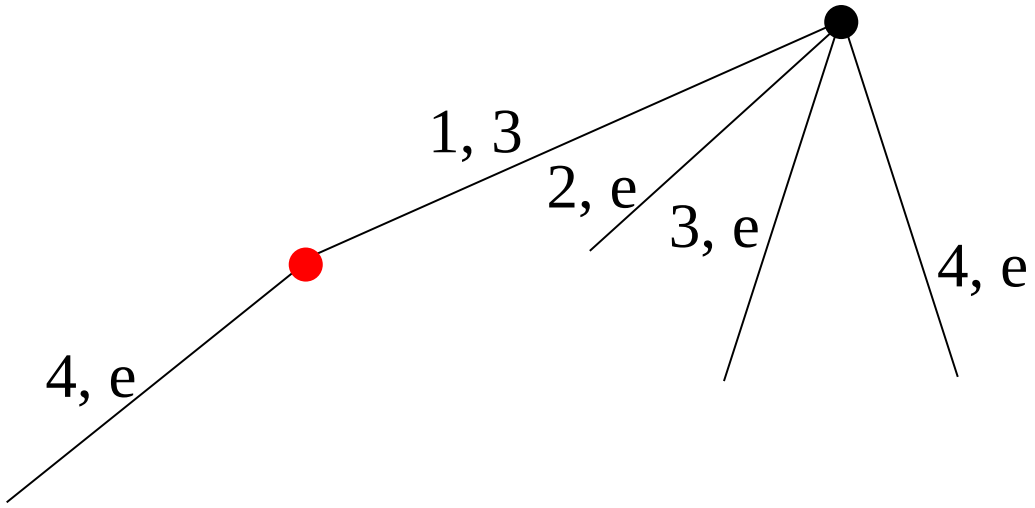

Example:

1	2	3	4	5	6	7	8	9	0	1	2	3	4	5	6
x	t	p	y	x	t	p	z	x	t	p	y	x	t	p	r

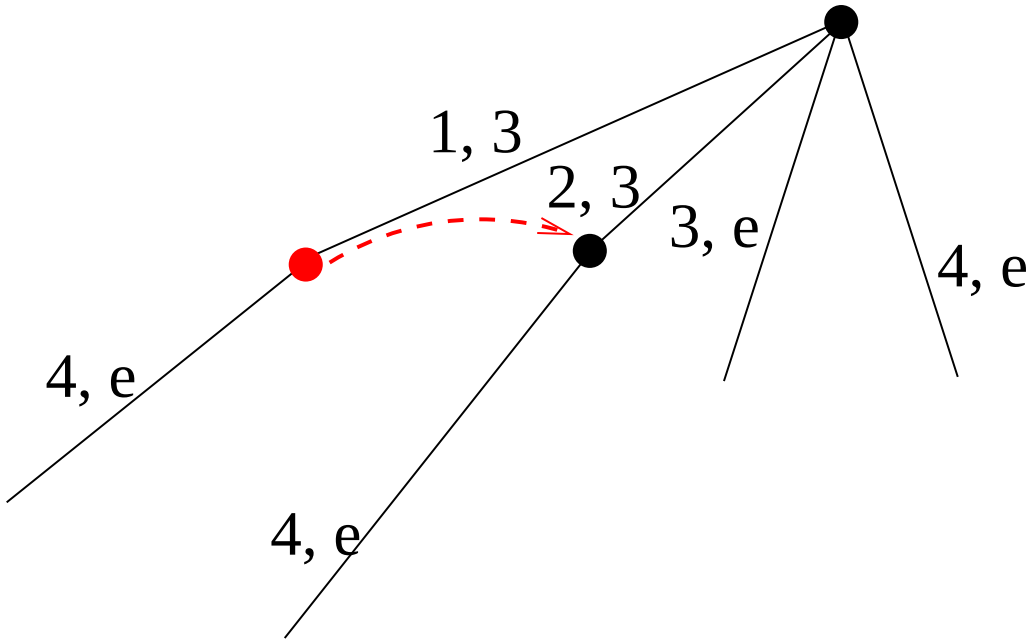

Example:

1	2	3	4	5	6	7	8	9	0	1	2	3	4	5	6
x	t	p	y	x	t	p	z	x	t	p	y	x	t	p	r

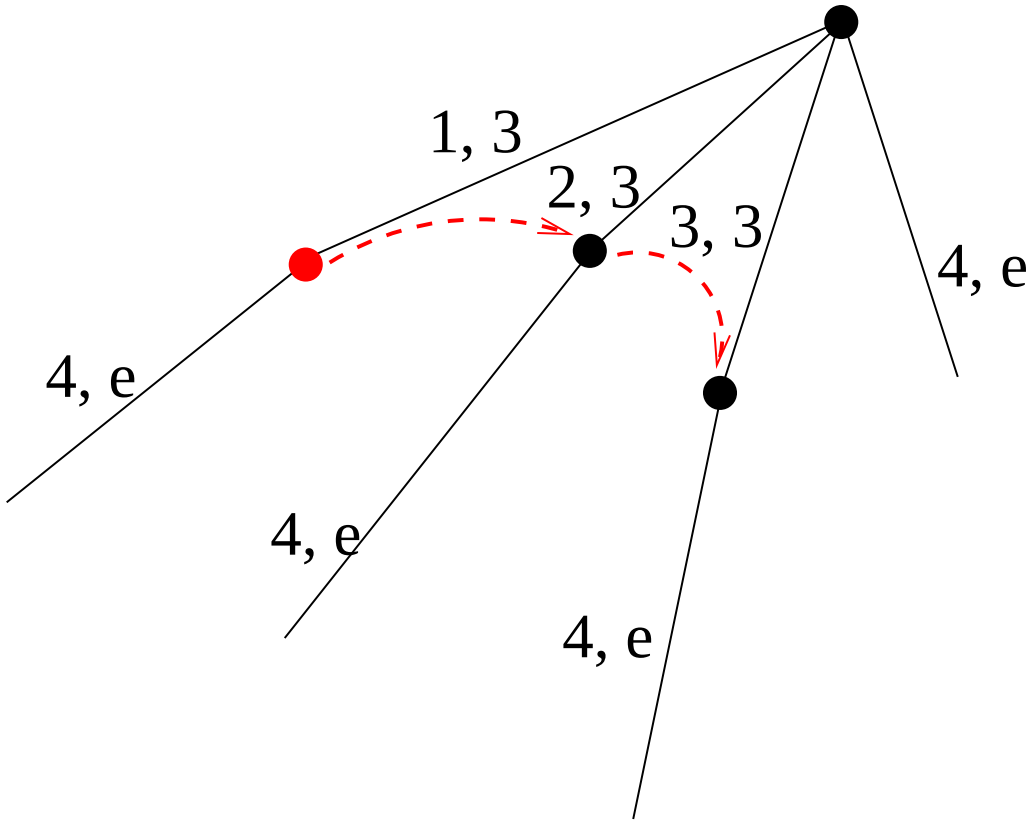

Example:

1	2	3	4	5	6	7	8	9	0	1	2	3	4	5	6
x	t	p	y	x	t	p	z	x	t	p	y	x	t	p	r

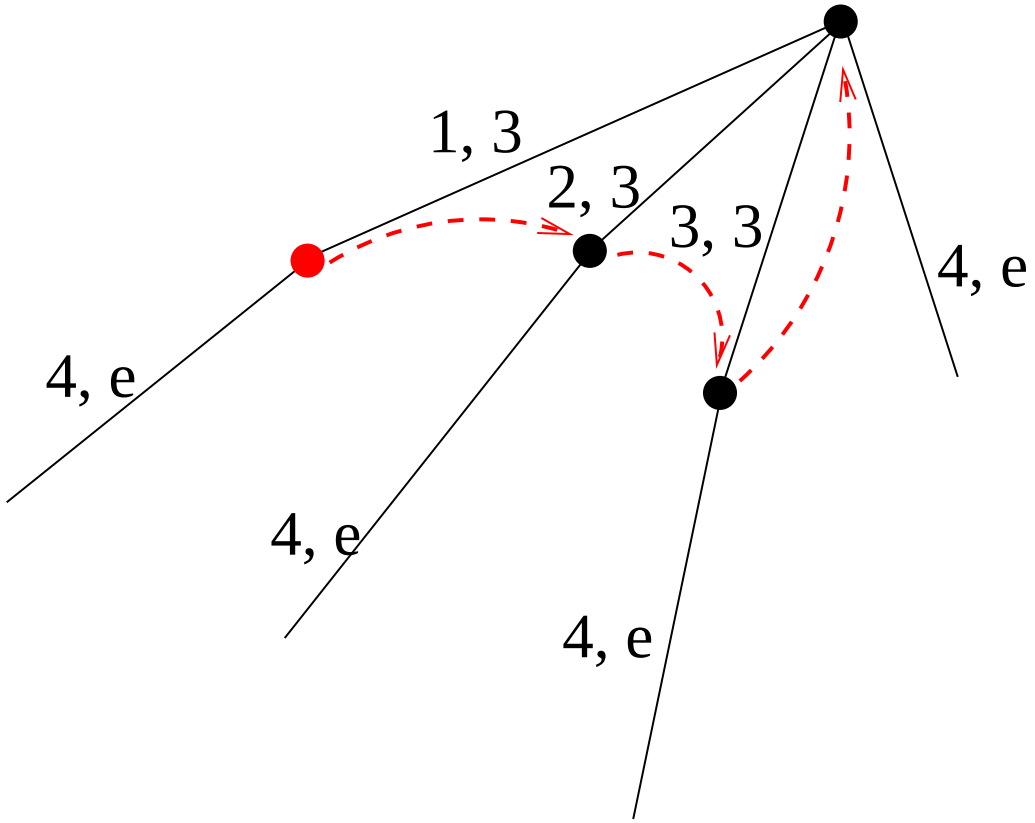

Example:

1	2	3	4	5	6	7	8	9	0	1	2	3	4	5	6
x	t	p	y	x	t	p	z	x	t	p	y	x	t	p	r

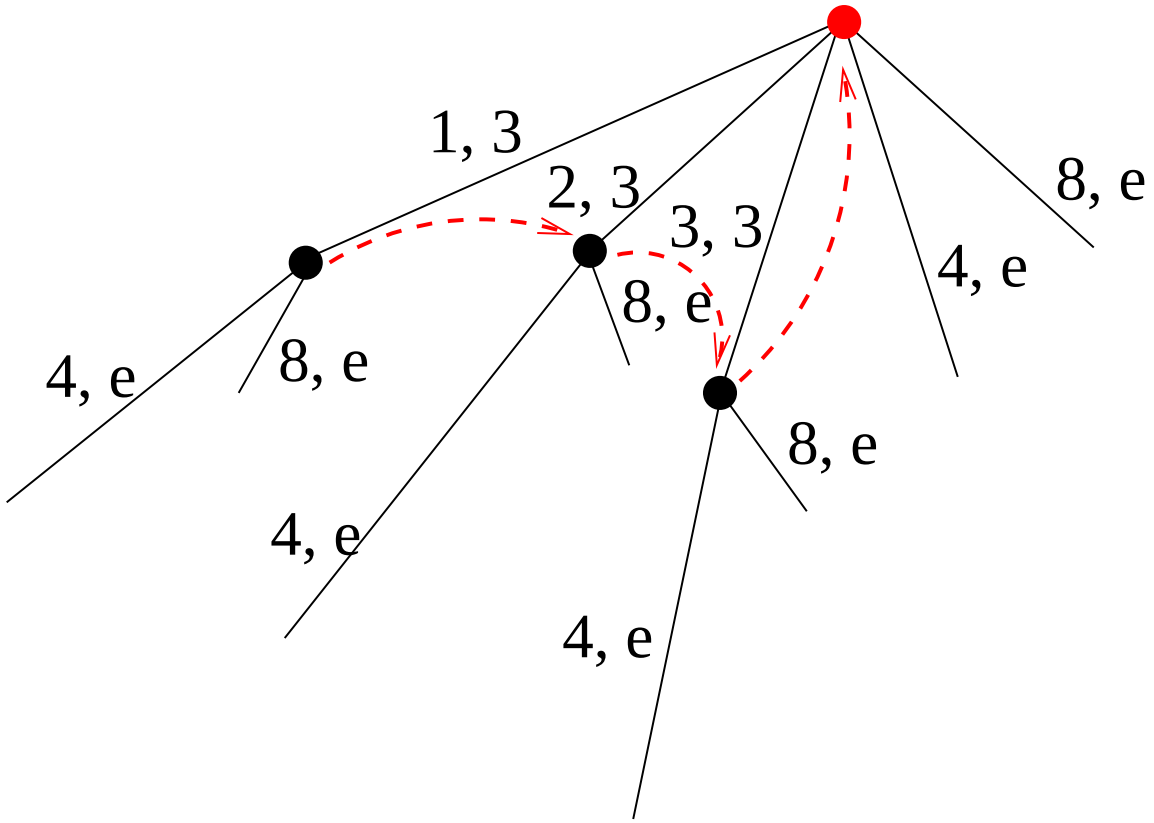

Example:

1	2	3	4	5	6	7	8	9	0	1	2	3	4	5	6
x	t	p	y	x	t	p	z	x	t	p	y	x	t	p	r

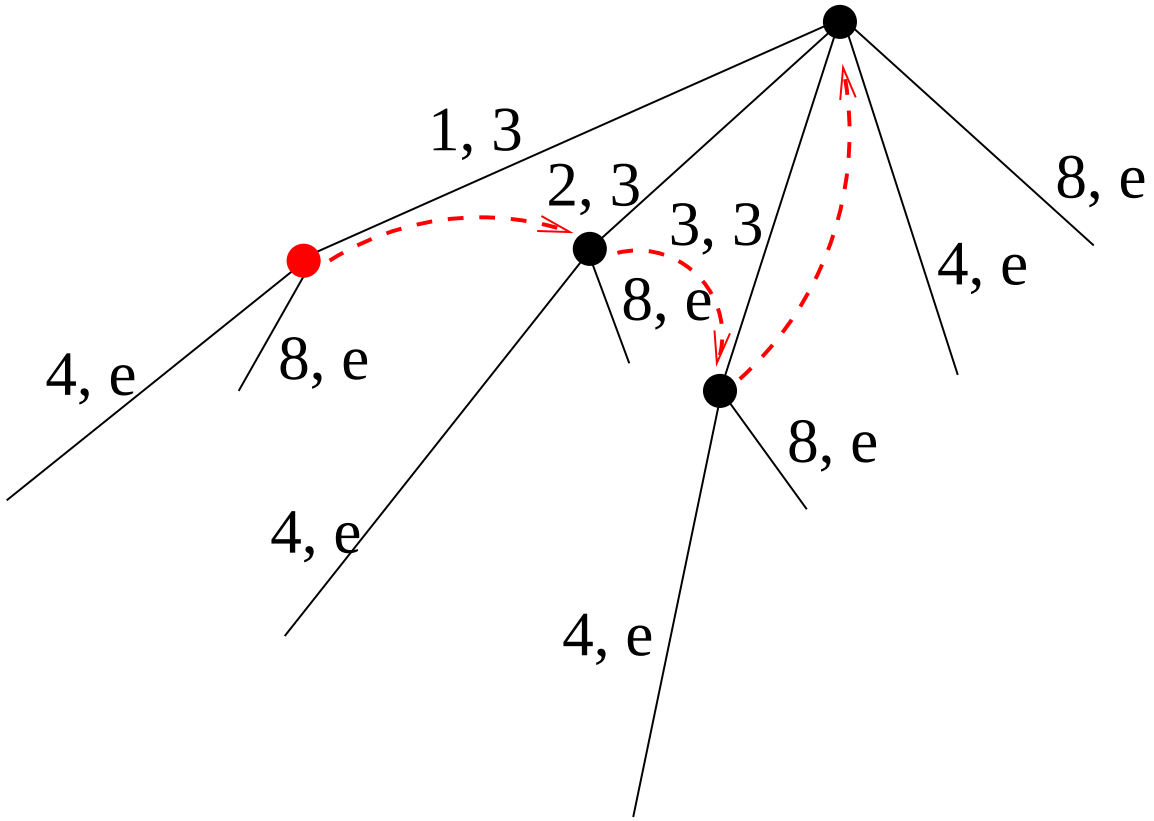

Example:

1	2	3	4	5	6	7	8	9	0	1	2	3	4	5	6
x	t	p	y	x	t	p	z	x	t	p	y	x	t	p	r

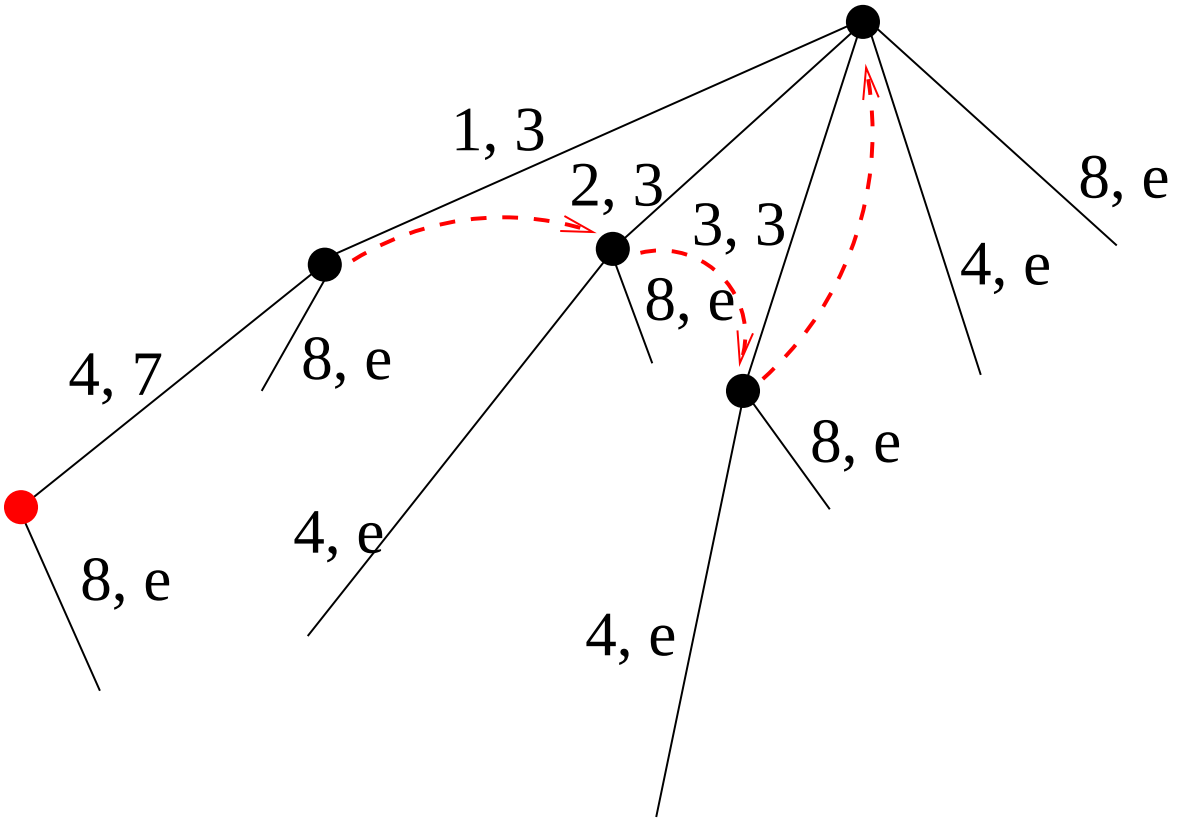

Example:

1	2	3	4	5	6	7	8	9	0	1	2	3	4	5	6
x	t	p	y	x	t	p	z	x	t	p	y	x	t	p	r

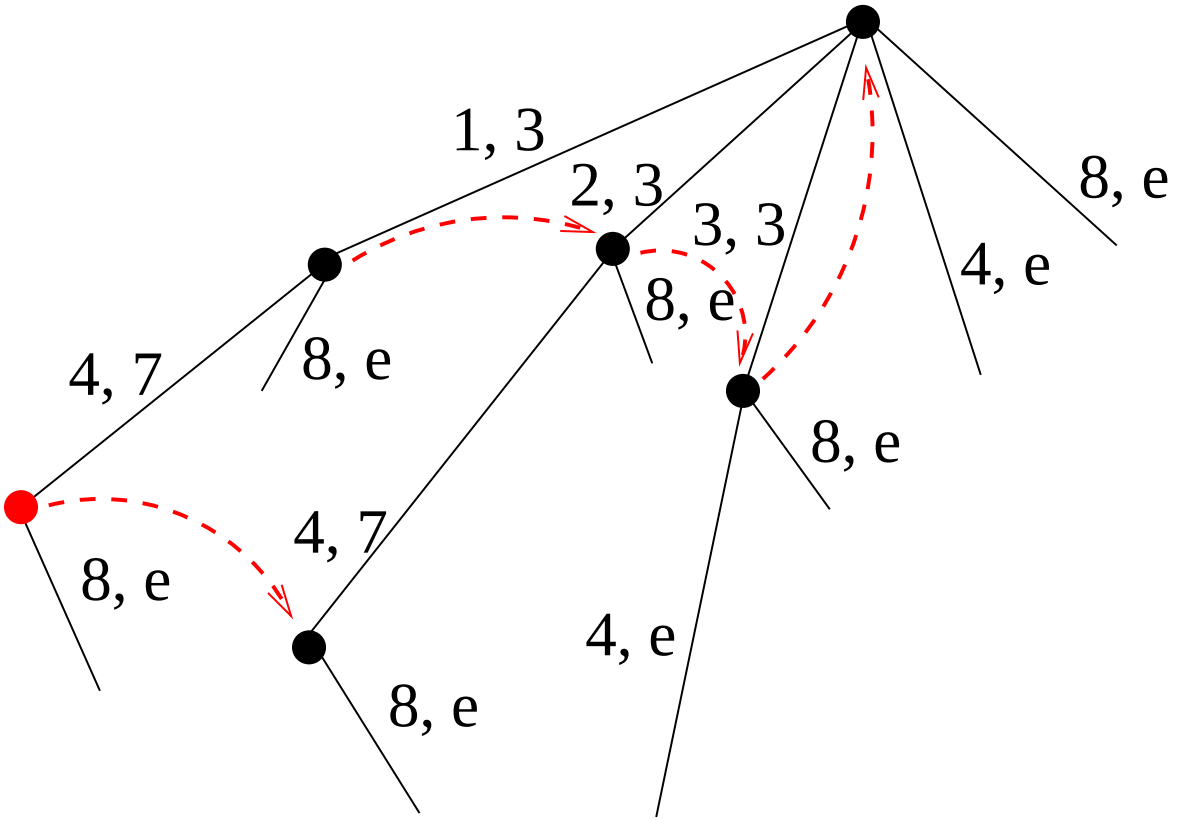

Example:

1	2	3	4	5	6	7	8	9	0	1	2	3	4	5	6
x	t	p	y	x	t	p	z	x	t	p	y	x	t	p	r

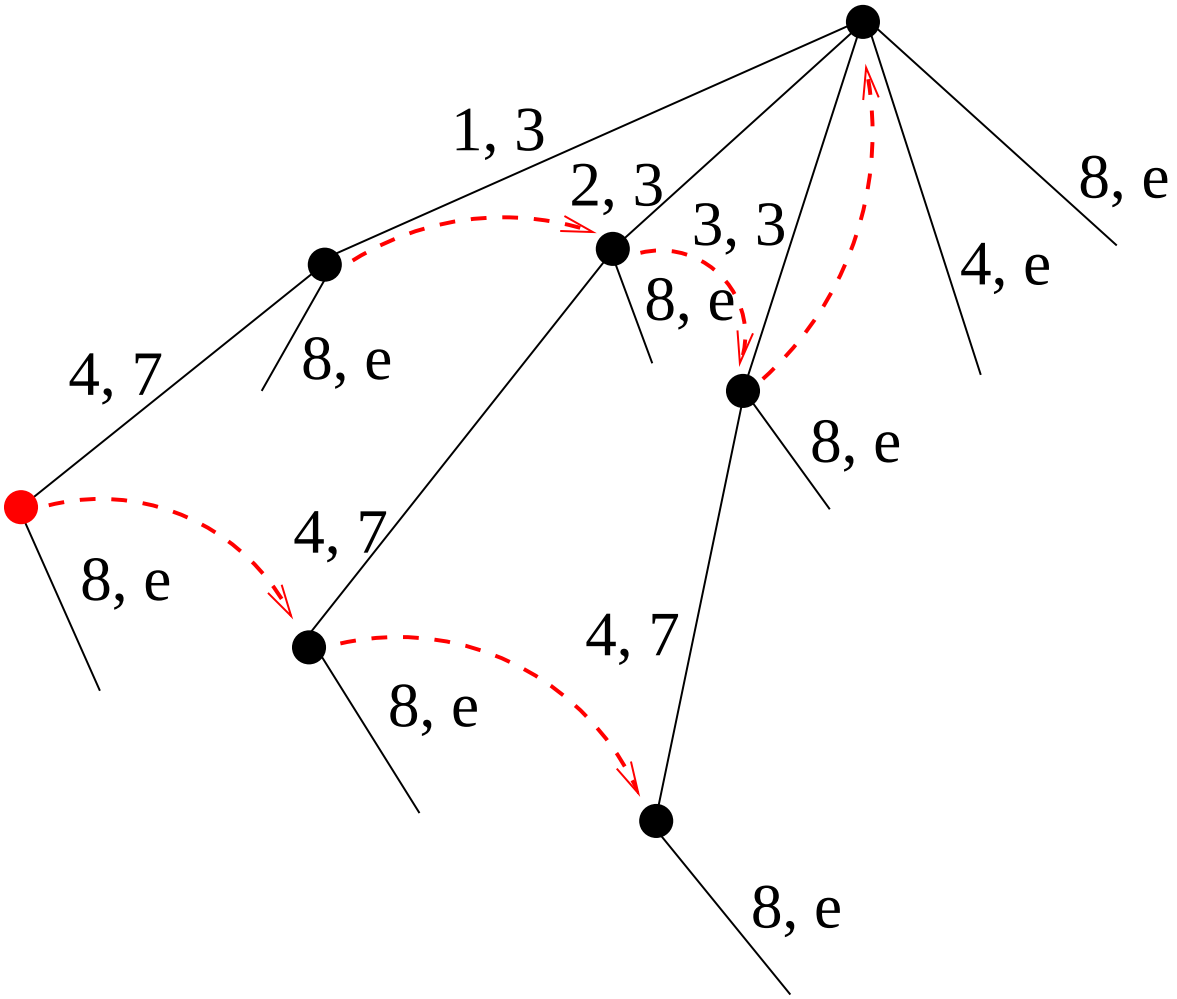

Example:

1	2	3	4	5	6	7	8	9	0	1	2	3	4	5	6
x	t	p	y	x	t	p	z	x	t	p	y	x	t	p	r

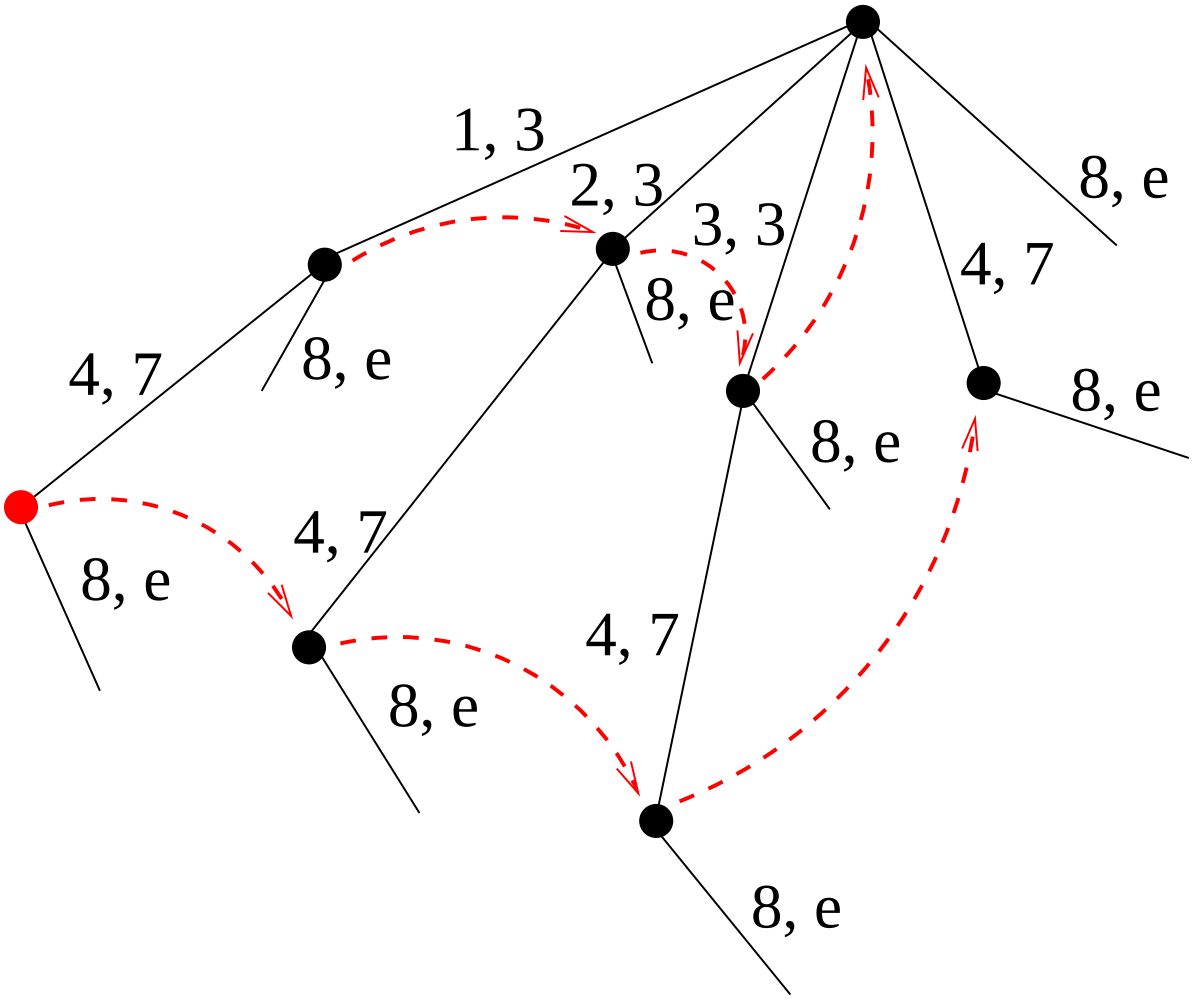

Example:

1	2	3	4	5	6	7	8	9	0	1	2	3	4	5	6
x	t	p	y	x	t	p	z	x	t	p	y	x	t	p	r

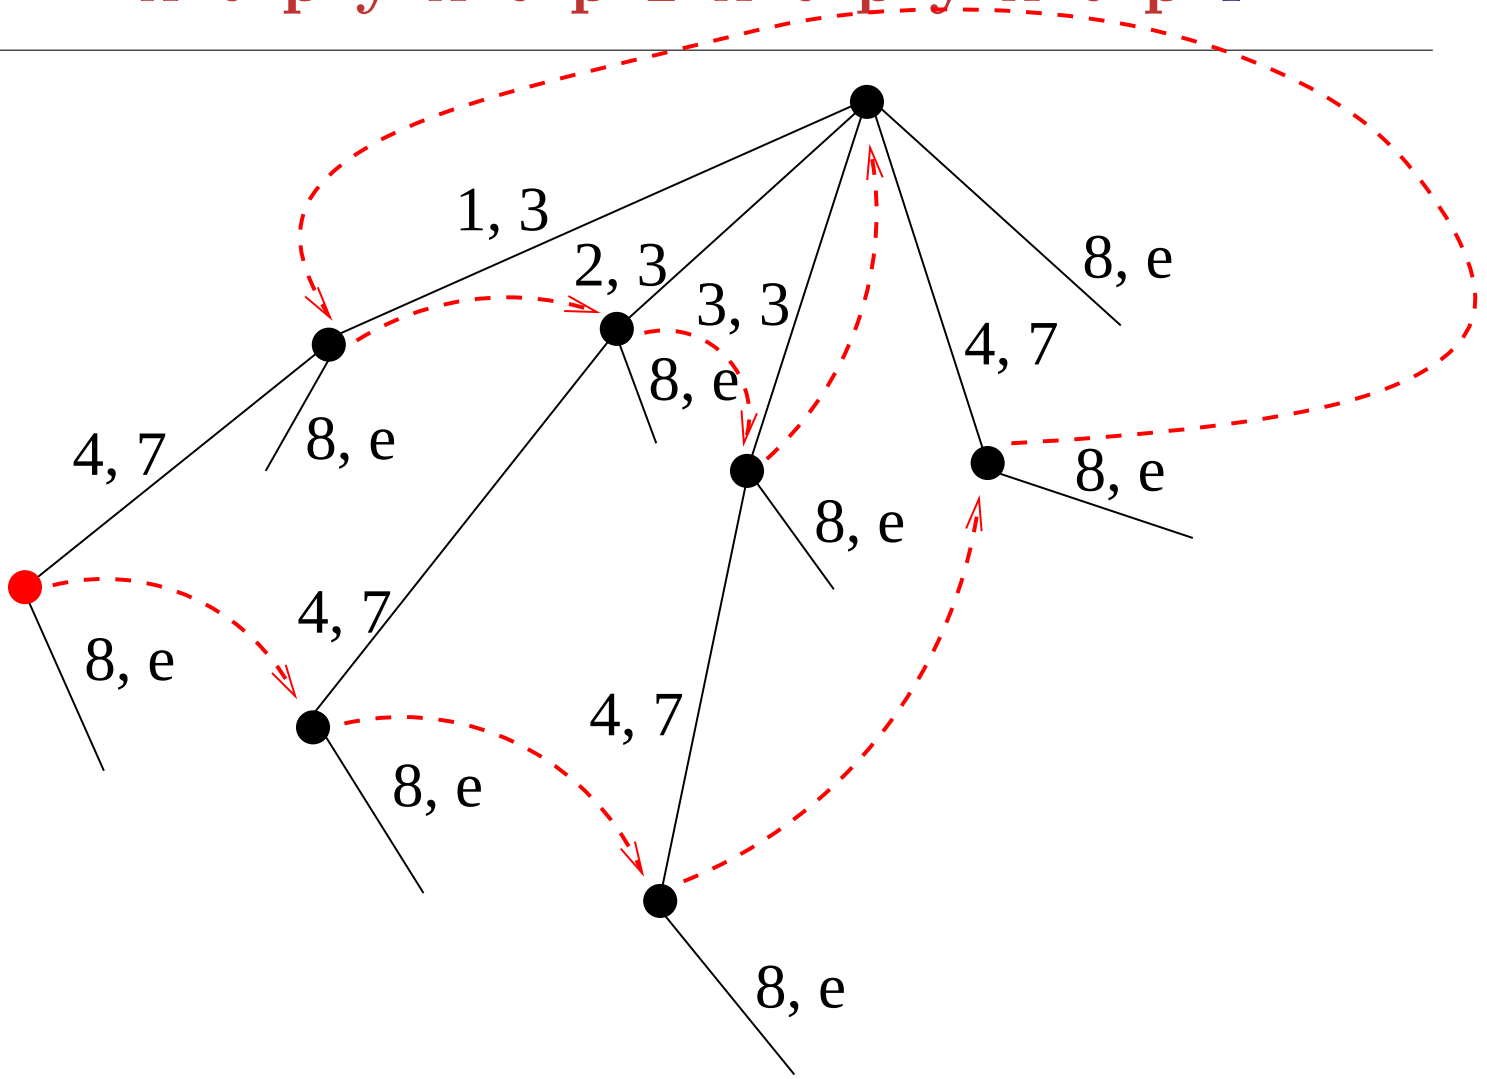

Example:

1	2	3	4	5	6	7	8	9	0	1	2	3	4	5	6
x	t	p	y	x	t	p	z	x	t	p	y	x	t	p	r

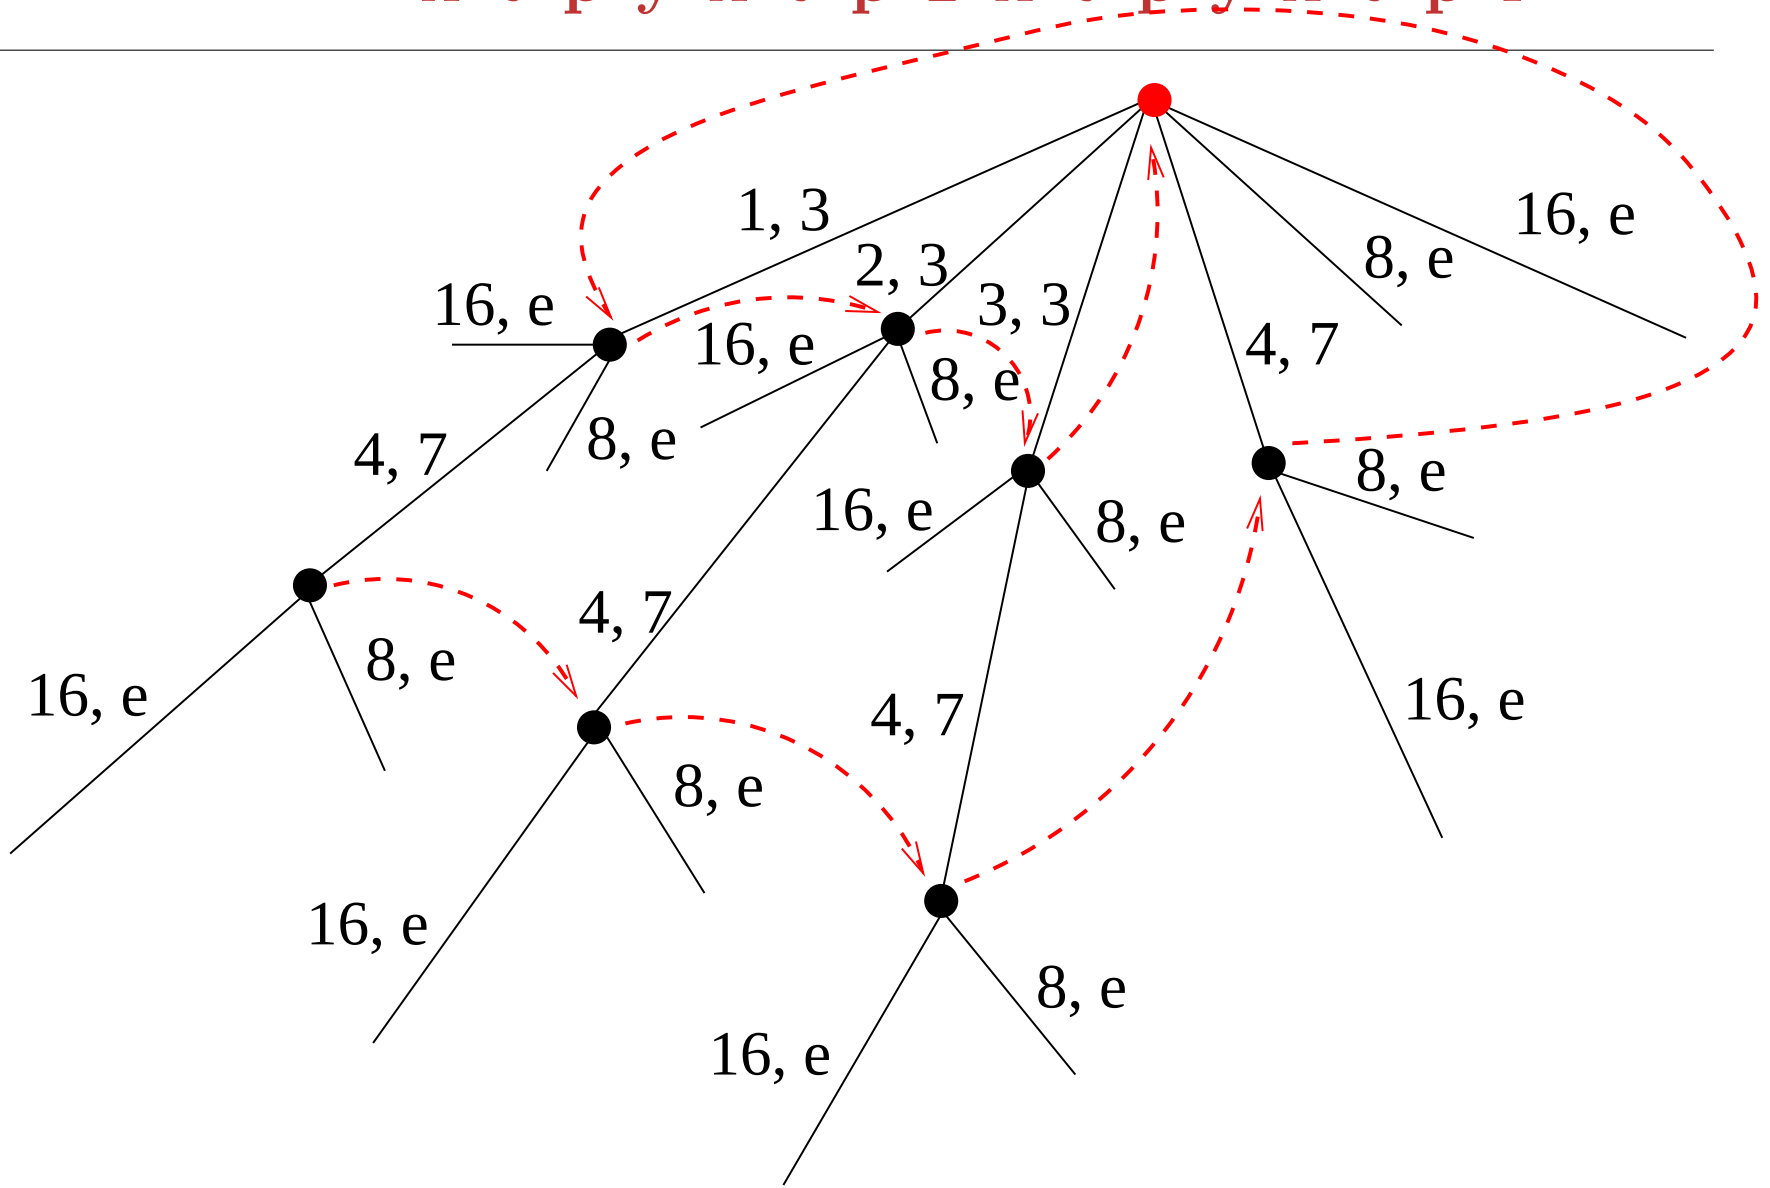

Searching for occurrences

- **Start at the root**
- **Search for match in tree**
- **Count the number of leaf nodes below point where word matches**
- **Linear in the length of the pattern, not the text**
- **Example: how many times does xt occur in text? t? py?**

Characteristics of suffix trees

- **One transition out of each state for any given symbol in Σ**
- **Size requirements quite large: $O(m|\Sigma|)$ for length m**
  - **m states, Σ continuations per state**
  - **Can hash, resulting in less space usage, but $O(n \log |\Sigma|)$ time to search, rather than $O(n)$**
- **With linear-time algorithm to build, construction not the problem**
- **Storage and use is the problem**

Suffix automata

- One space saving idea: turn tree into directed acyclic graph (DAG)
- Convert tree to finite-state automaton with re-entrancies
  - Tree is a special case of FSA: start state at root, final states at leaves
  - Further, tree is deterministic FSA
  - Minimize the tree to a minimal automaton
- Potentially massive space savings on final graph

Example suffix tree

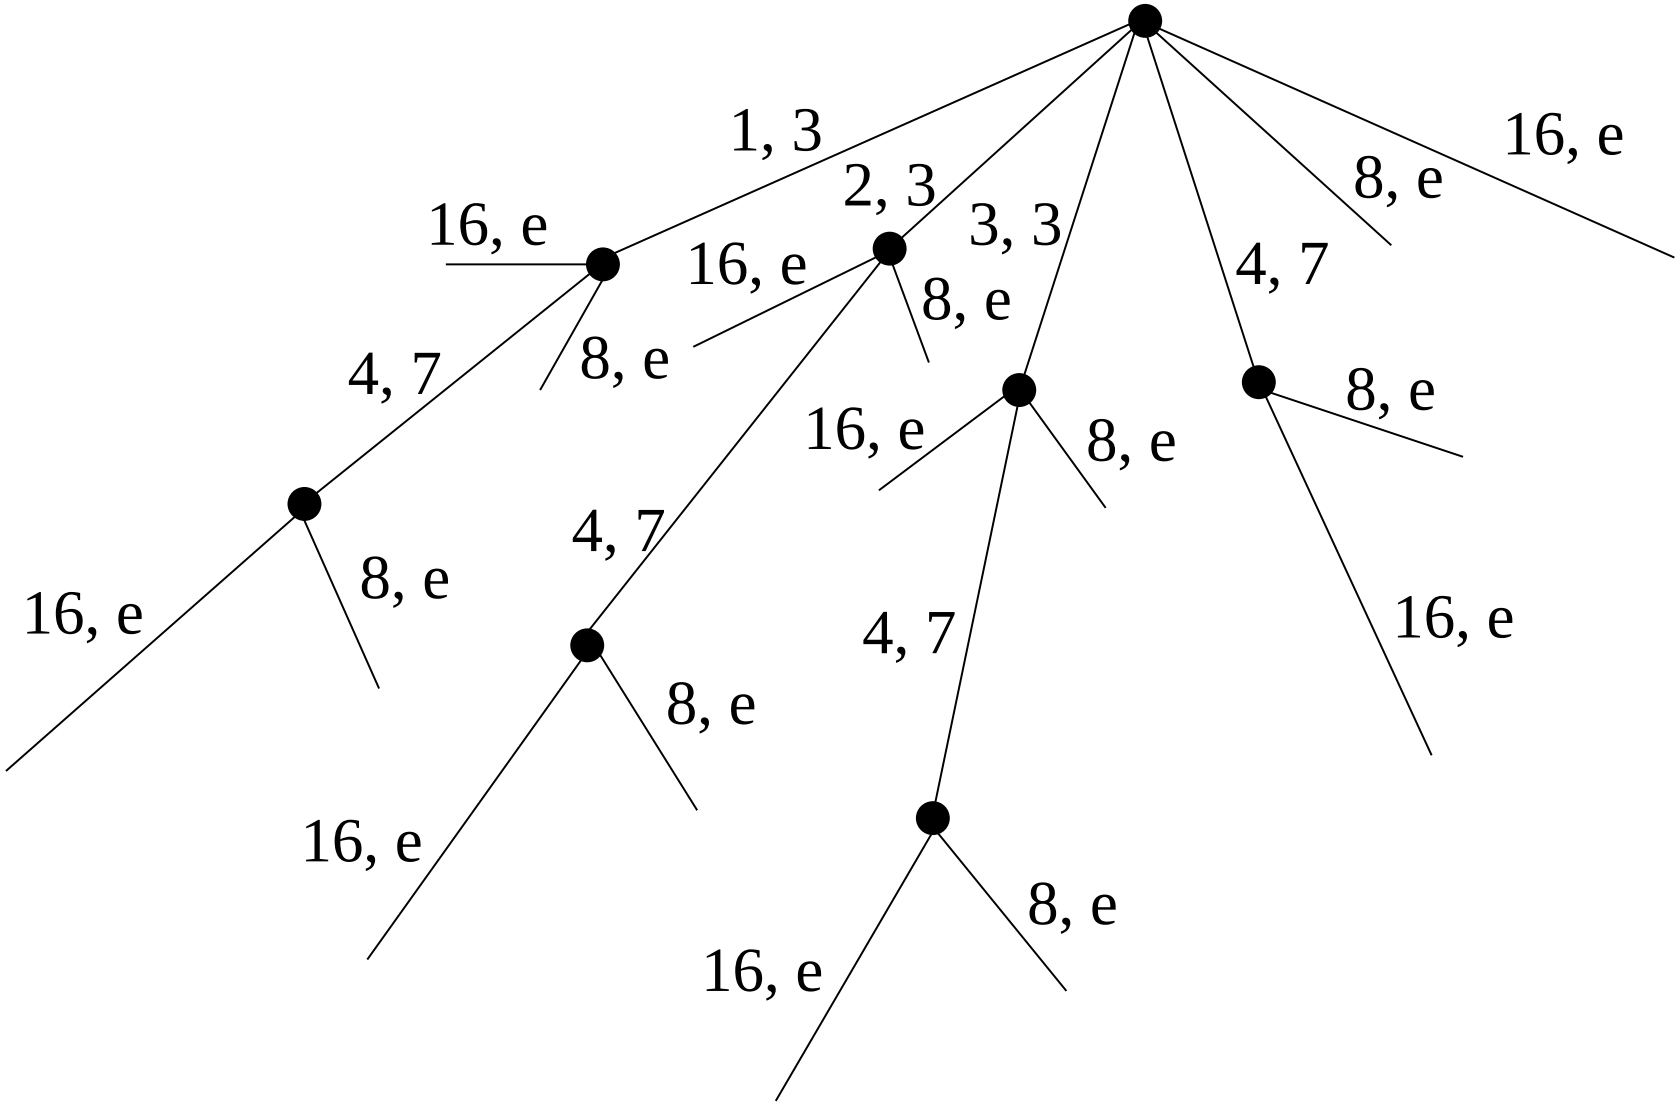

Change to suffix automata: directed edges, final states

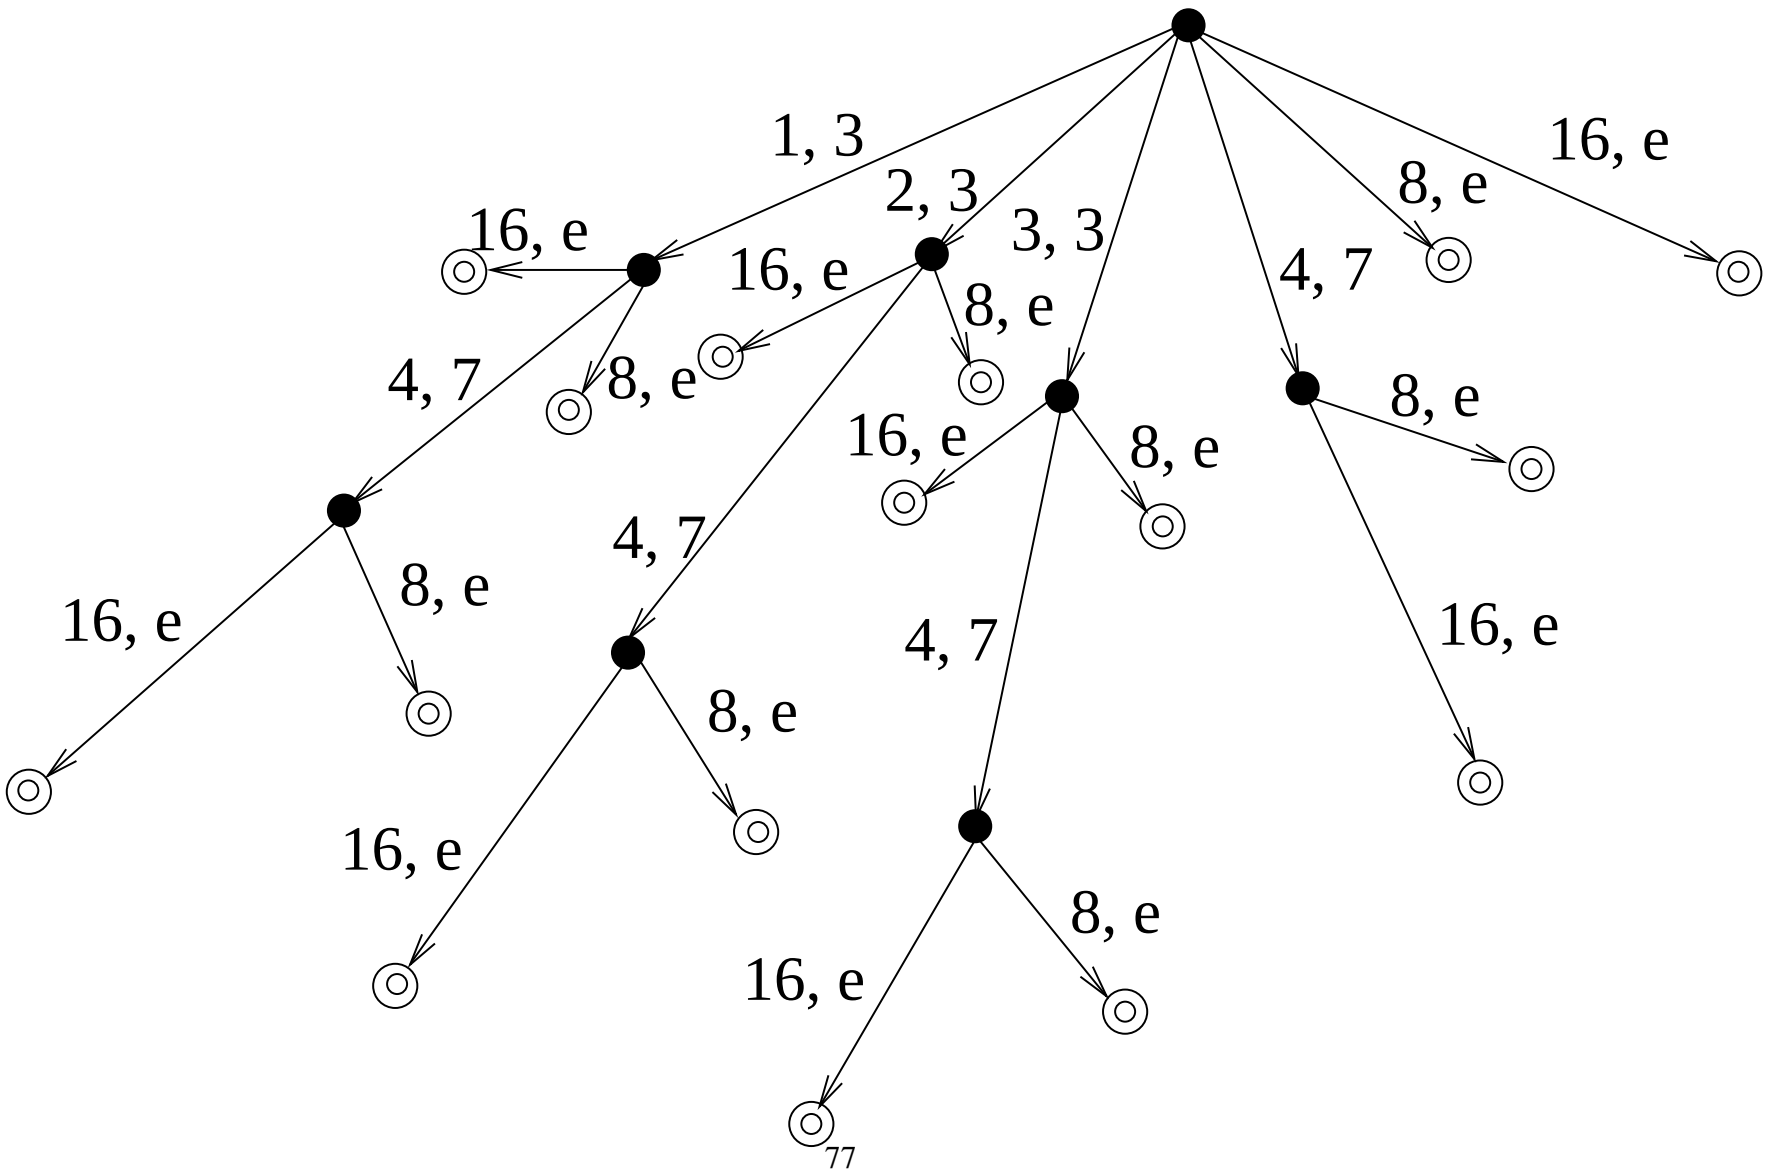

Change to suffix automata: collapse states

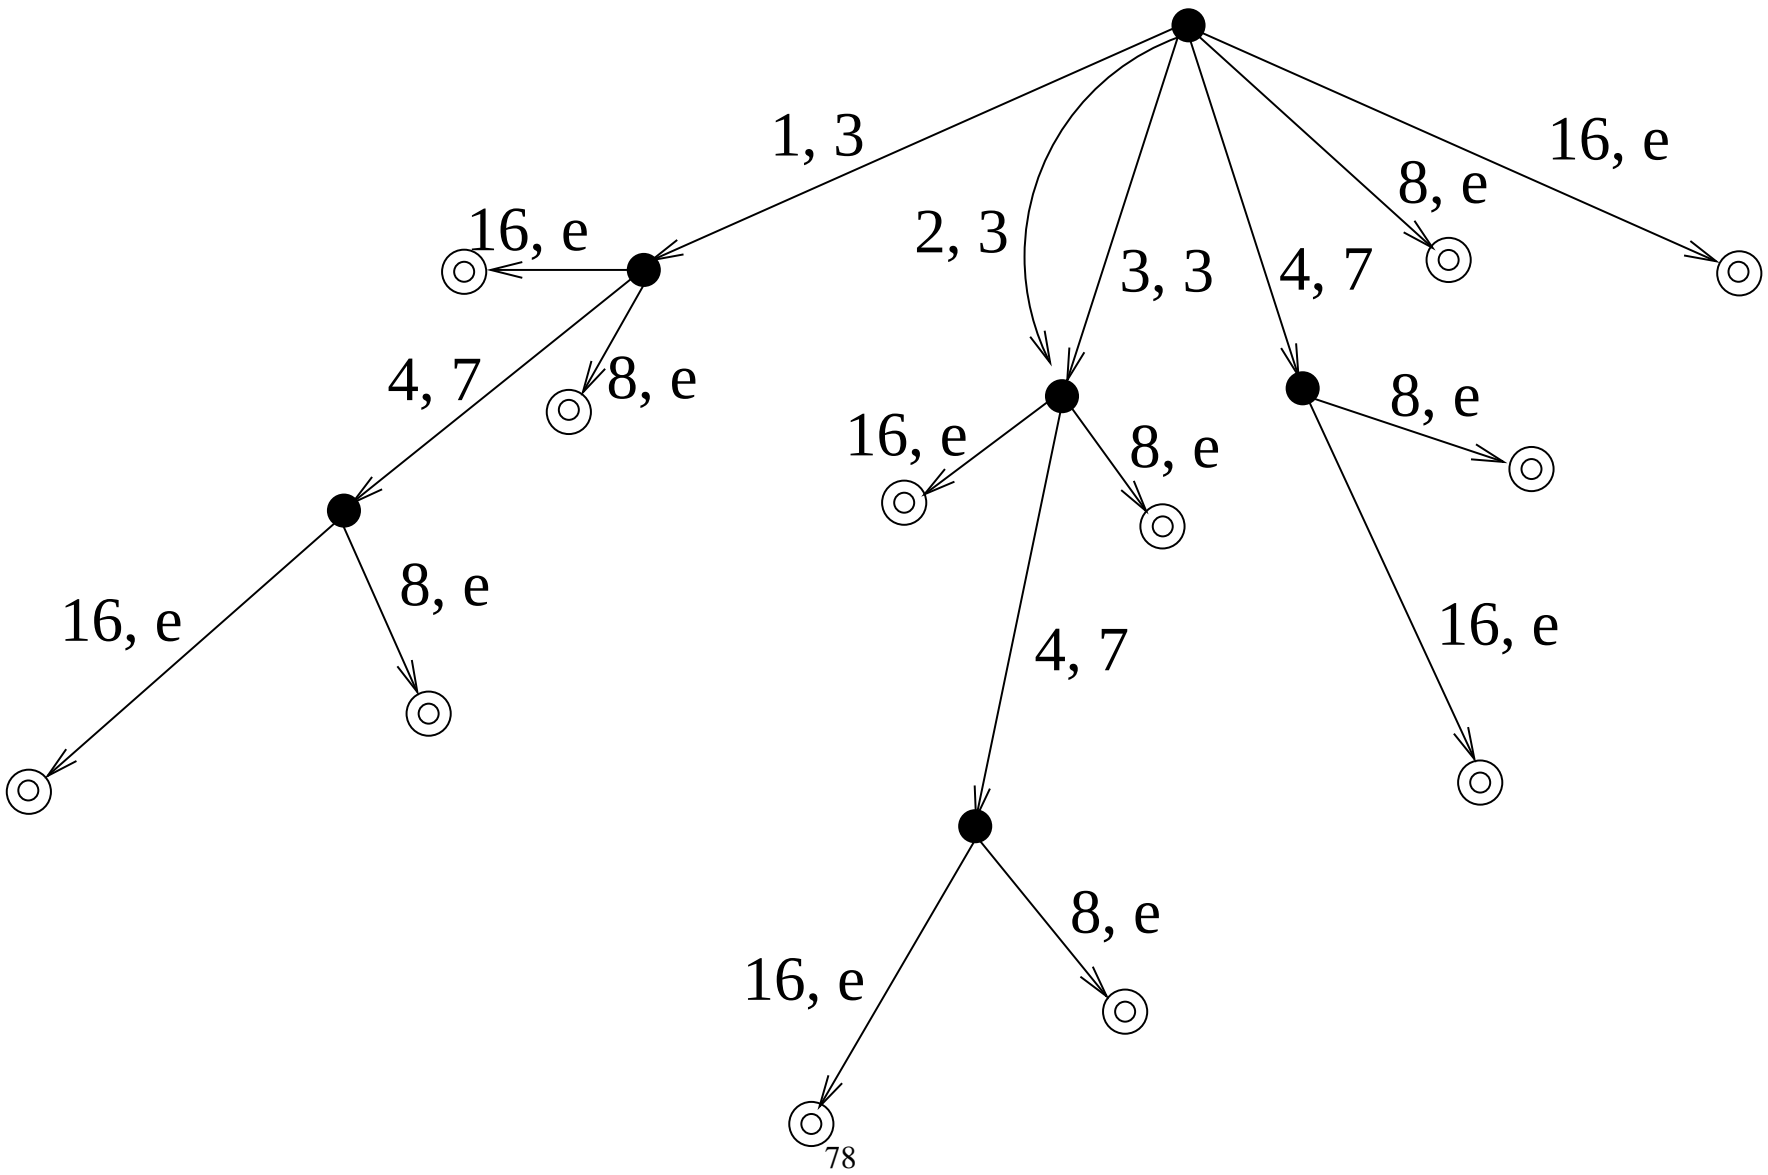

Change to suffix automata: collapse states

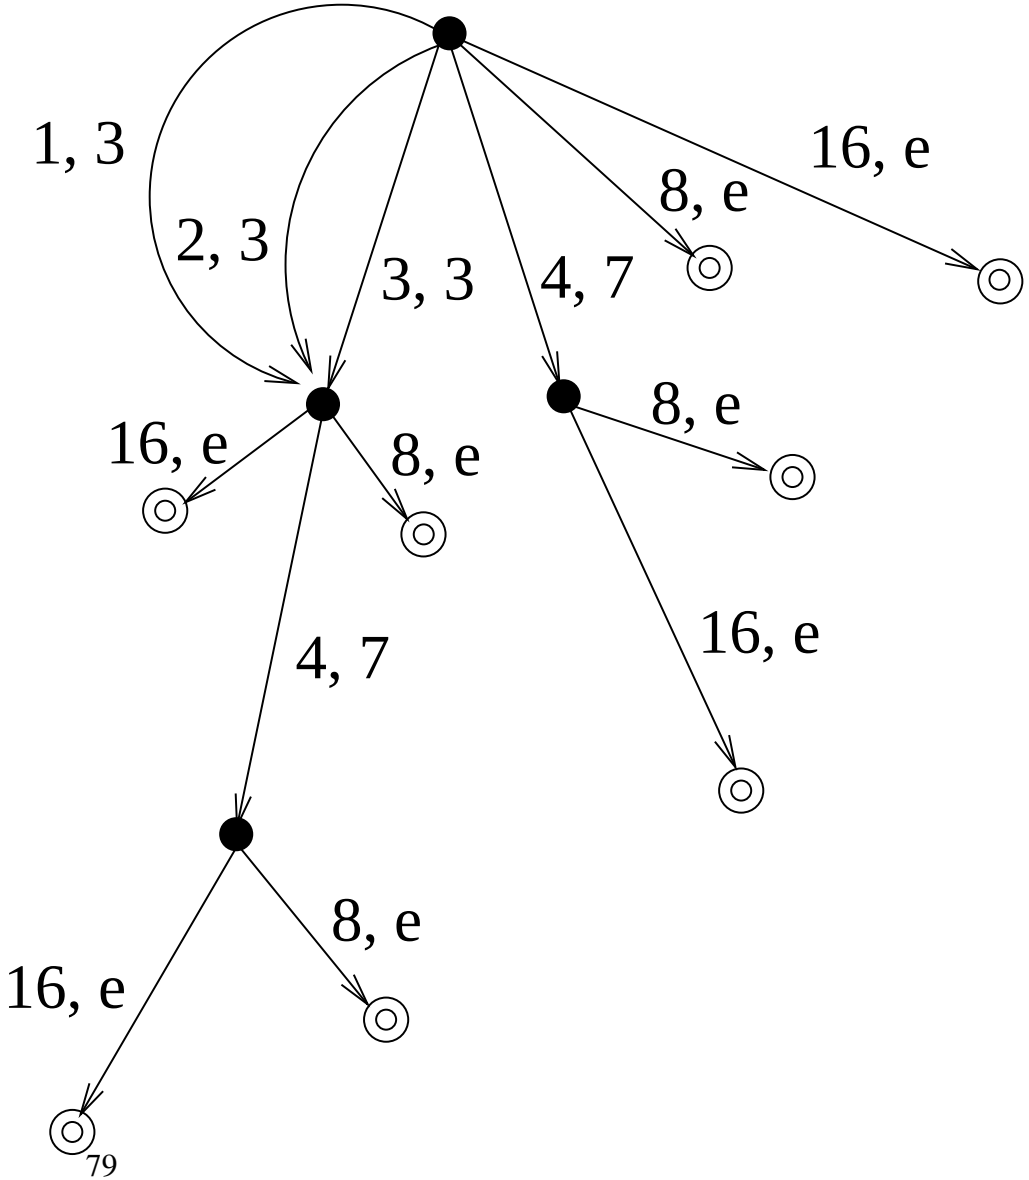

Change to suffix automata: collapse states

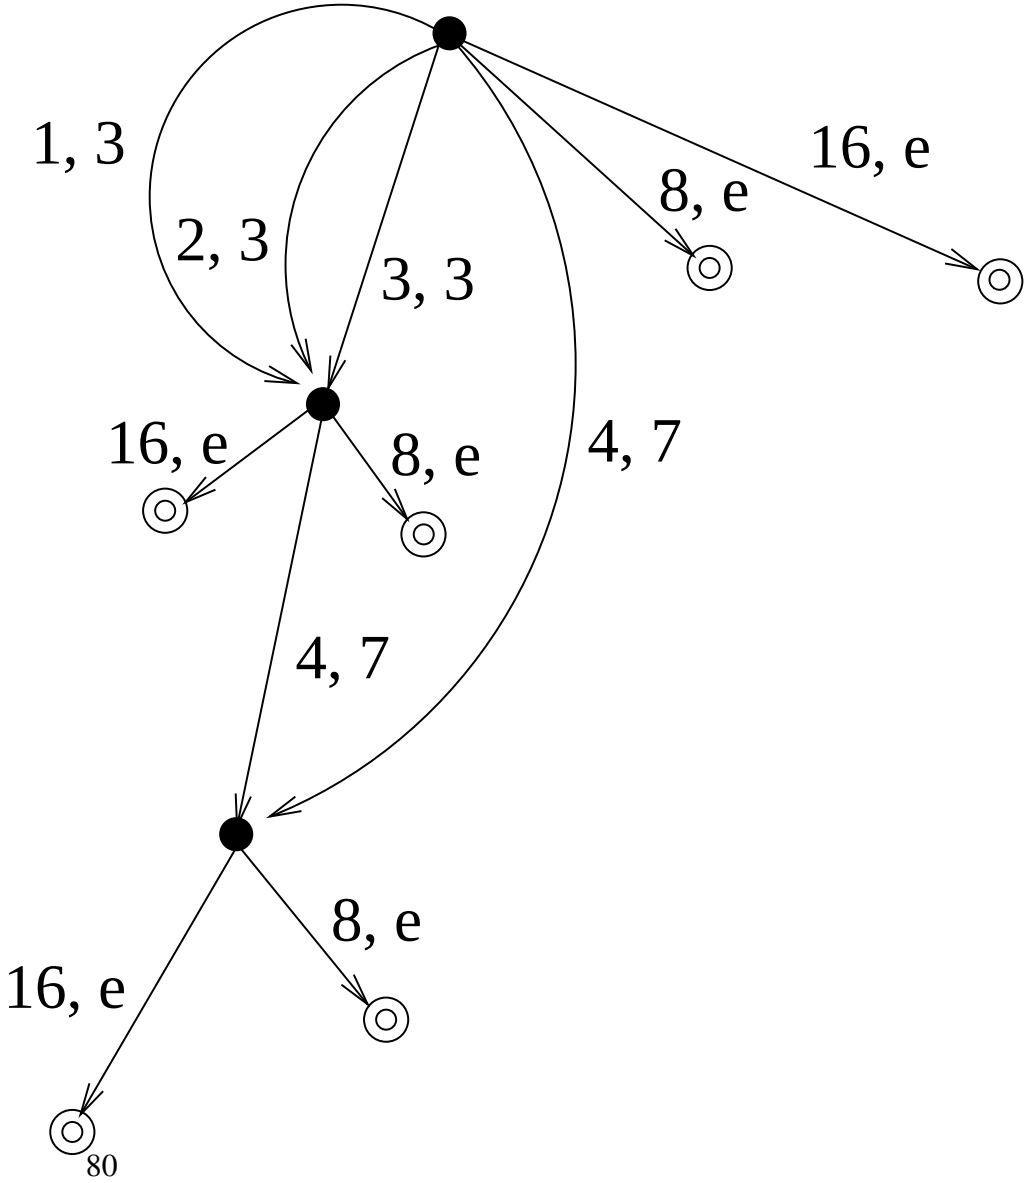

Deterministic suffix automata

- **Large space savings. In this case:**
  - **from 16 final states, 8 internal states, 23 arcs**
  - **to 6 final states, 3 internal states, 11 arcs**
- **Can be achieved with well-known graph minimization algorithms**
- **One particular problem in exact match:**
  - **For a given final state, may have reached there via various paths**
  - **Hence extra work to identify indices of matches**
- **Can get additional space savings with ϵ transitions**

Suffix automata with no ϵ transitions

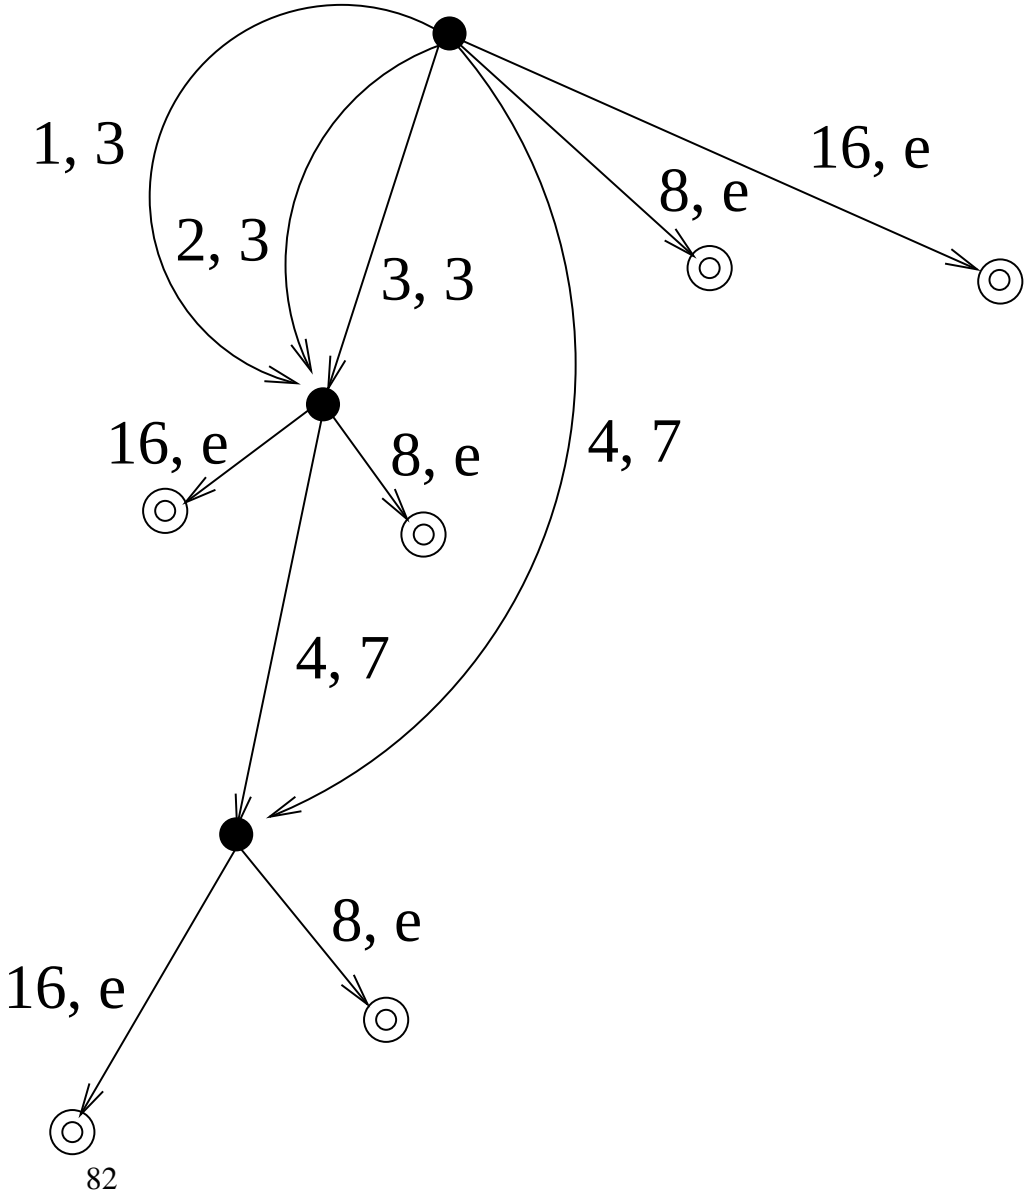

Suffix automata with ϵ transitions

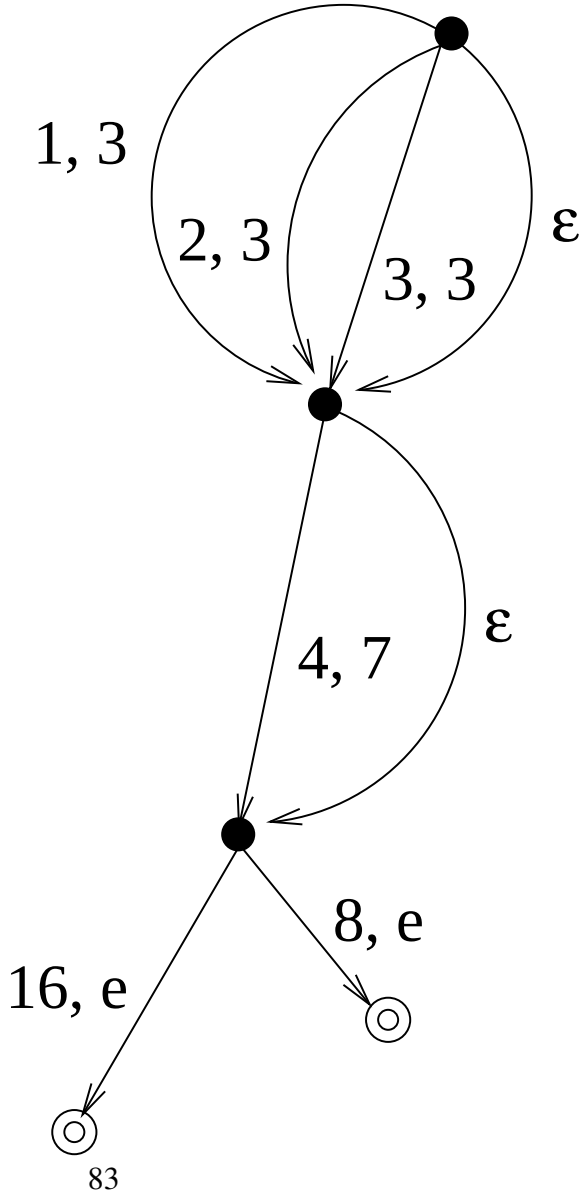

Suffix array

- Another clever space savings idea is the suffix array
- Note that, for text of length m, there are exactly m suffixes
- Thus, in an array of m ints, we can store the suffix indices in alphabetical order
- For any given pattern, all of the suffixes beginning with the pattern would be adjacent in the array
- Much better space usage, not bad exact match time

Suffix array:	1	2	3	4	5	6	7	8	9	0	1	2	3	4	5	6
	x	t	p	y	x	t	p	z	x	t	p	y	x	t	p	r

Suffix idx	Suffix	String pos
1	pr	15
2	pyxtpr	11
3	pyxtpzxtpyxtpr	3
4	pzxtpyxtpr	7
5	r	16
6	tpr	14
7	tpyxtpr	10
8	tpyxtpzxtpyxtpr	2
9	tpzxtpyxtpr	6
10	xtpr	13
11	xtpyxtpr	9
12	xtpyxtpzxtpyxtpr	1
13	xtpzxtpyxtpr	5
14	yxtpr	12
15	yxtpzxtpyxtpr	4
16	zxtpyxtpr	8

Building suffix array

- Each state in suffix tree has at most one arc for each letter
- If we sort the arcs leaving each node in alphabetical order
  - Follow a pre-order (depth first) traversal of the tree
  - Emit indices of leaf nodes as encountered

Example suffix tree

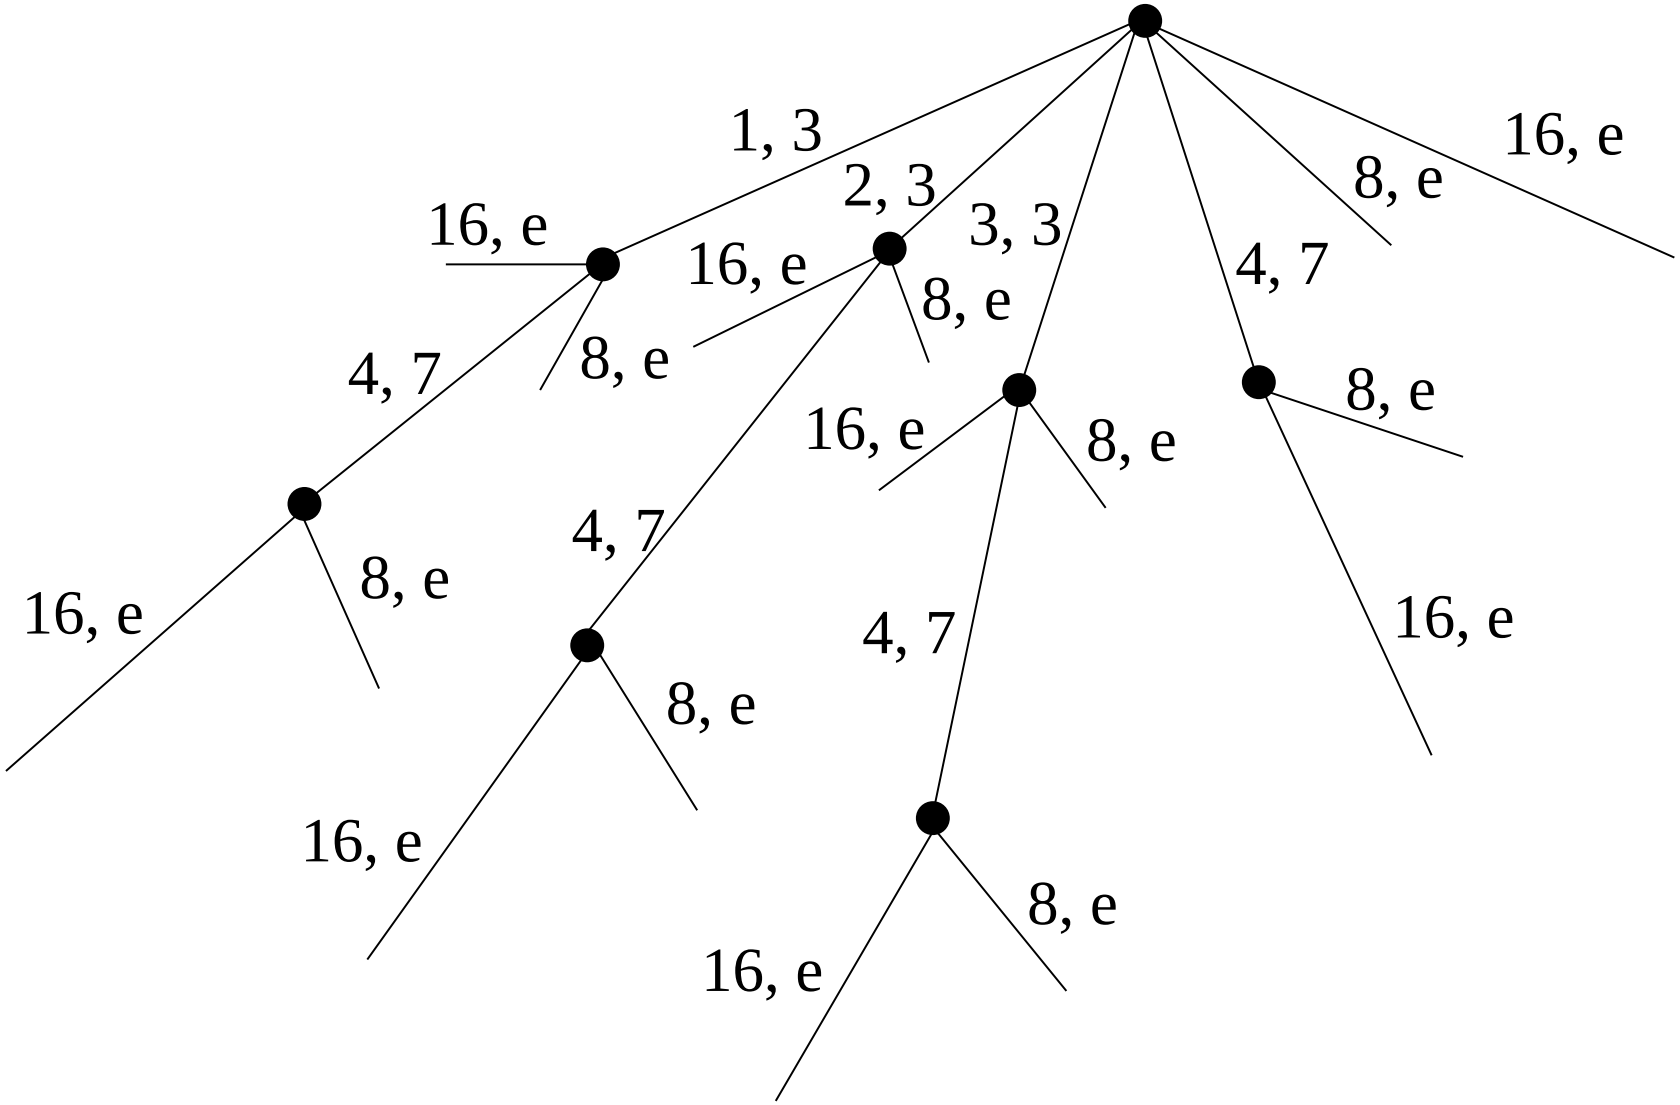

Suffix tree, letters again

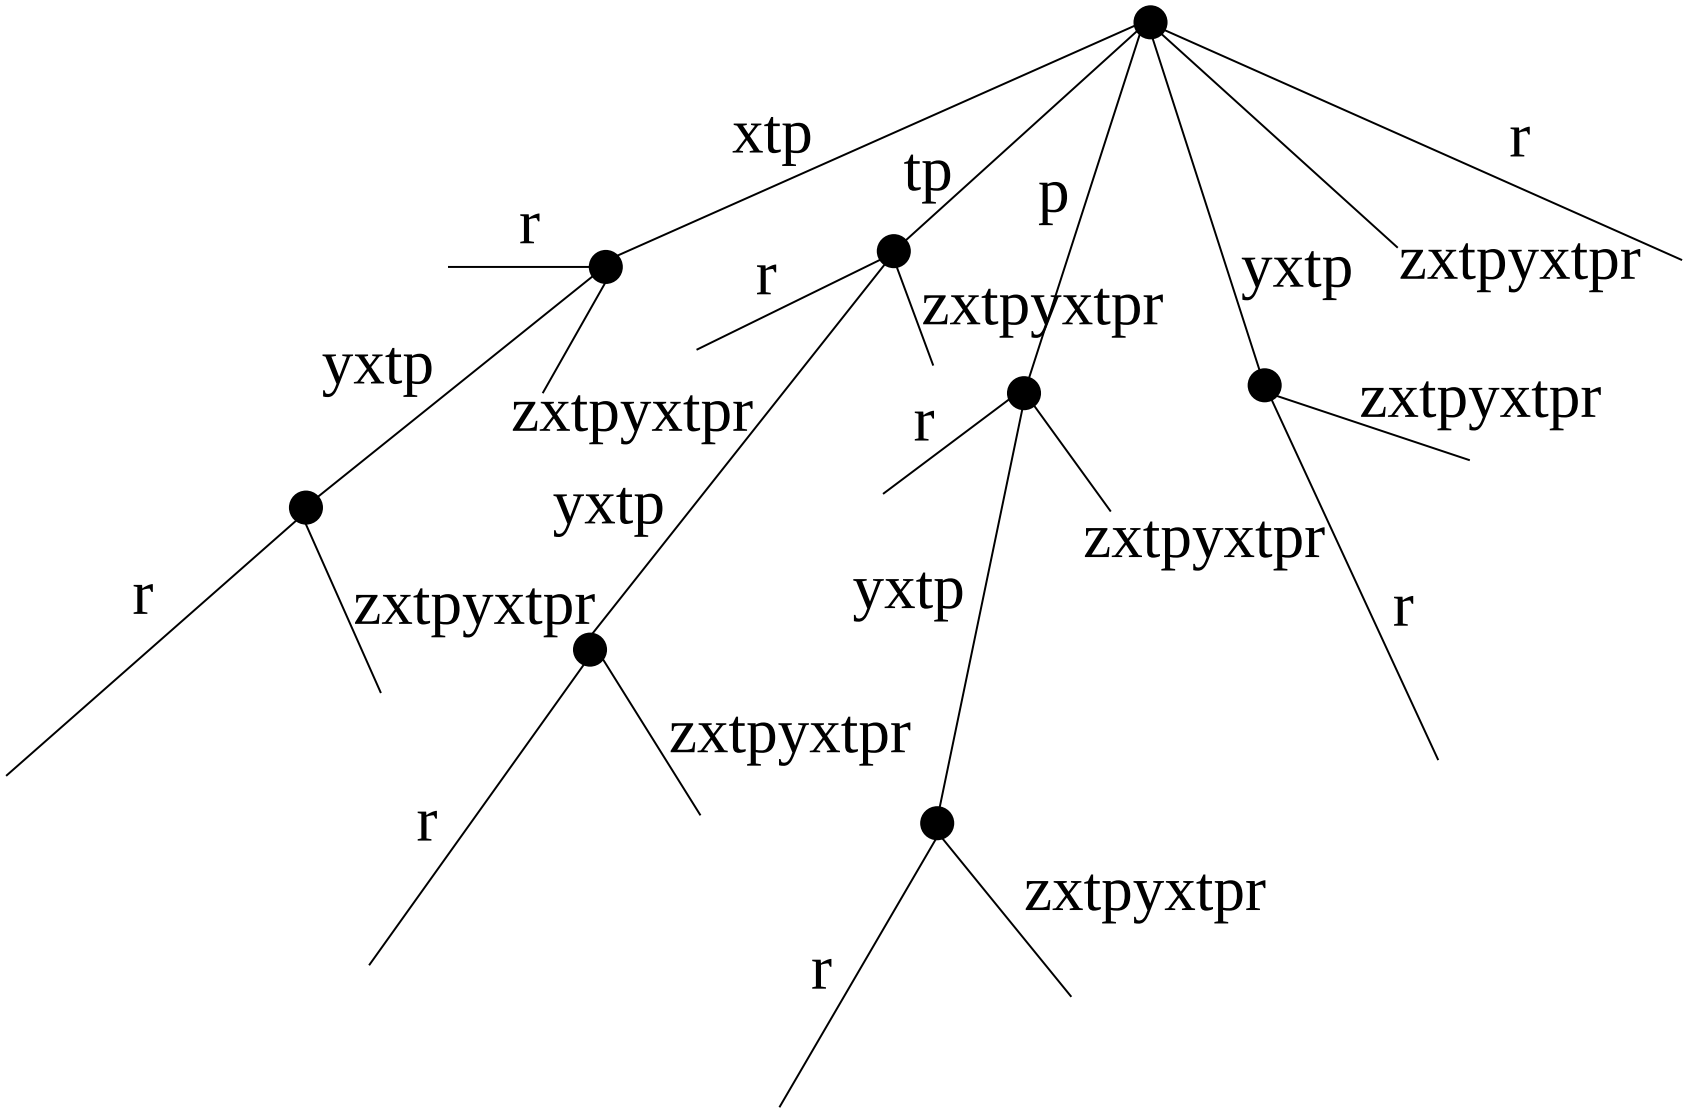

Suffix tree, node order

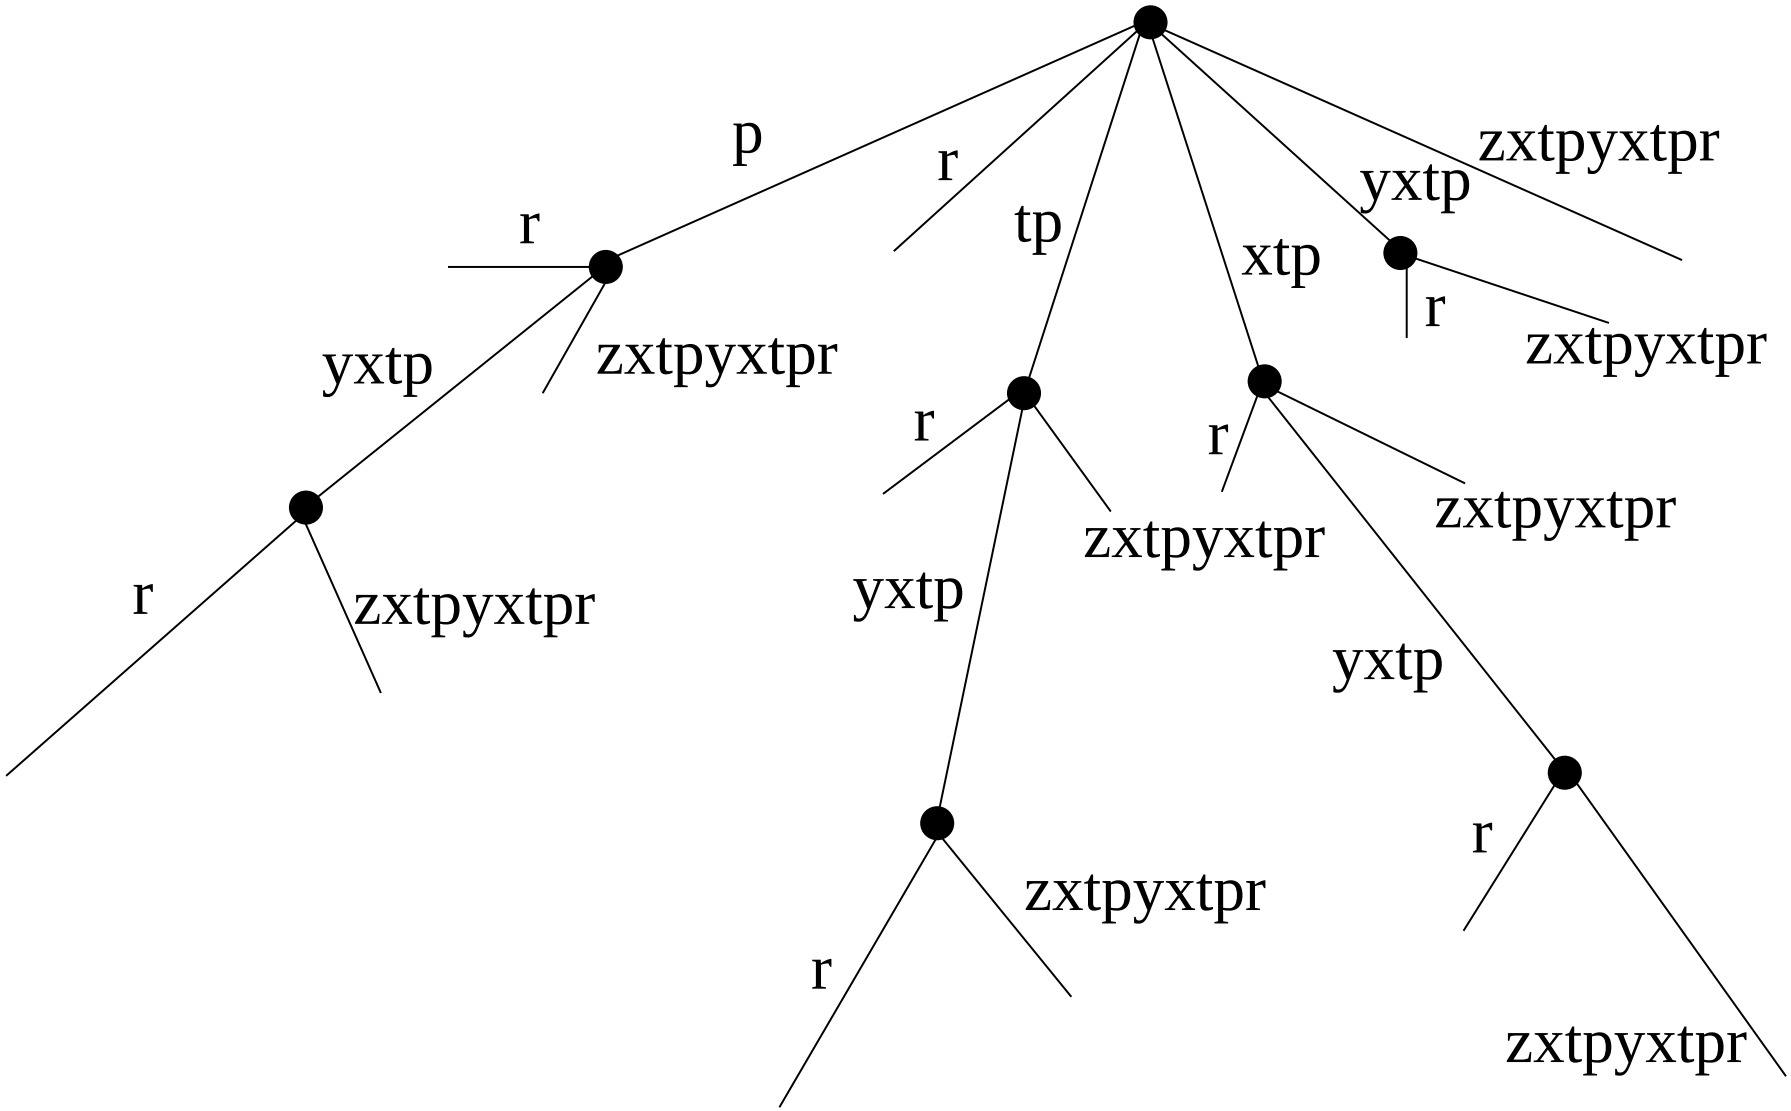

Suffix tree, pre-order enumeration

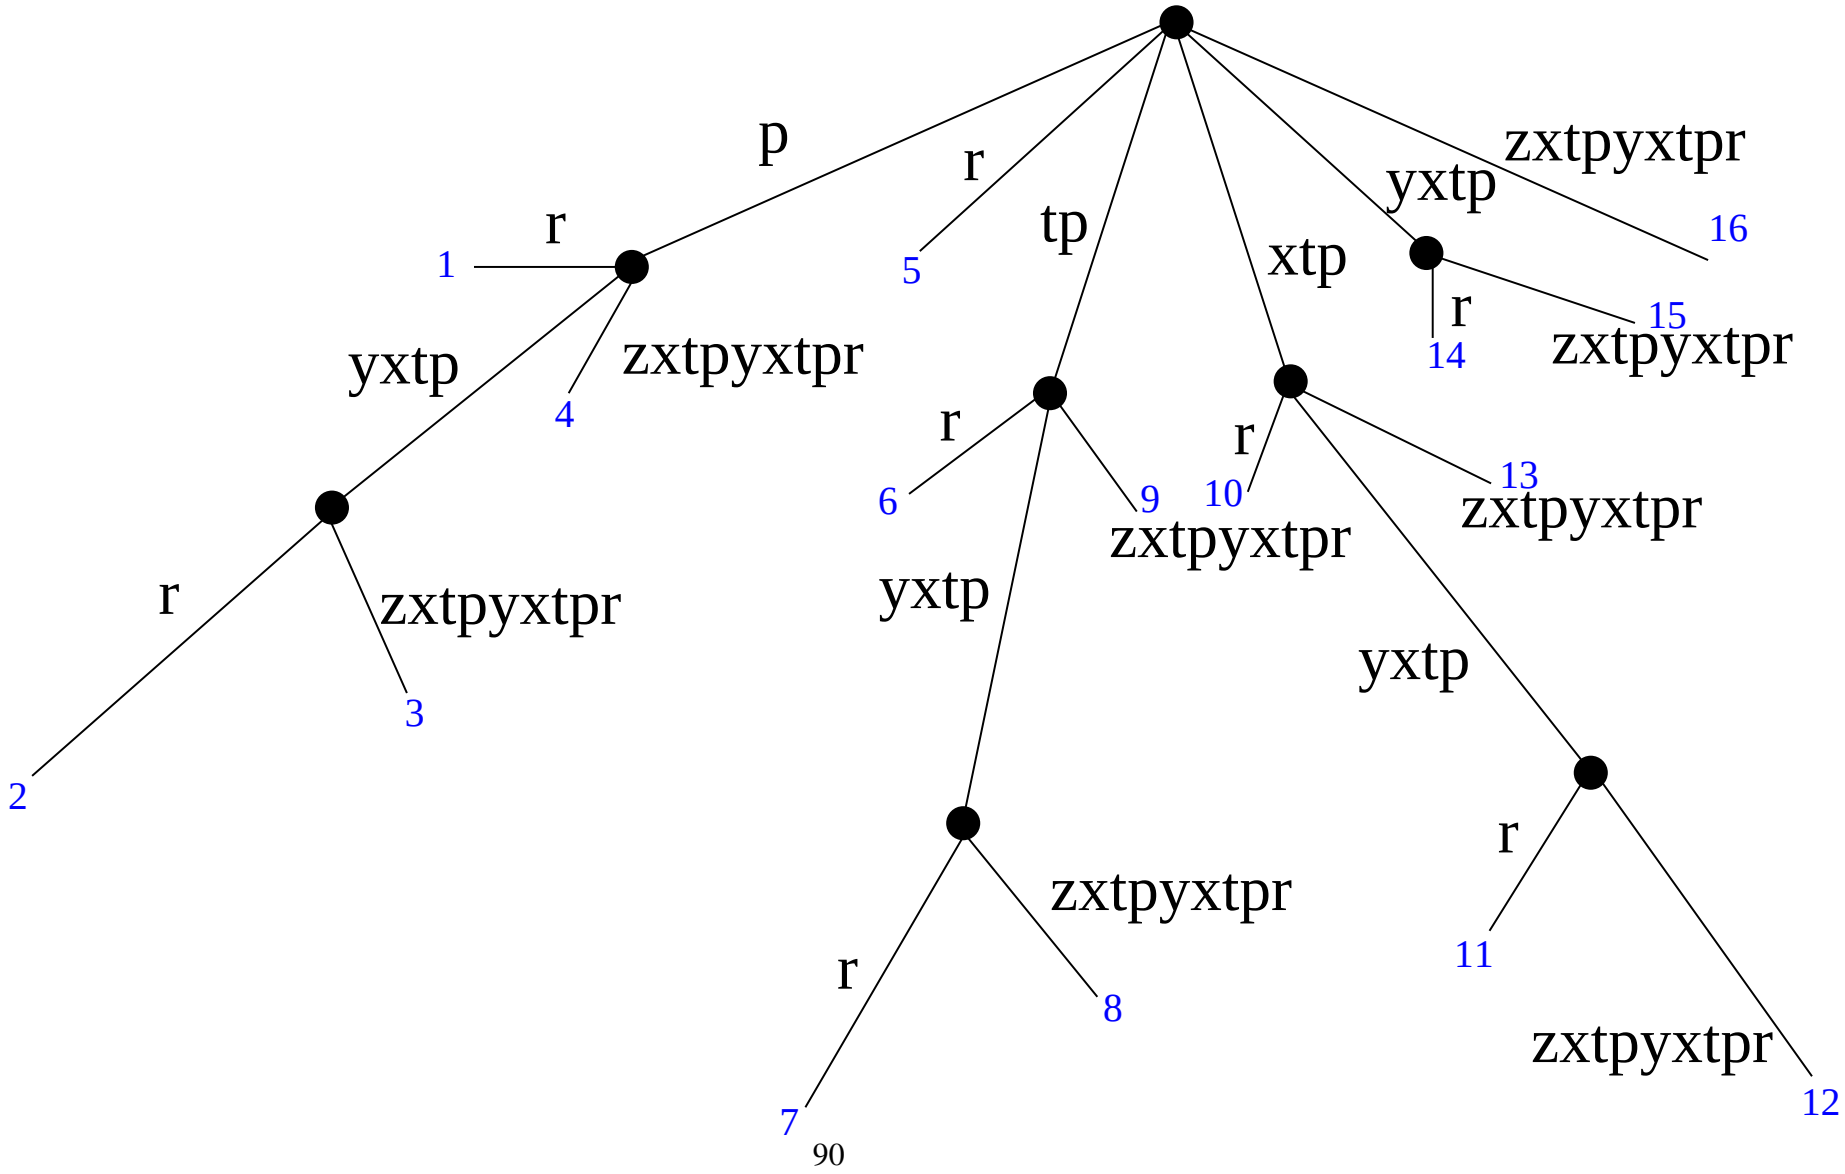

Searching with suffix array

- Perform binary search on the suffix array: $O(n \log m)$
- Can do better than this with clever preprocessing
  - Want to avoid comparisons that are not necessary
- Suppose we have one match in the suffix array
  - The neighbors may or may not be matches
  - How many characters do we need to check?
  - If we know the *longest common prefix* (LCP) of the neighbors, we know where to begin
- Can build a binary search tree with pre-compiled LCP values at each node

Longest common prefix

- Have enough information to calculate pair-wise LCP in suffix tree
- In pre-order traversal of suffix tree to build the suffix array:
  - The LCP between suffix idx i and suffix idx $i+1$ is the depth in the tree of the lowest common ancestor node of the suffixes
  - (Recall that the depth of a node is the length of the string labeling arcs reaching that node)
- Critical lemma for easy LCP accumulation in binary tree:

  For any i, j such that $j > i+1$,

$$\text{LCP}(i, j) = \min_{i \leq k \leq j-1} \text{LCP}(k, k+1)$$

- In binary tree, take the smallest value from the two children

Suffix tree, pre-order enumeration

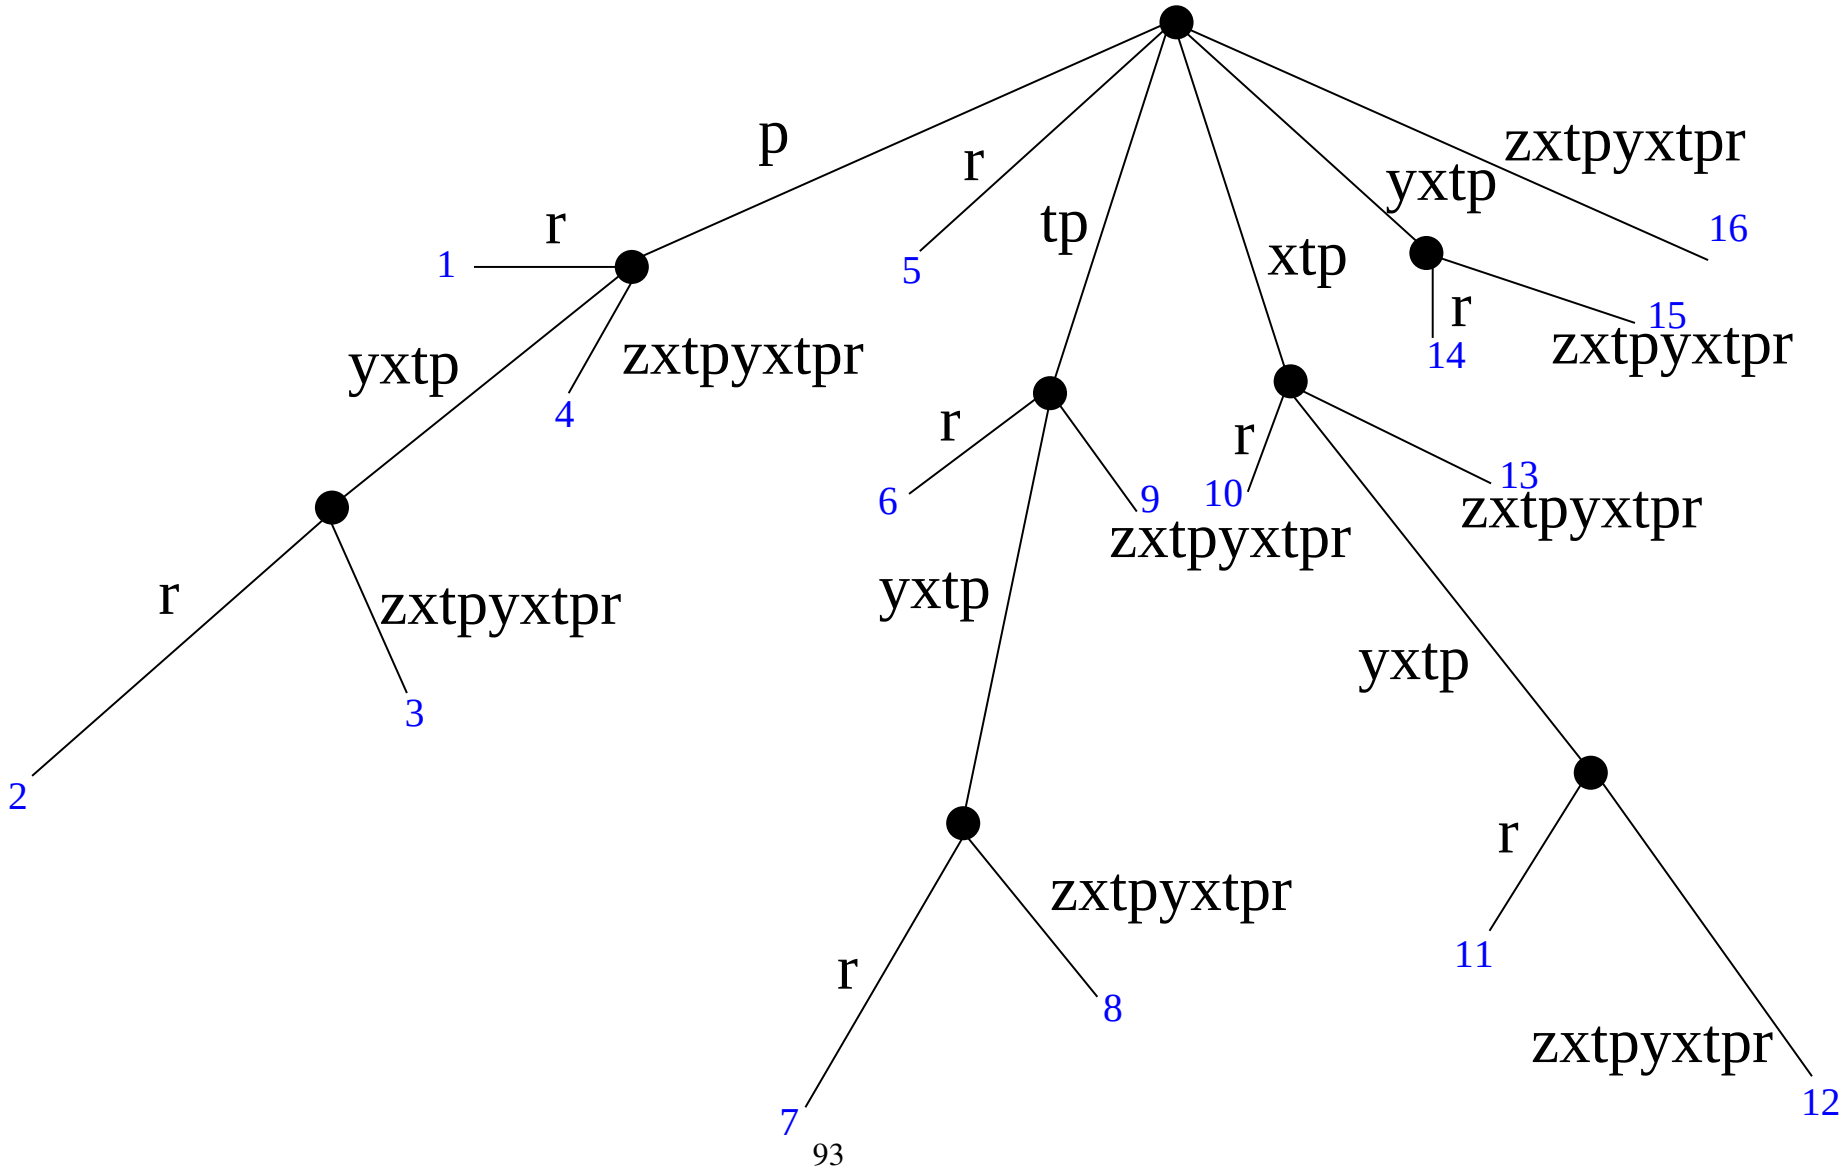

Suffix array:	**1**	**2**	**3**	**4**	**5**	**6**	**7**	**8**	**9**	**0**	**1**	**2**	**3**	**4**	**5**	**6**
	x	**t**	**p**	**y**	**x**	**t**	**p**	**z**	**x**	**t**	**p**	**y**	**x**	**t**	**p**	**r**

Suffix idx	Suffix	String pos	LCP($i, i+1$)
1	pr	15	1
2	pyxtpr	11	5
3	pyxtpzxtpyxtpr	3	1
4	pzxtpyxtpr	7	0
5	r	16	0
6	tpr	14	2
7	tpyxtpr	10	6
8	tpyxtpzxtpyxtpr	2	2
9	tpzxtpyxtpr	6	0
10	xtpr	13	3
11	xtpyxtpr	9	7
12	xtpyxtpzxtpyxtpr	1	3
13	xtpzxtpyxtpr	5	0
14	yxtpr	12	4
15	yxtpzxtpyxtpr	4	0
16	zxtpyxtpr	8	

Building binary tree

- Index node with first and last suffix array indices
- Split in half, point to left and right children
- Need to store two LCP numbers at each node
  - Note that the internal LCP does not include rightmost suffix
  - Will need to include rightmost suffix if combining as left child
  - Keep LCP both with and without rightmost suffix
- Use LCP to avoid comparing parts of pattern already found to match

Binary tree

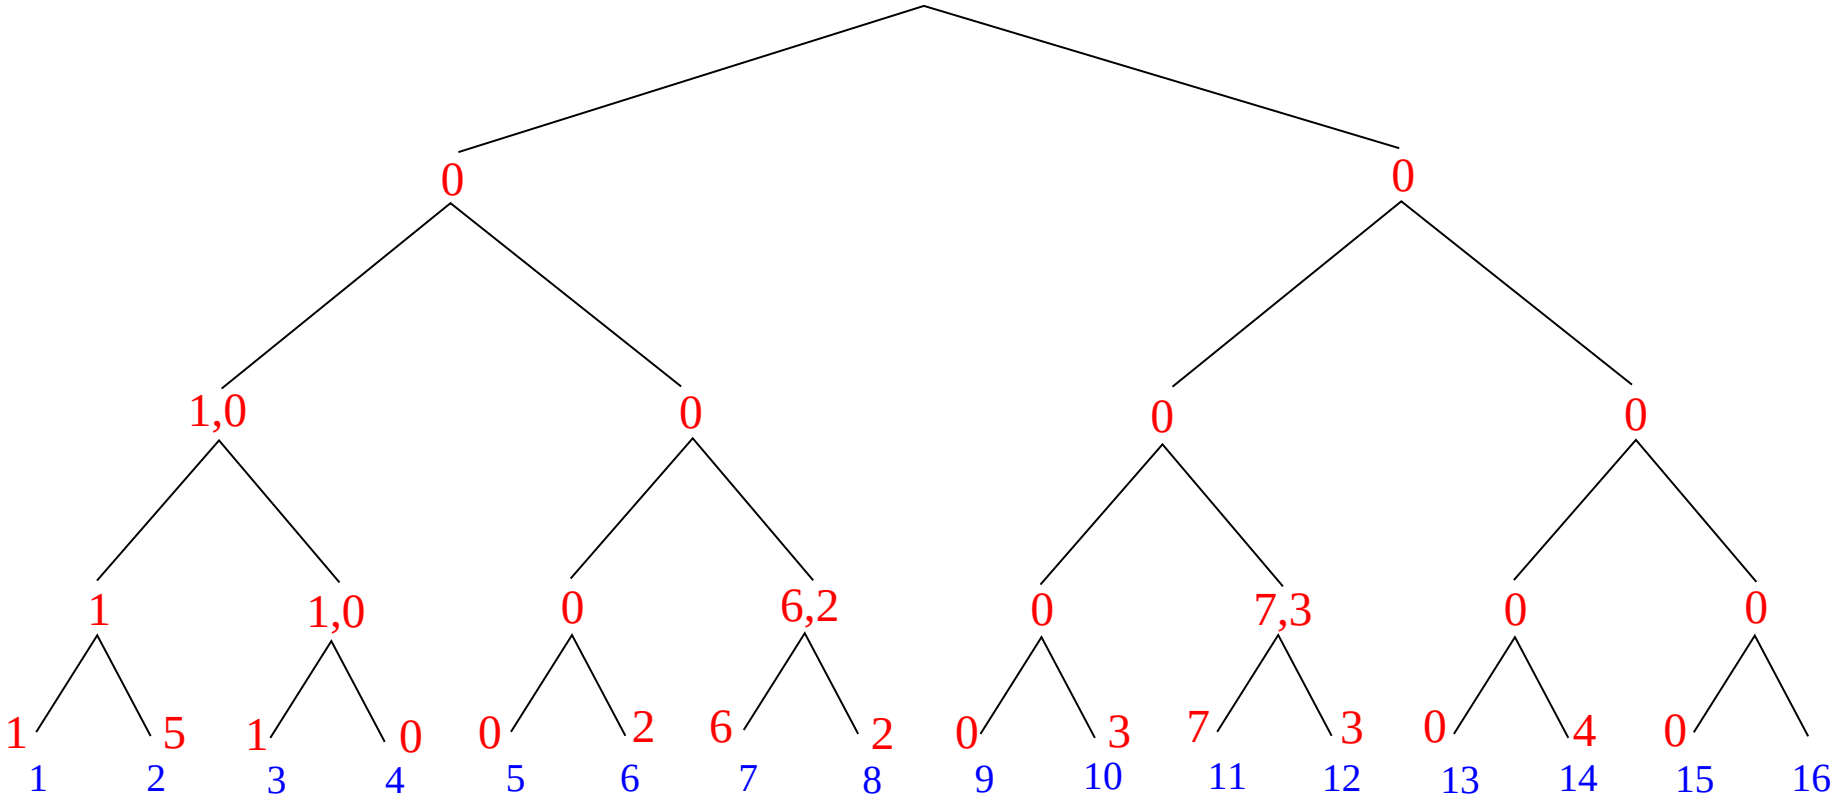
0
0
1,0
0
0
0
1
1,0
0
6,2
0
7,3
0
0
1
5
1
0
0
2
6
2
0
3
7
3
0
4
0
1
2
3
4
5
6
7
8
9
10
11
12
13
14
15
16

Binary search

- Start at root
- Compare pattern with rightmost suffix of left child of current node
  - If pattern less-than suffix, go to left node and continue
  - If suffix less-than pattern, go to right node and continue
  - If pattern matches suffix, move down tree (rightmost child)
    - Descend to first node with LCP $\geq n$ (length of pattern)
    - All suffixes beneath node match pattern
    - move to the left of the leftmost and right of rightmost to find further matches